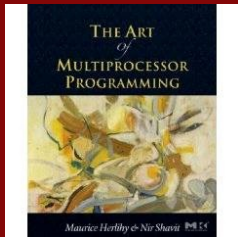




MPCS 52060 - Parallel Programming

M4: Concurrent Data Structures (Part 1)



Original slides from “The Art of Multiprocessor Programming” by Maurice Herlihy & Nir Shavit with modifications by Lamont Samuels

More Low-Level Synchronization Primitives

Motivation for Semaphores

- What if we wanted to control access to shared resource?
 - For example, A system can only handle a certain number of users concurrently signed on. After the maximum number of logged in users is reached, then others must wait until others logout.
- A **semaphore** is a synchronization primitive used to control access to a shared resource by multiple threads.
 - It has a capacity (c), which allows for having at most (c) threads in a critical section. Unlike with locks, where only one thread can be a critical section at a time.
 - You can also think of them as a way to control how many resources are available of a particular entity by allowing resources to be concurrently acquired and released in a safe way.
 - Tracks only how many resources are free; it does not keep track of which of the resources are free

Semaphore Pseudo-Implementation

- The capacity variable of a semaphore is an integer value that cannot be directly accessed.
 - Go does not have a semaphore construct so the below examples are pseudocode similar to the implementations of semaphores in other languages:
- **Creation:** Must initialize it to some capacity integer value

```
var sema *Semaphore  
// NewSemaphore allocates and  
// returns a *Semaphore with its  
// internal capacity initialized  
sema = NewSemaphore(0)
```

- **Behaviors:** It has two main operations (methods in our case) that modify the integer capacity value

```
// Decrement semaphore  
sema.Down()  
// Increment semaphore  
sema.Up()
```

Semaphore Pseudo-Implementation

```
func (s *Semaphore) Down() {  
    // Wait if the value of semaphore s is less than or equal to 0  
    // Decrement the value of semaphore s by one  
}  
func (s *Semaphore) Up() {  
    //Increment the value of semaphore s by 1  
    //If there are 1 or more threads waiting, wake one up  
}
```

Semaphores & Mutual Exclusion

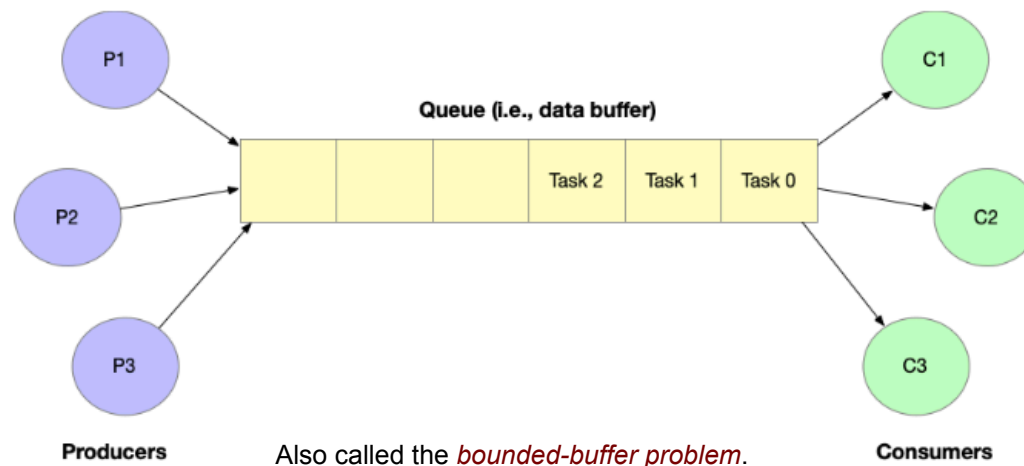
- There are two well known types of semaphores
 - **Binary semaphore** - acts like a mutex by setting its capacity initial value to 1.

```
var mutex_sema *Semaphore
mutex_sema = NewSemaphore(1)
mutex_sema.Down()
//critical section
mutex_sema.Up()
```

- **Counting semaphore** - initialize the semaphore to be equal to the number of available resources.

Real-World Example: Producer and Consumer Problem

- One or more threads generate tasks (**producers**) and one or more threads receive and process them (**consumers**).
- Producers and consumers communicate using a queue of maximum size N and must adhere to the following conditions
 - Consumers must wait for a producer to produce a task if the queue is empty.
 - Producer must wait for the consumer to consume a task if the queue is full.



Real-World Example: Producer and Consumer Problem

//Producer

for {

 //Generate Task

 sema_emptyCount.Down()

 sema_mutex.Down()

 //Put task in Queue

 sema_mutex.Up()

 sema_fullCount.Up()

}

//Consumer

for {

 sema_fullCount.Down()

 sema_mutex.Down()

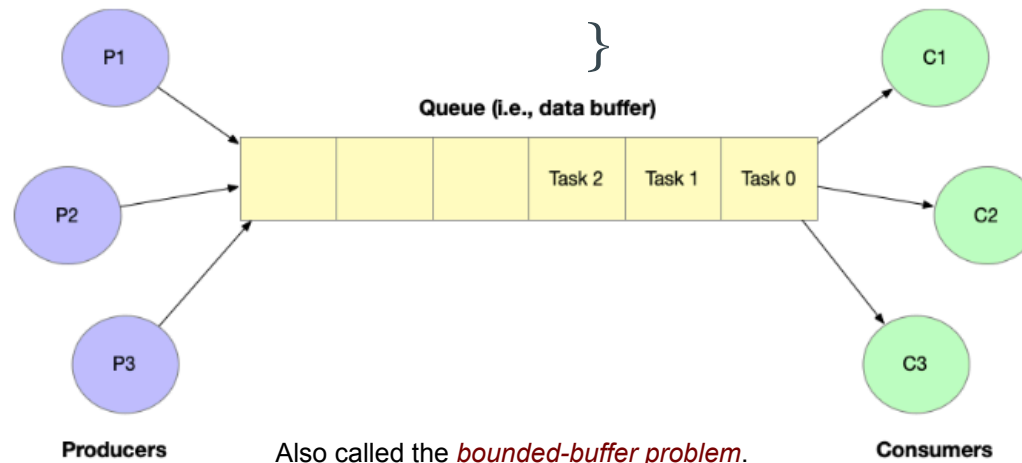
 //Remove task from Queue

 sema_mutex.Up()

 sema_emptyCount.Up()

 //Process Task

}



Condition Variables

- A data object that allows a thread to suspend execution until a certain event or condition occurs.
- When the event or condition occurs another thread can signal the thread to “wake up.”
- A condition variable is always associated with a mutex.

```
// lock mutex
if (condition has occurred) {
    // signal thread(s);
} else {
    // 1. Wait until another thread signals to wake up by unlocking
    // the mutex and block (e.g., sleep, or spin etc.);
    // 2. After the signal happens then the thread wakes, requires the lock
    // and checks to make sure the condition is still true.
}
// unlock mutex
```

Condition Variables in Go

- `sync.Cond` represents conditional variables in Go.
- **Creation:** `NewCond(l Locker) *Cond`
- **Operations** on condition variables:
 - `func (c *Cond) Wait()`: suspends the calling thread and releases the monitor lock. When it resumes, reacquire the lock. Called when condition is not true.
 - `func (c *Cond) Signal()`: resumes one thread waiting in `wait()` if any. Called when condition becomes true and wants to wake up one waiting thread.
 - `func (c *Cond) Broadcast()`: resumes all threads waiting in `wait()`. Called when condition becomes true and wants to wake up all waiting threads.

Demo: Condition Variables

Concurrent Data Structures

Concurrent Data Structures

- We assume
 - shared-memory multiprocessors environment
 - concurrently execute multiple threads which communicate and synchronize through data structures in shared memory

Concurrent Data Structures

- Far more difficult to design than sequential ones
 - Correctness
 - Primary source of difficulty is concurrency
 - The steps of different threads can be interleaved arbitrarily
 - Scalability (performance)
- We will look at
 - Concurrent Linked List/Queue/Stack

Main performance issue of lock based system

- Sequential bottleneck
 - At any point in time, at most one lock-protected operation is doing useful work.
- Memory contention
 - Overhead in traffic as a result of multiple threads concurrently attempting to access the same memory location.
- Blocking
 - If thread that currently holds the lock is delayed, then all other threads attempting to access are also delayed.
 - Implementation of locks is known as a blocking algorithm
 - Consider non-blocking (lock-free) algorithm

Nonblocking algorithms

- implemented by a hardware operation
 - atomically combines a load and a store
 - Ex) compare-and-swap(CAS)
- lock-free
 - if there is guaranteed system-wide progress;
 - while a given thread might be blocked by other threads, all CPUs can continue doing other useful work without stalls.
- wait-free
 - if there is also guaranteed per-thread progress.
 - in addition to all CPUs continuing to do useful work, no computation can ever be blocked by another computation.

Linked List

- Illustrate these patterns ...
- Using a list-based Set
 - Common application
 - Building block for other apps

Set Interface

- Unordered collection of items
- No duplicates
- Methods
 - add(x) put x in set
 - remove(x) take x out of set
 - contains(x) tests if x in set


List-Based Sets

```
public interface Set<T> {  
    public boolean add(T x);  
    public boolean remove(T x);  
    public boolean contains(T x);  
}
```

List Node

```
public class Node {  
    public T item;    // item of interest  
    public int key;    // usually hash code  
    public Node next; // reference to next node  
}
```

The List-Based Set



Sorted with Sentinel nodes
(min & max possible keys)

Sequential List Based Set

Add()




Remove()




Sequential List Based Set


Add()



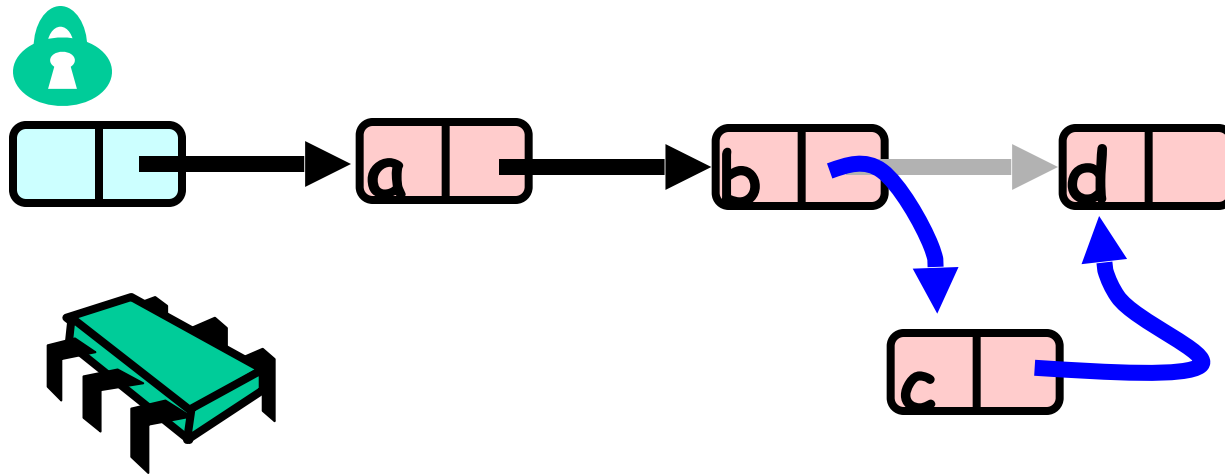
Remove()




Course Grained Locking



Course Grained Locking



Course Grained Locking



Simple but **hotspot + bottleneck**

Coarse-Grained Synchronization

- Sequential bottleneck
 - Threads "stand in line"
- Adding more threads
 - Does not improve throughput
 - Struggle to keep it from getting worse
- So why even use a multiprocessor?
 - Well, some apps inherently parallel ...

Coarse-Grained Synchronization (Linked List)

```
public class CoarseList<T> {  
    private Node head;  
    private Node tail;  
    private Lock lock = new ReentrantLock();  
  
    public CoarseList() {  
        // Add sentinels to start and end  
        head = new Node(Integer.MIN_VALUE);  
        tail = new Node(Integer.MAX_VALUE);  
        head.next = this.tail;  
    }  
}
```

```
public boolean add(T item) {  
    Node pred, curr;  
    int key = item.hashCode();  
    lock.lock();  
    try {  
        pred = head;  
        curr = pred.next;  
        while (curr.key < key) {  
            pred = curr;  
            curr = curr.next;  
        }  
        if (key == curr.key) {  
            return false;  
        } else {  
            Node node = new Node(item);  
            node.next = curr;  
            pred.next = node;  
            return true;  
        }  
    } finally {  
        lock.unlock();  
    }  
}
```

```

public boolean remove(T item) {
    Node pred, curr;
    int key = item.hashCode();
    lock.lock();
    try {
        pred = this.head;
        curr = pred.next;
        while (curr.key < key) {
            pred = curr;
            curr = curr.next;
        }
        if (key == curr.key)
            pred.next = curr.next;
        return true;
    } else {
        return false;
    }
} finally {
    lock.unlock();
}
}

```

```

public boolean contains(T item) {
    Node pred, curr;
    int key = item.hashCode();
    lock.lock();
    try {
        pred = head;
        curr = pred.next;
        while (curr.key < key) {
            pred = curr;
            curr = curr.next;
        }
        return (key == curr.key);
    } finally {
        lock.unlock();
    }
}

```

Coarse-Grained Locking

- Easy, same as synchronized methods
 - "One lock to rule them all ..."
- Simple, clearly correct
 - Deserves respect!
- Works poorly with contention

Performance Improvement

- For highly-concurrent objects
- Goal:
 - Concurrent access
 - More threads, more throughput

First: Fine-Grained Synchronization

- Instead of using a single lock ..
- Split object into
 - Independently-synchronized components
- Methods conflict when they access
 - The same component ...
 - At the same time

Second: Optimistic Synchronization

- Search without locking ...
- If you find it, lock and check ...
 - OK: we are done
 - Oops: start over
- Evaluation
 - Usually cheaper than locking
 - Mistakes are expensive

Third: Lazy Synchronization

- Postpone hard work
- Removing components is tricky
 - Logical removal
 - Mark component to be deleted
 - Physical removal
 - Do what needs to be done

Fourth: Lock-Free Synchronization

- Don't use locks at all
 - Use `compareAndSet()` & relatives ...
- Advantages
 - No Scheduler Assumptions/Support
- Disadvantages
 - Complex
 - Sometimes high overhead

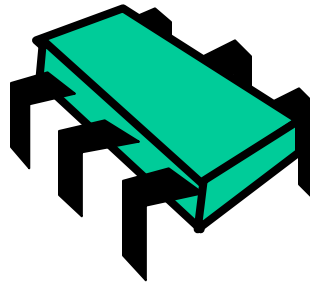
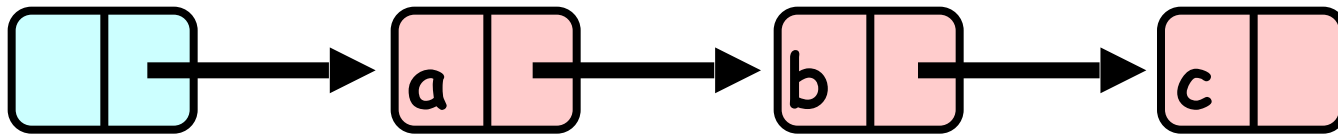
Fine-grained Locking

- Requires careful thought
- Split object into pieces
 - Each piece has own lock
 - Methods that work on disjoint pieces need not exclude each other

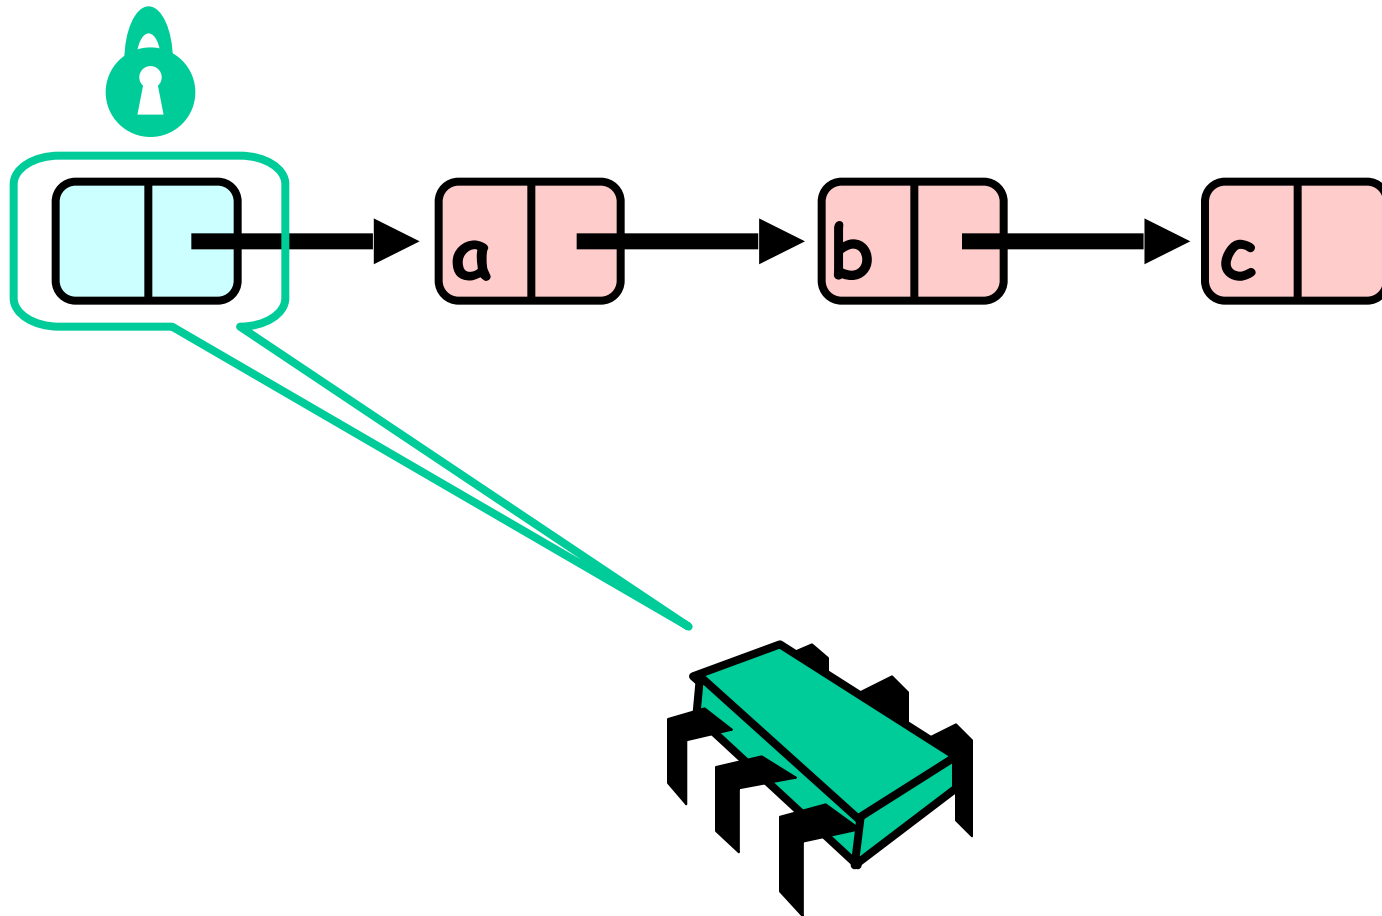
Fine-grained Locking

- Use multiple locks of small granularity to protect different parts of the data structure
- Goal
 - To allow concurrent operations to proceed in parallel when they do not access the same parts of the data structure

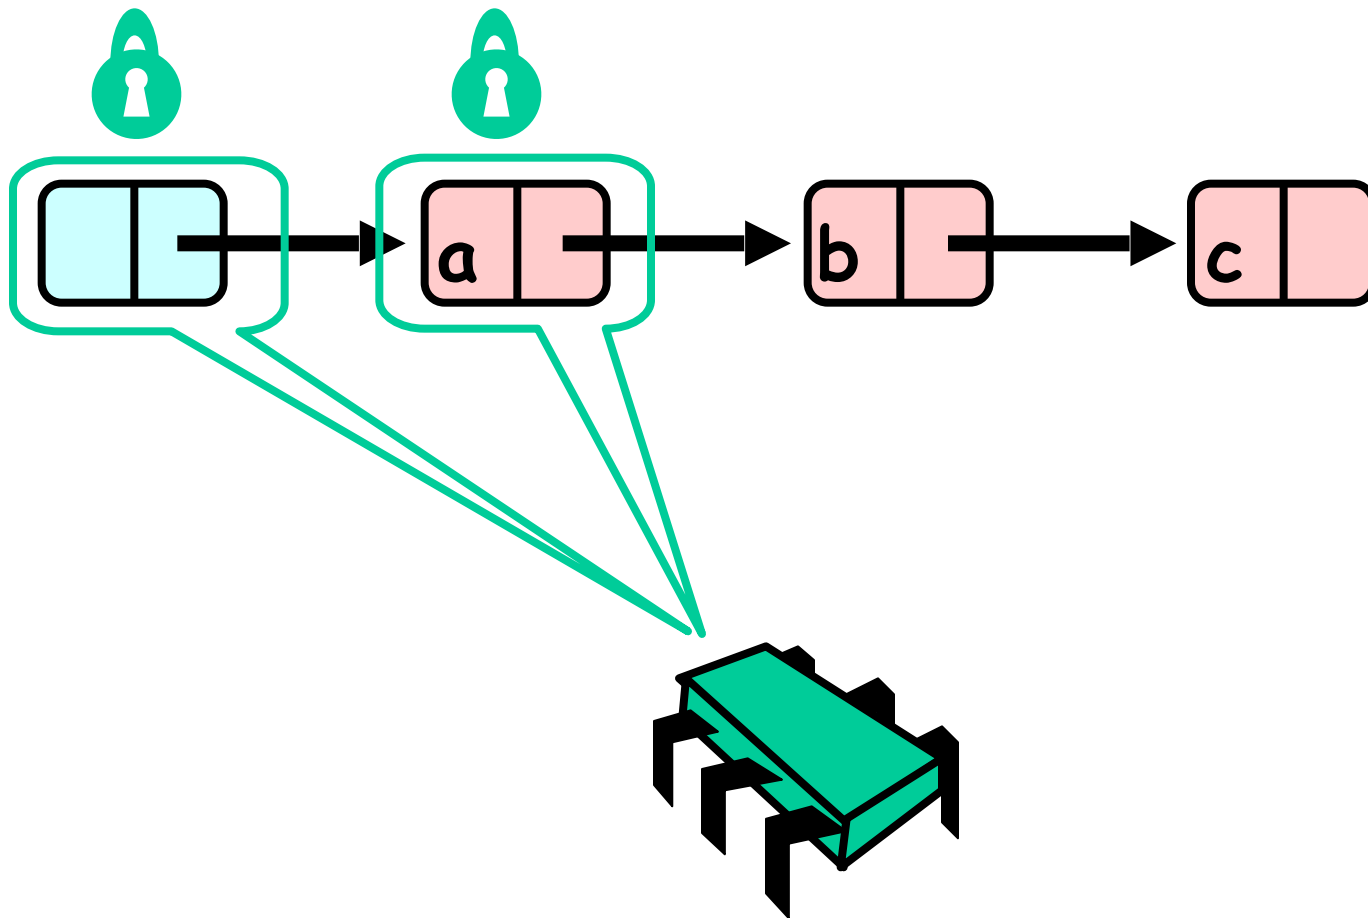
Hand-over-Hand locking



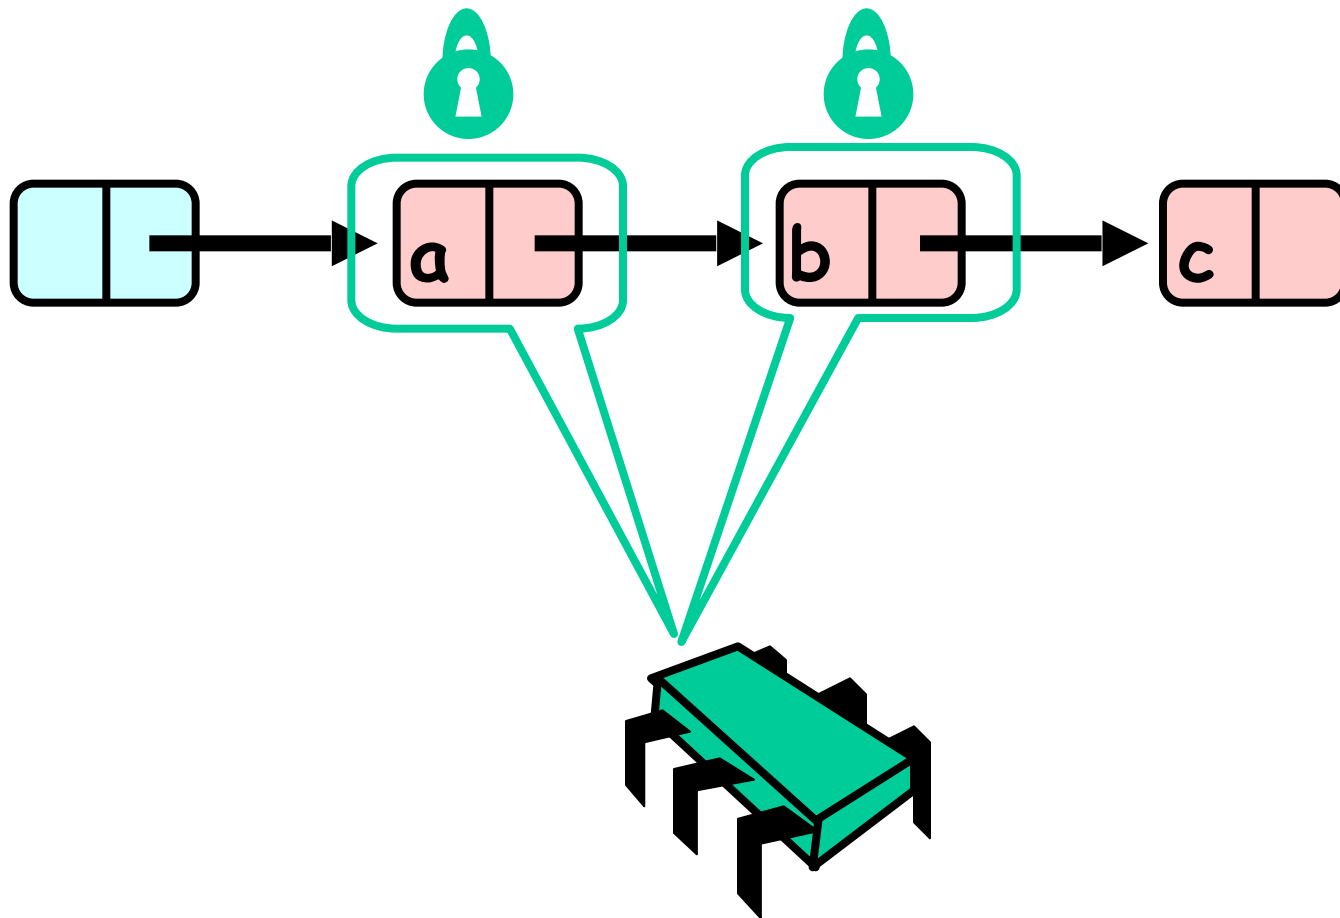
Hand-over-Hand locking




Hand-over-Hand locking




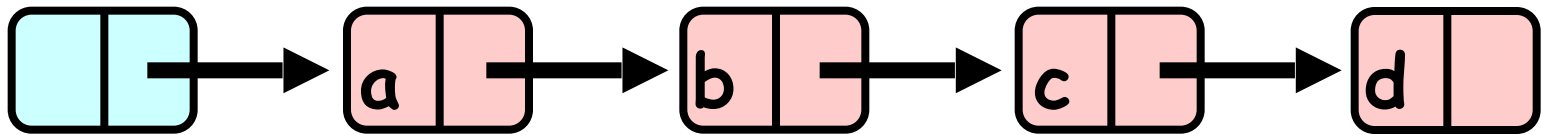
Hand-over-Hand locking



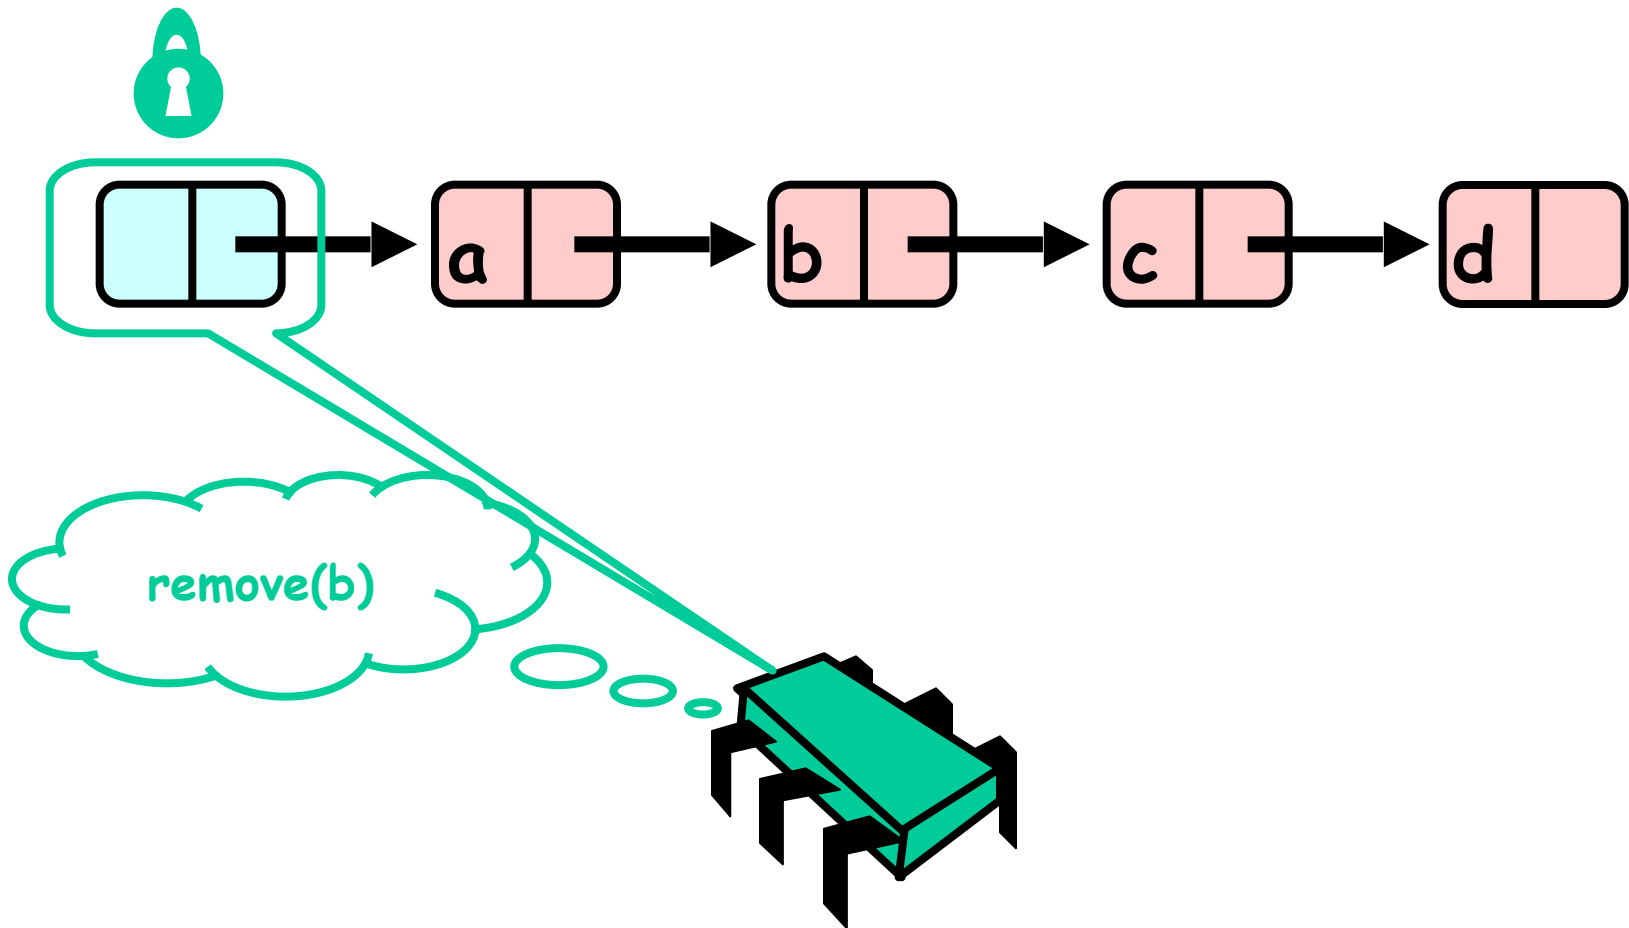
Hand-over-Hand locking



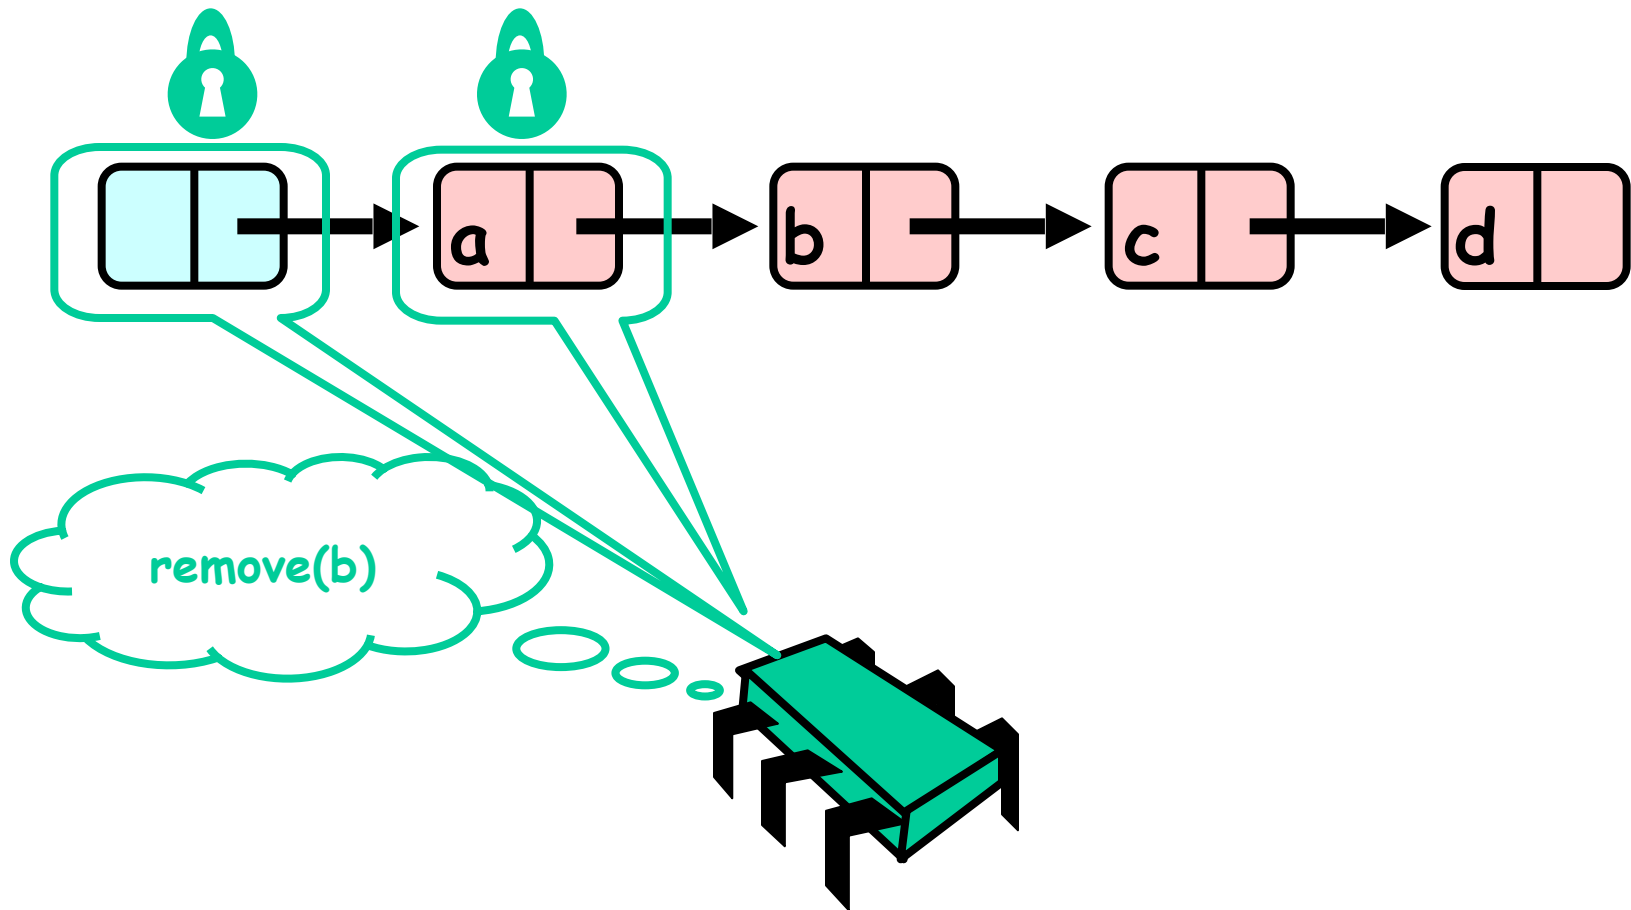
Removing a Node



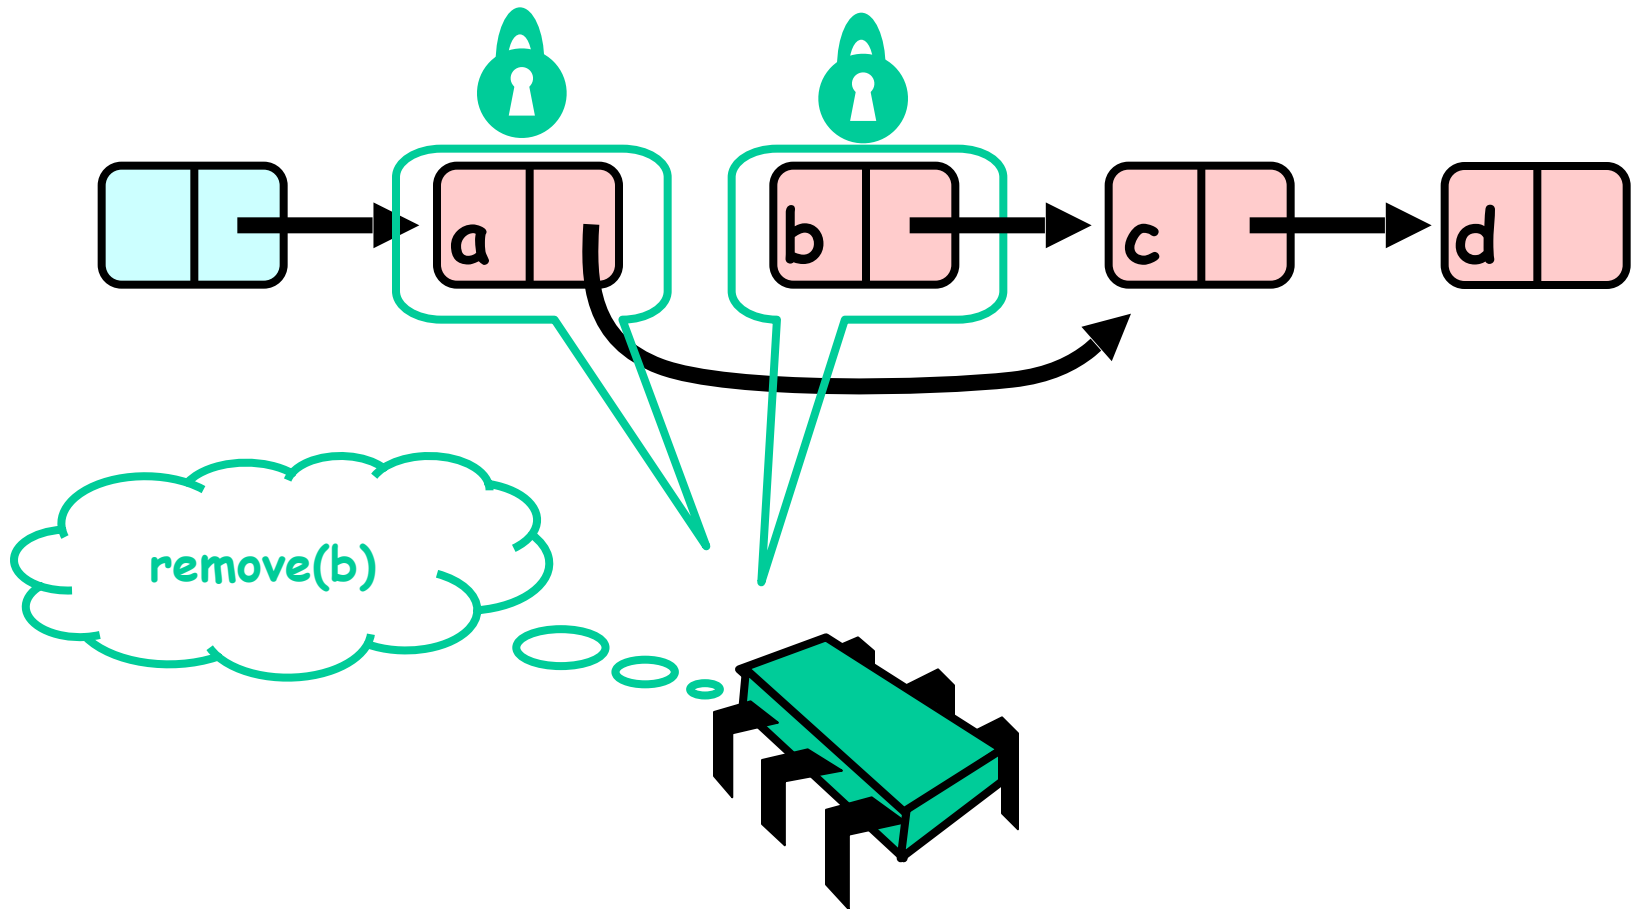
Removing a Node



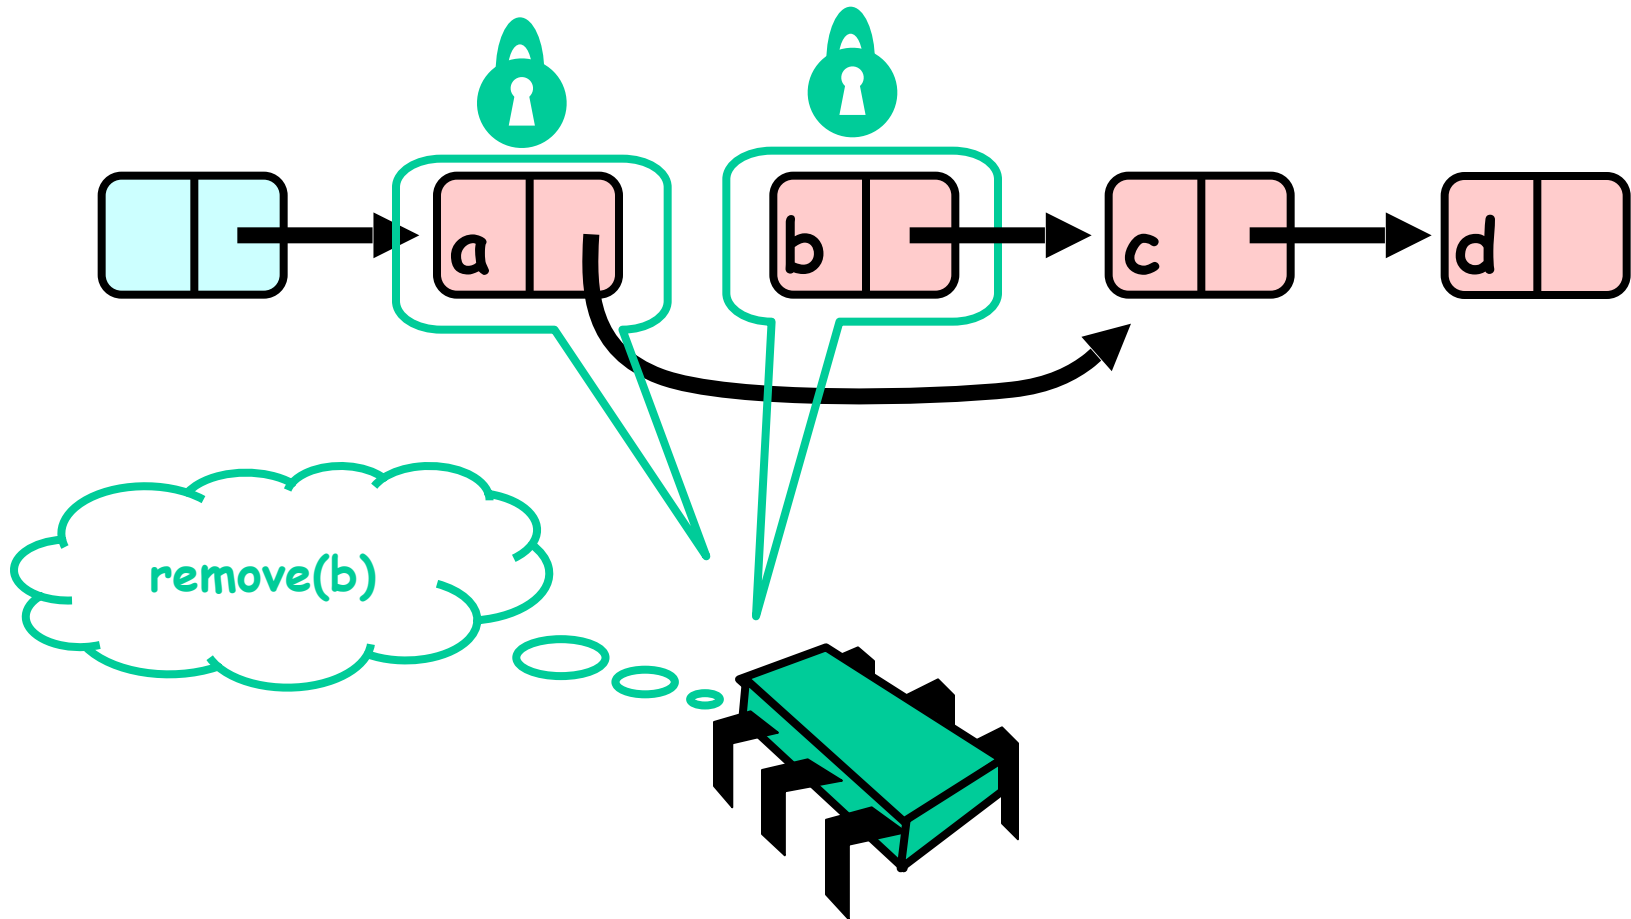
Removing a Node



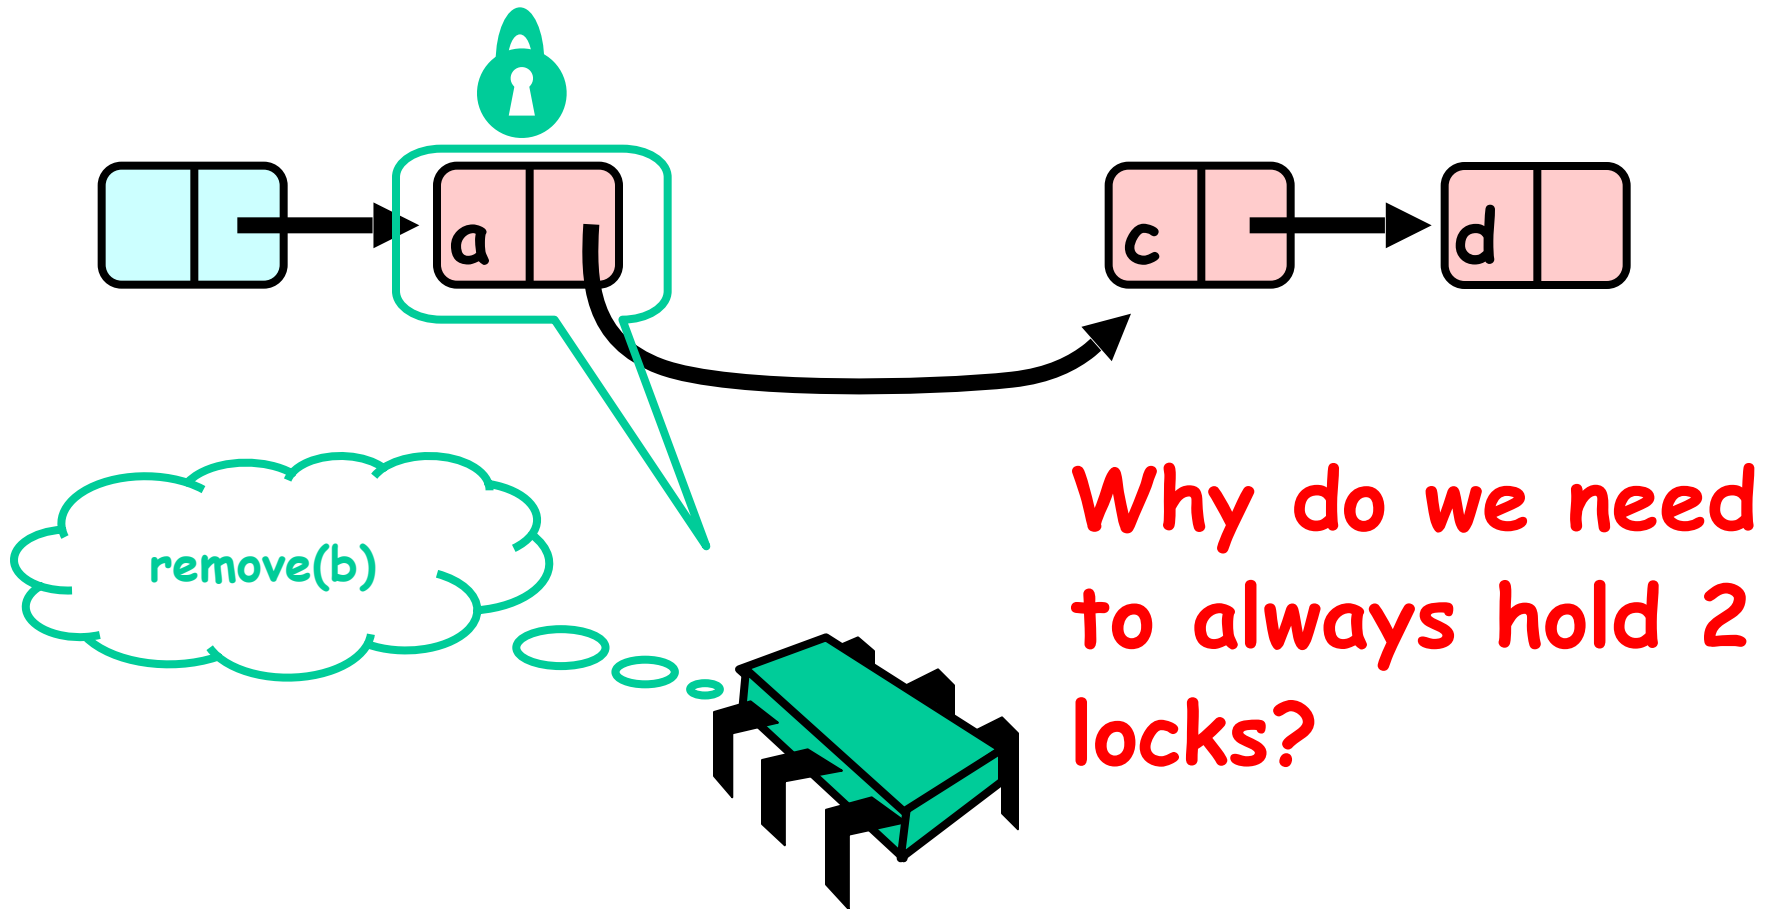
Removing a Node



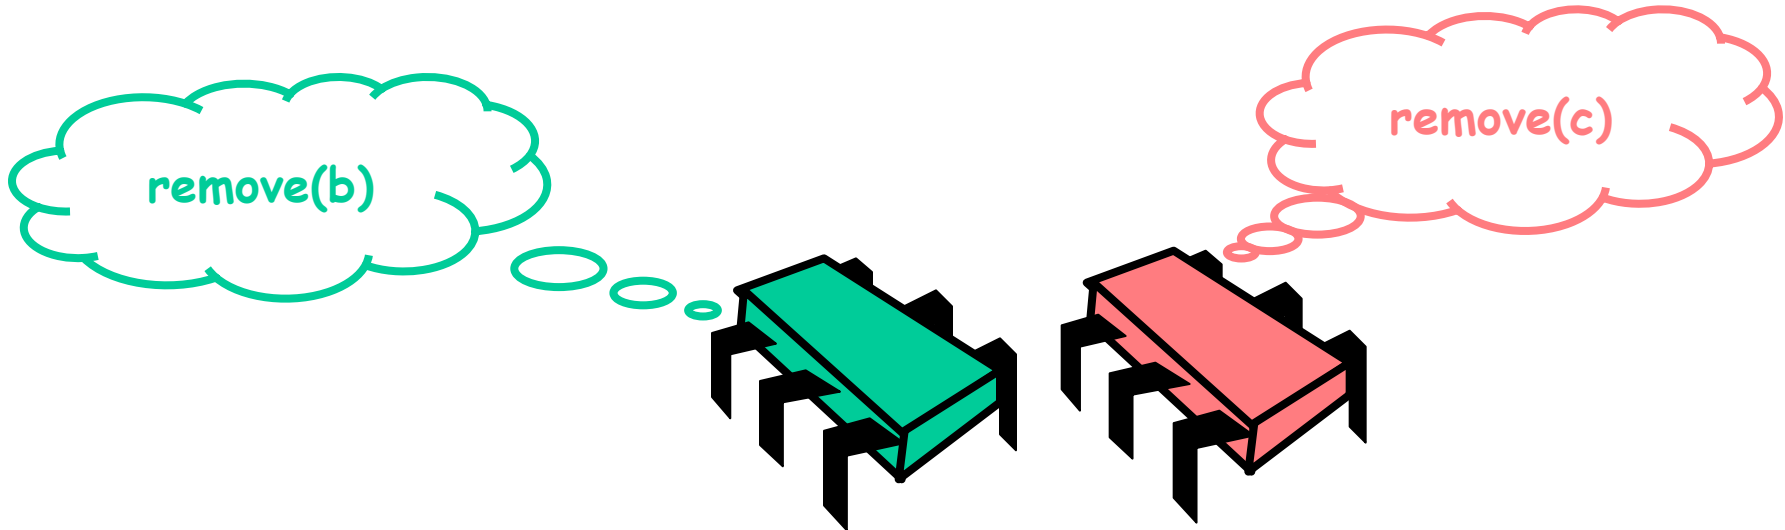
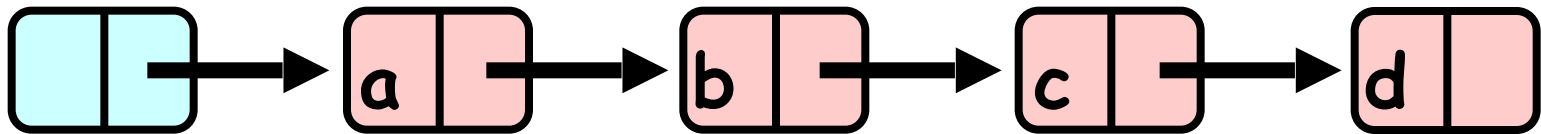
Removing a Node




Removing a Node




Concurrent Removes




Concurrent Removes




Concurrent Removes




Concurrent Removes



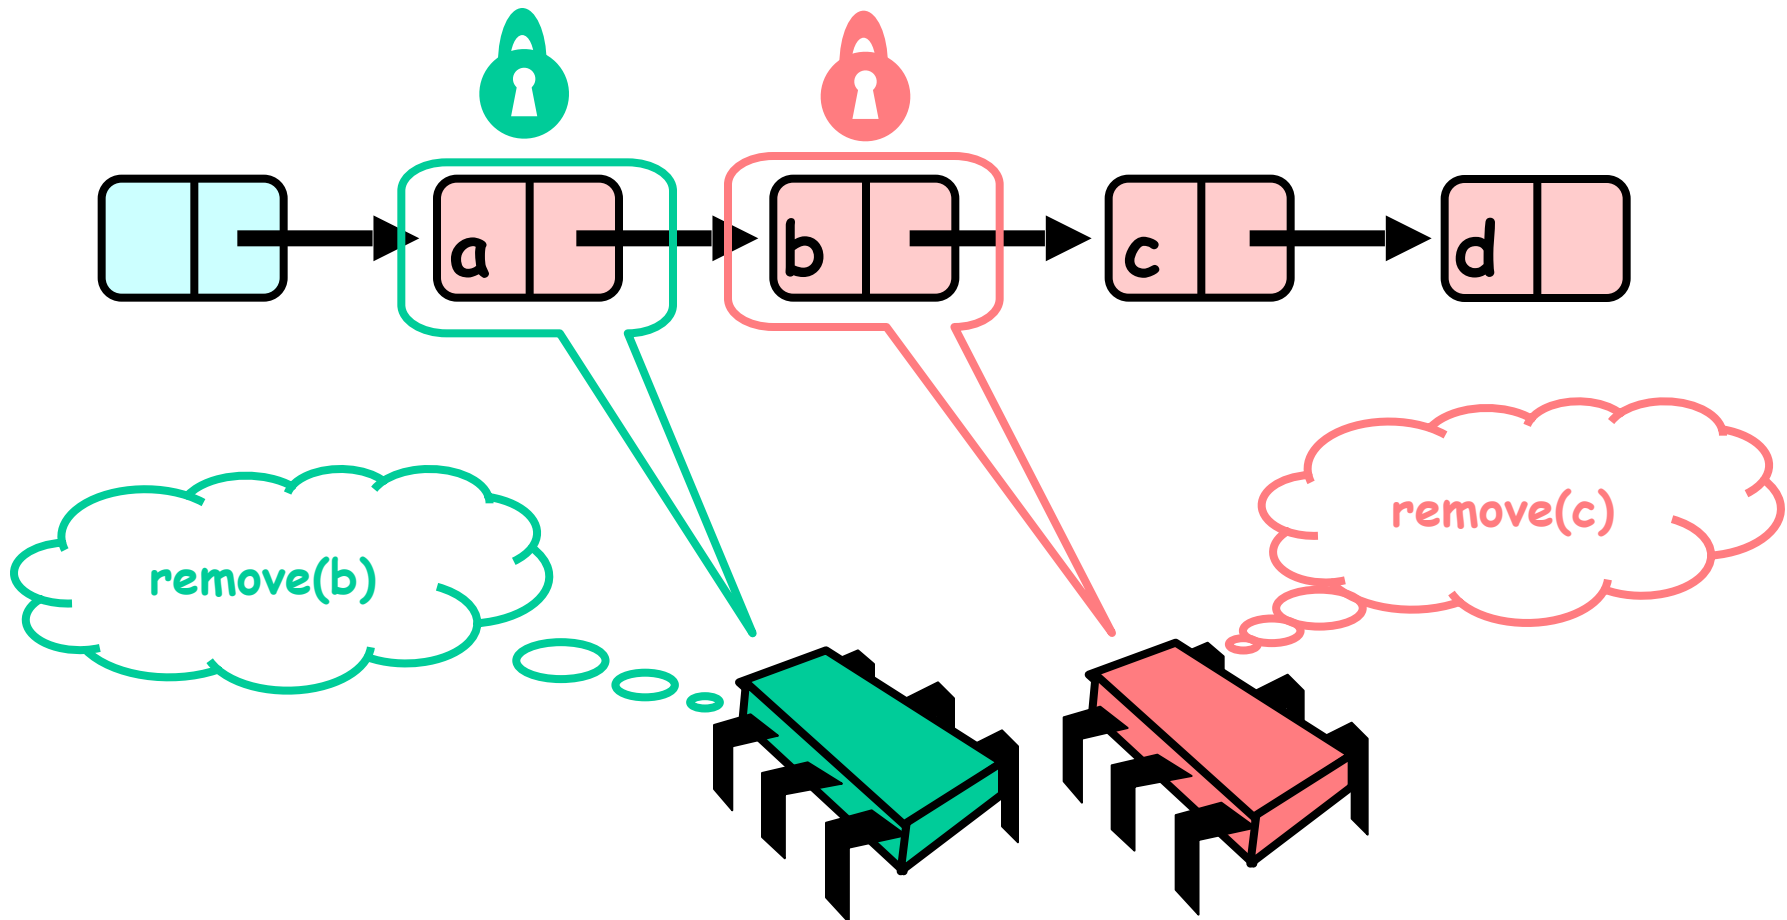
Concurrent Removes




Concurrent Removes



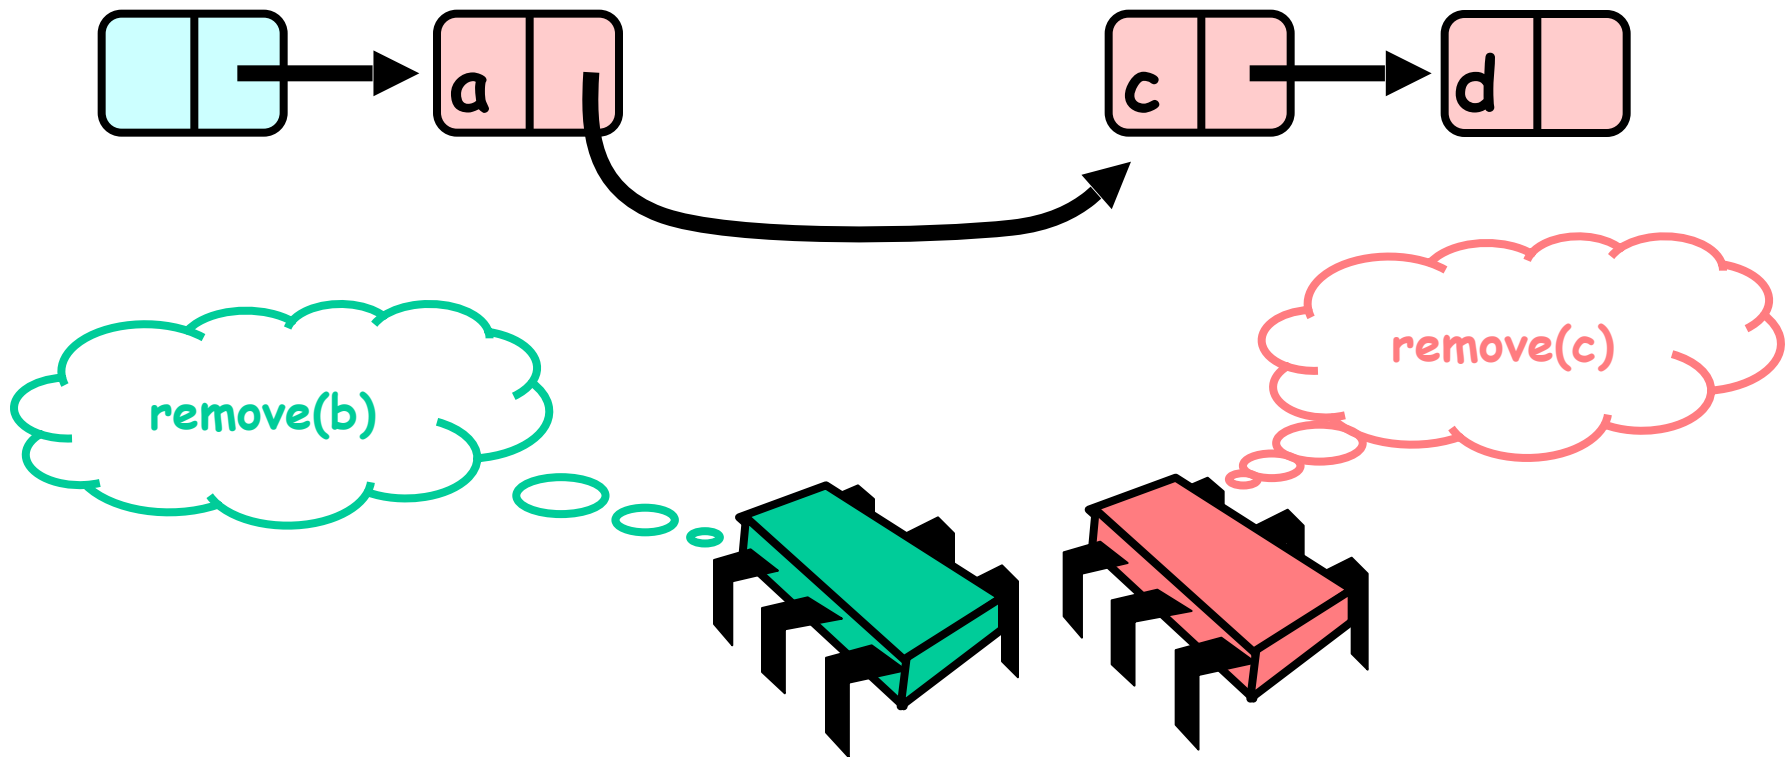
Concurrent Removes



Concurrent Removes

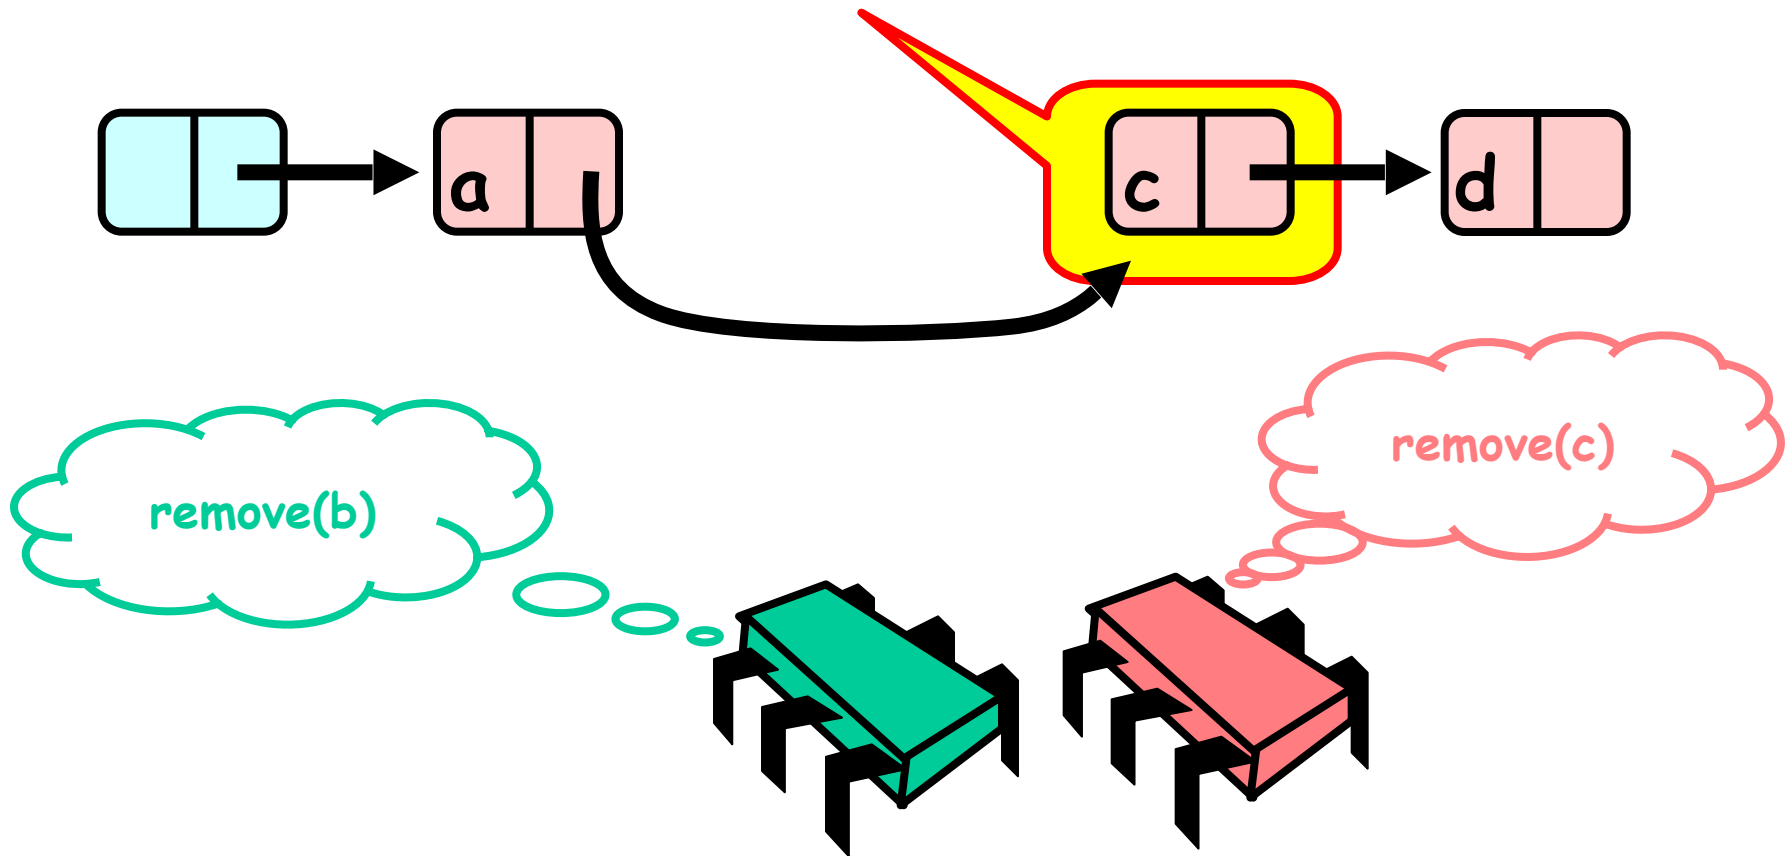


Uh, Oh



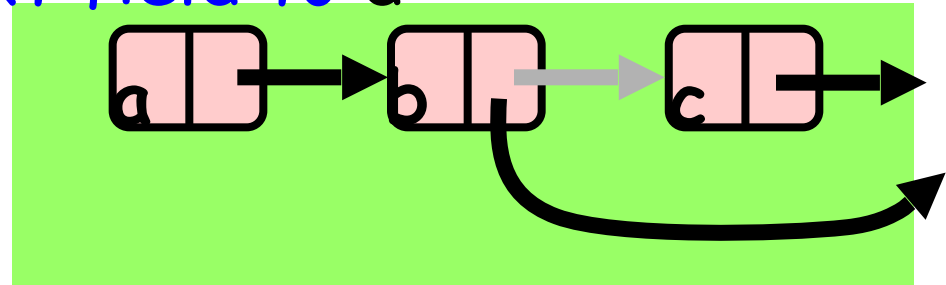
Uh, Oh

Bad news, C not removed

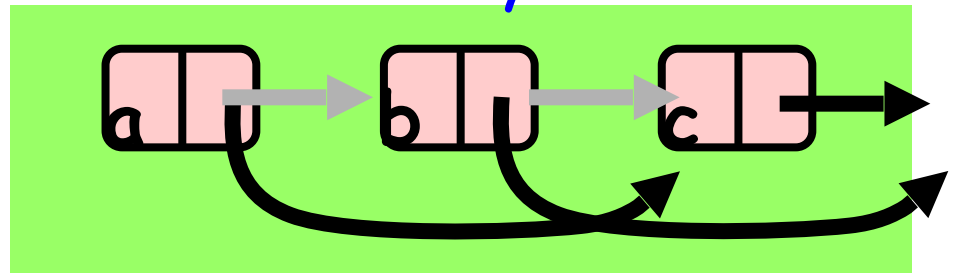


Problem

- To delete node *c*
 - Swing node *b*'s next field to *d*



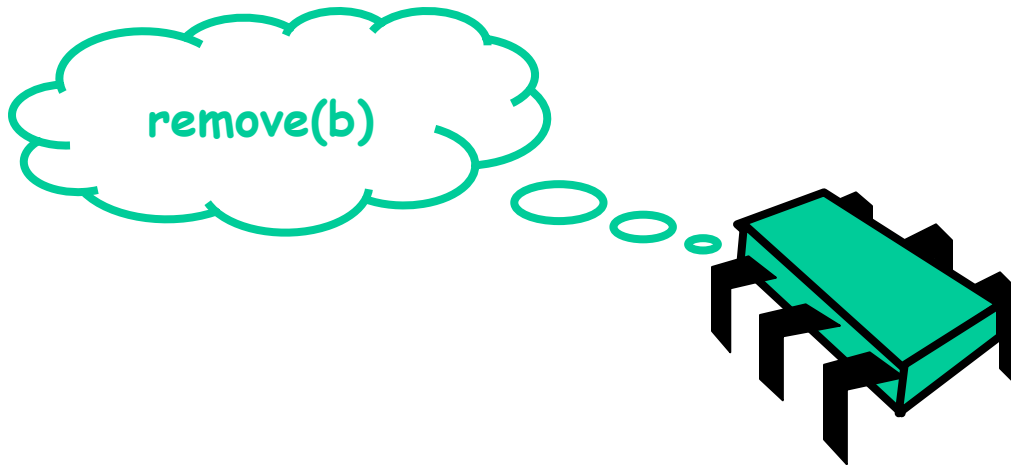
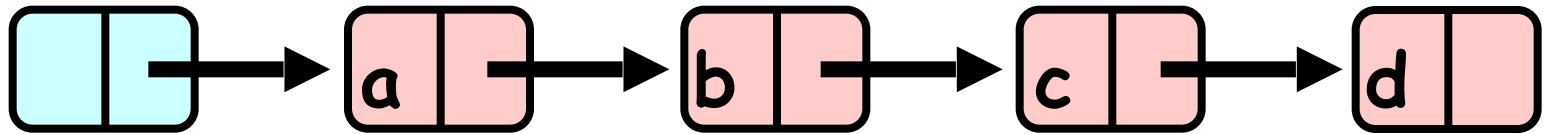
- Problem is,
 - Someone deleting *b* concurrently could direct a pointer to *c*



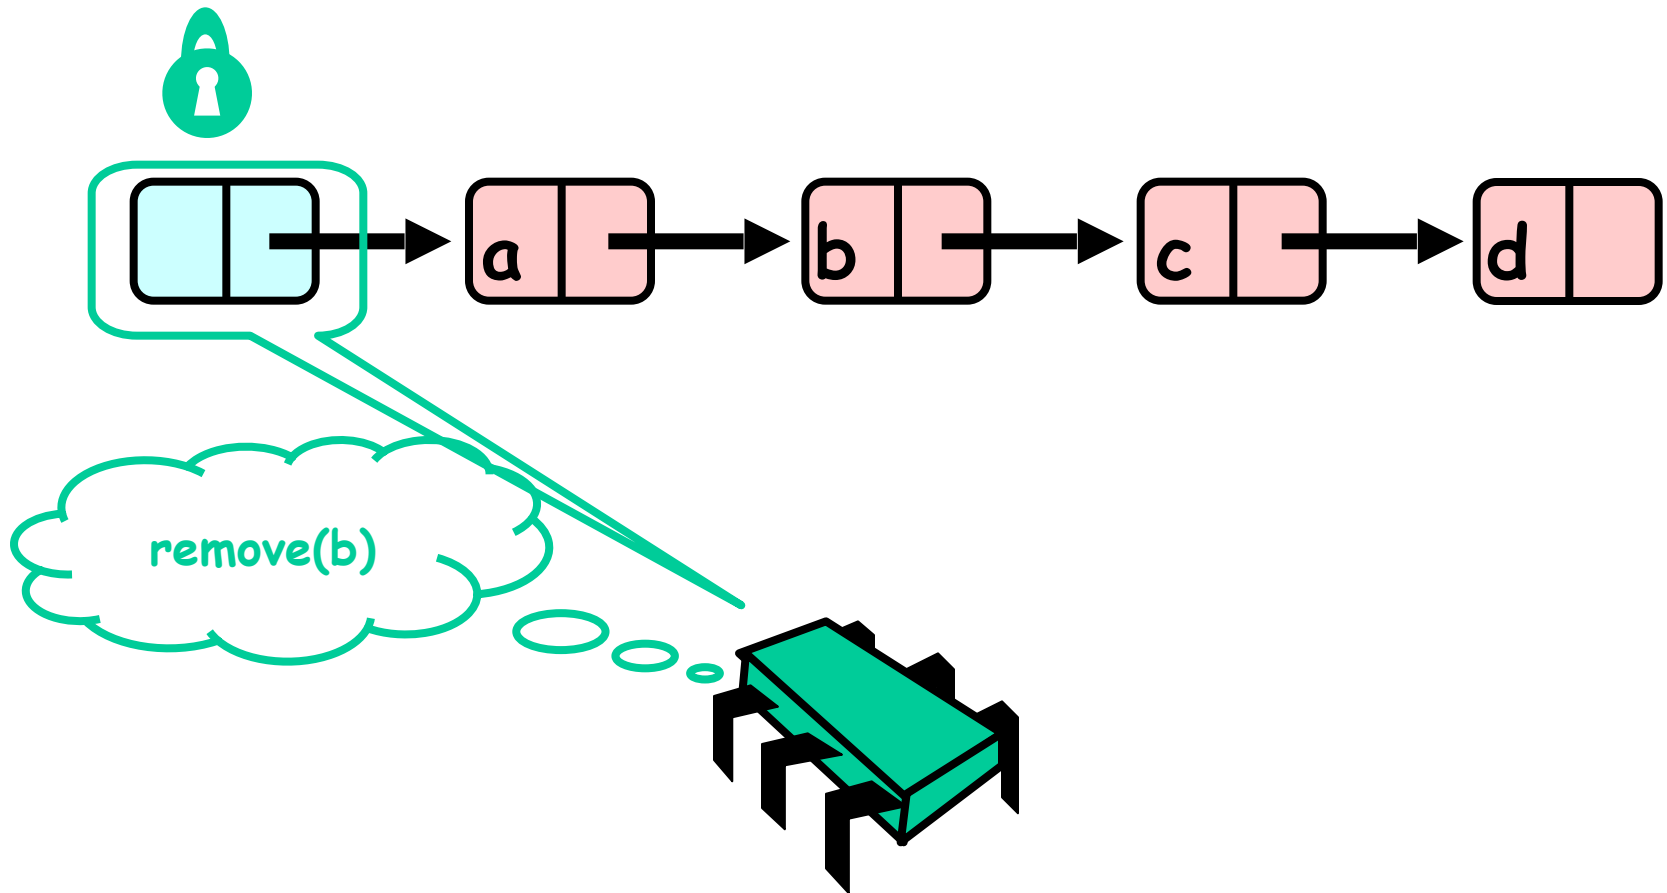
Insight

- If a node is locked
 - No one can delete node's successor
- If a thread locks
 - Node to be deleted
 - And its predecessor
 - Then it works

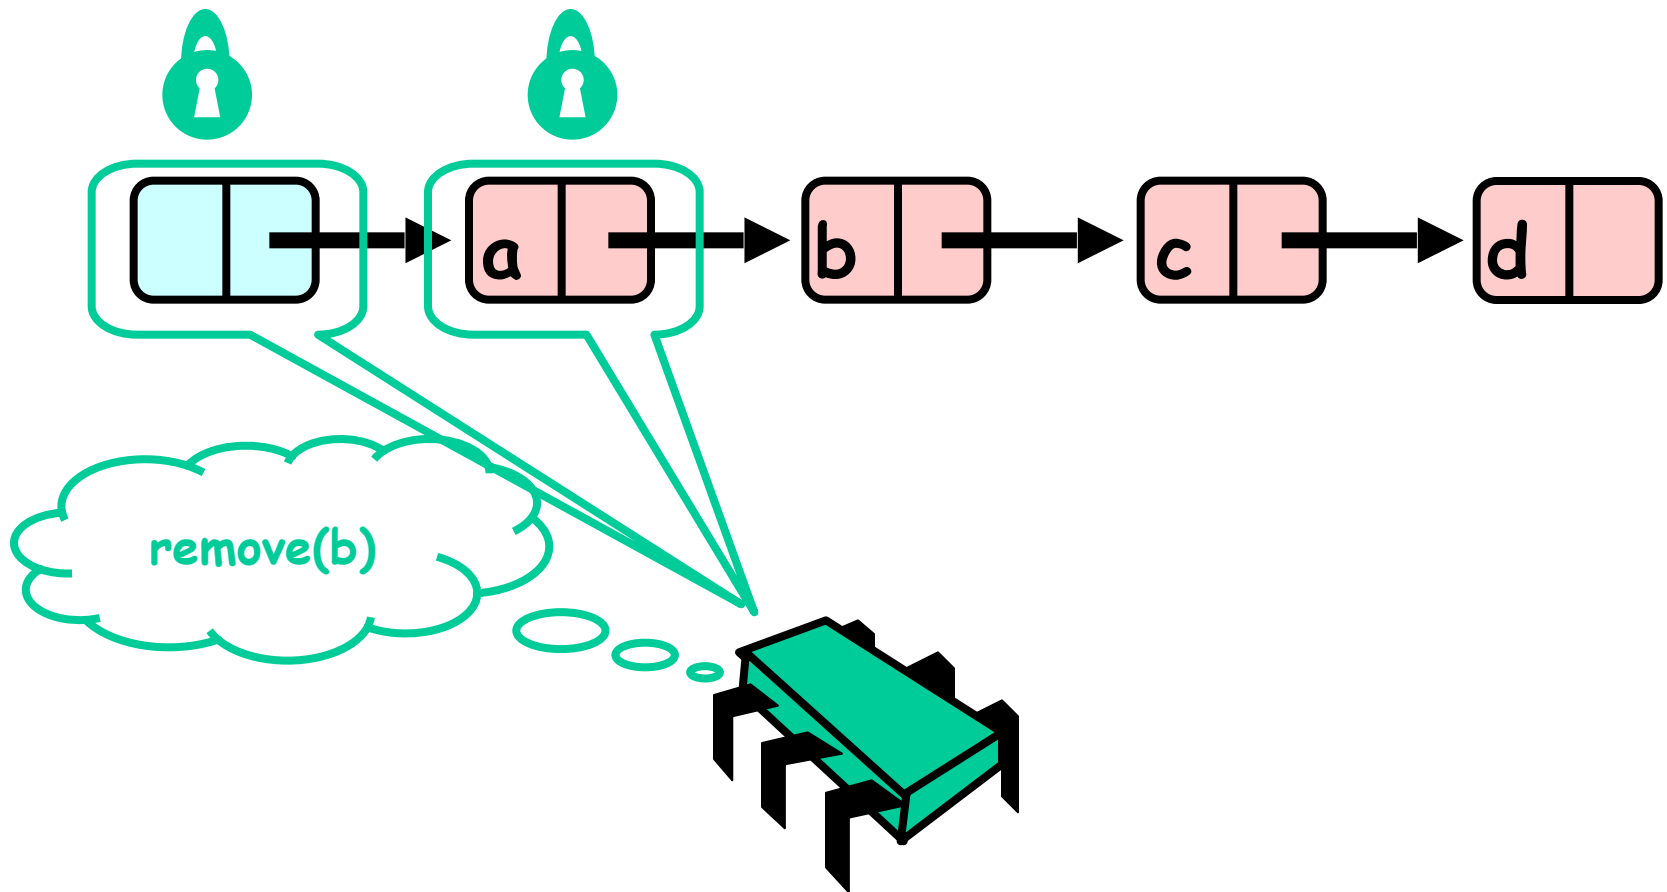
Hand-Over-Hand Again



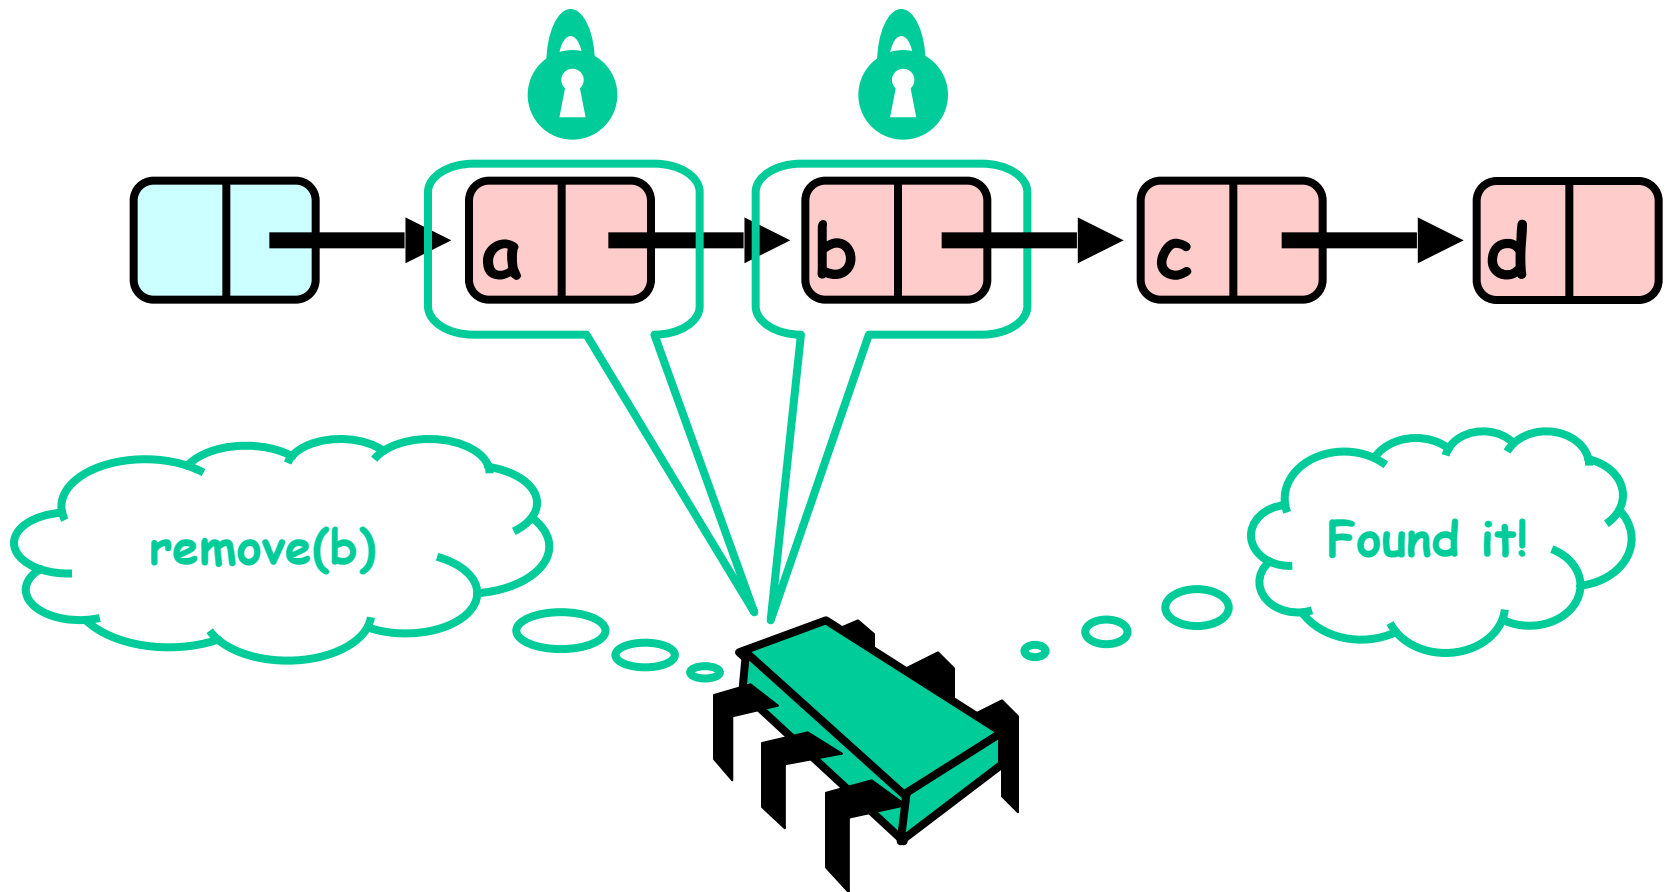
Hand-Over-Hand Again



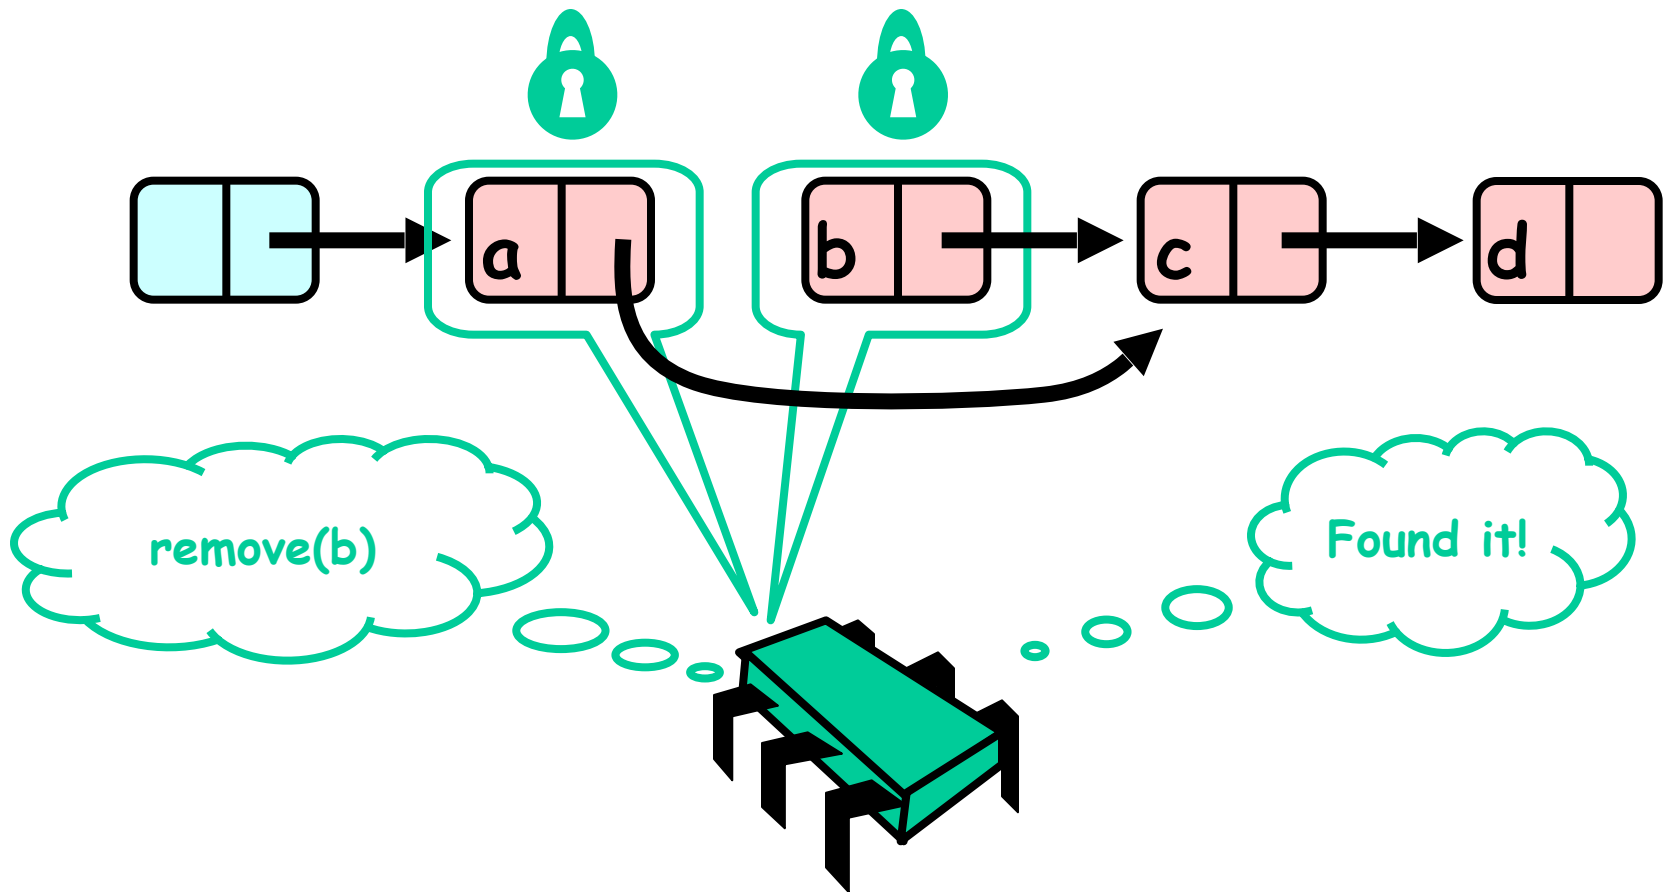
Hand-Over-Hand Again



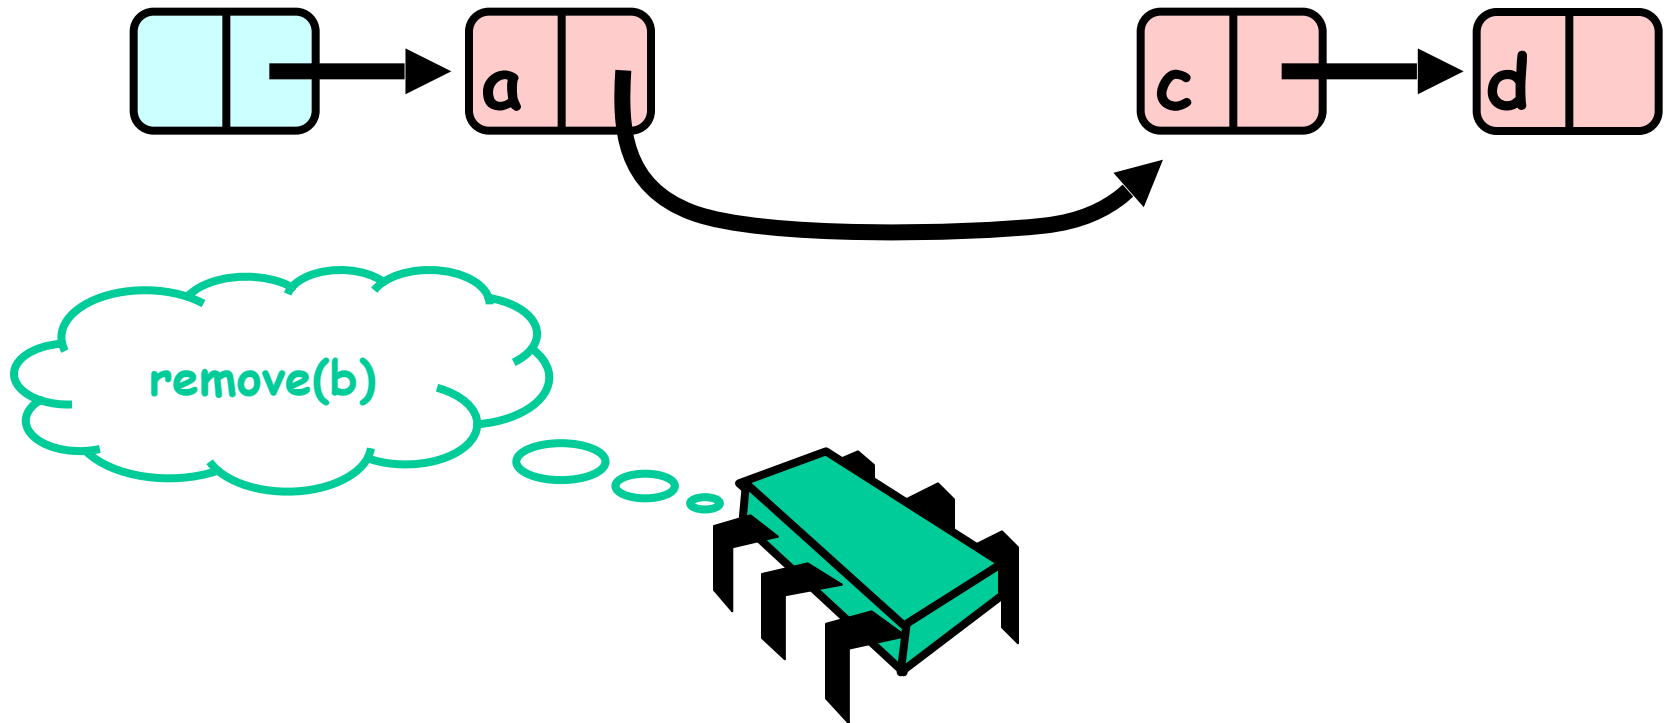
Hand-Over-Hand Again



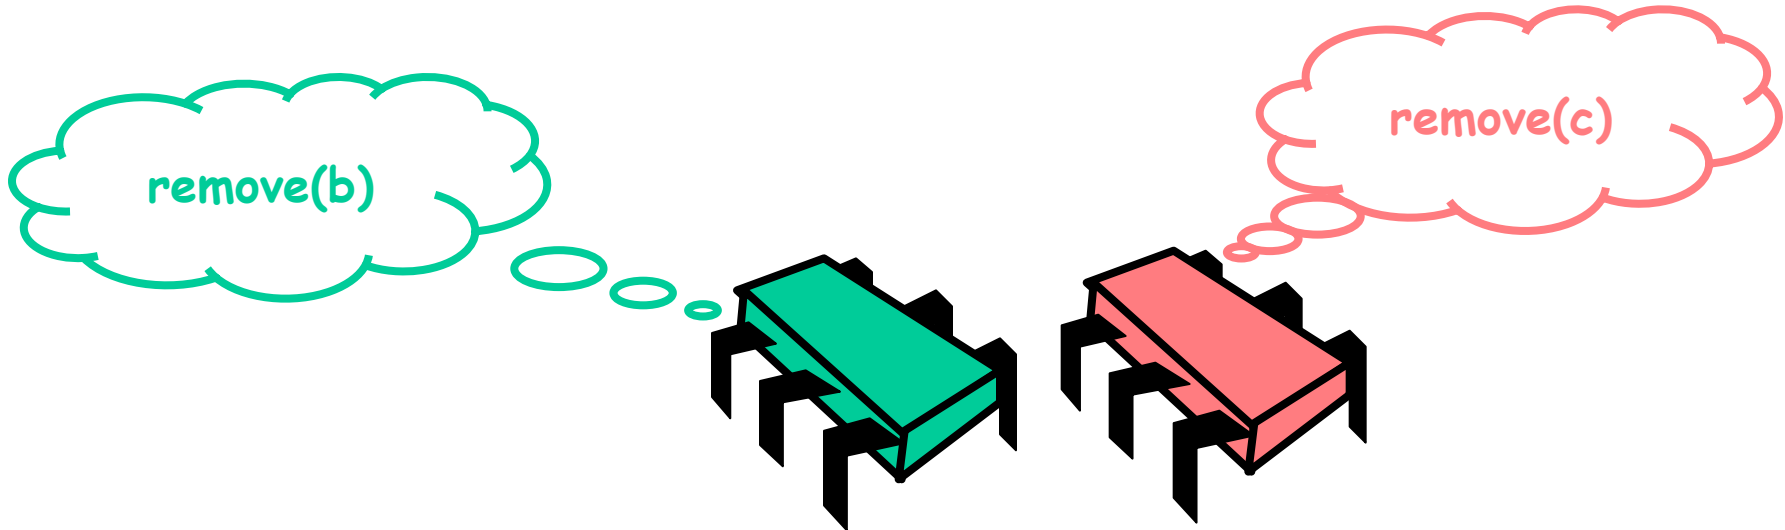
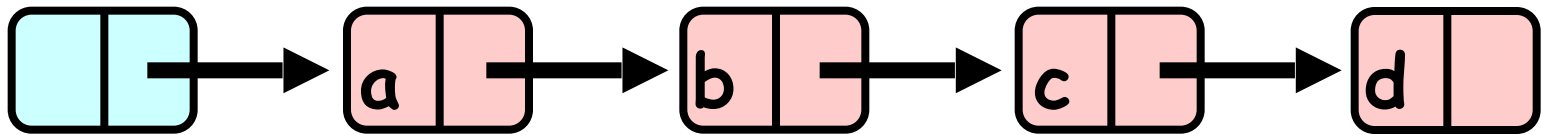
Hand-Over-Hand Again



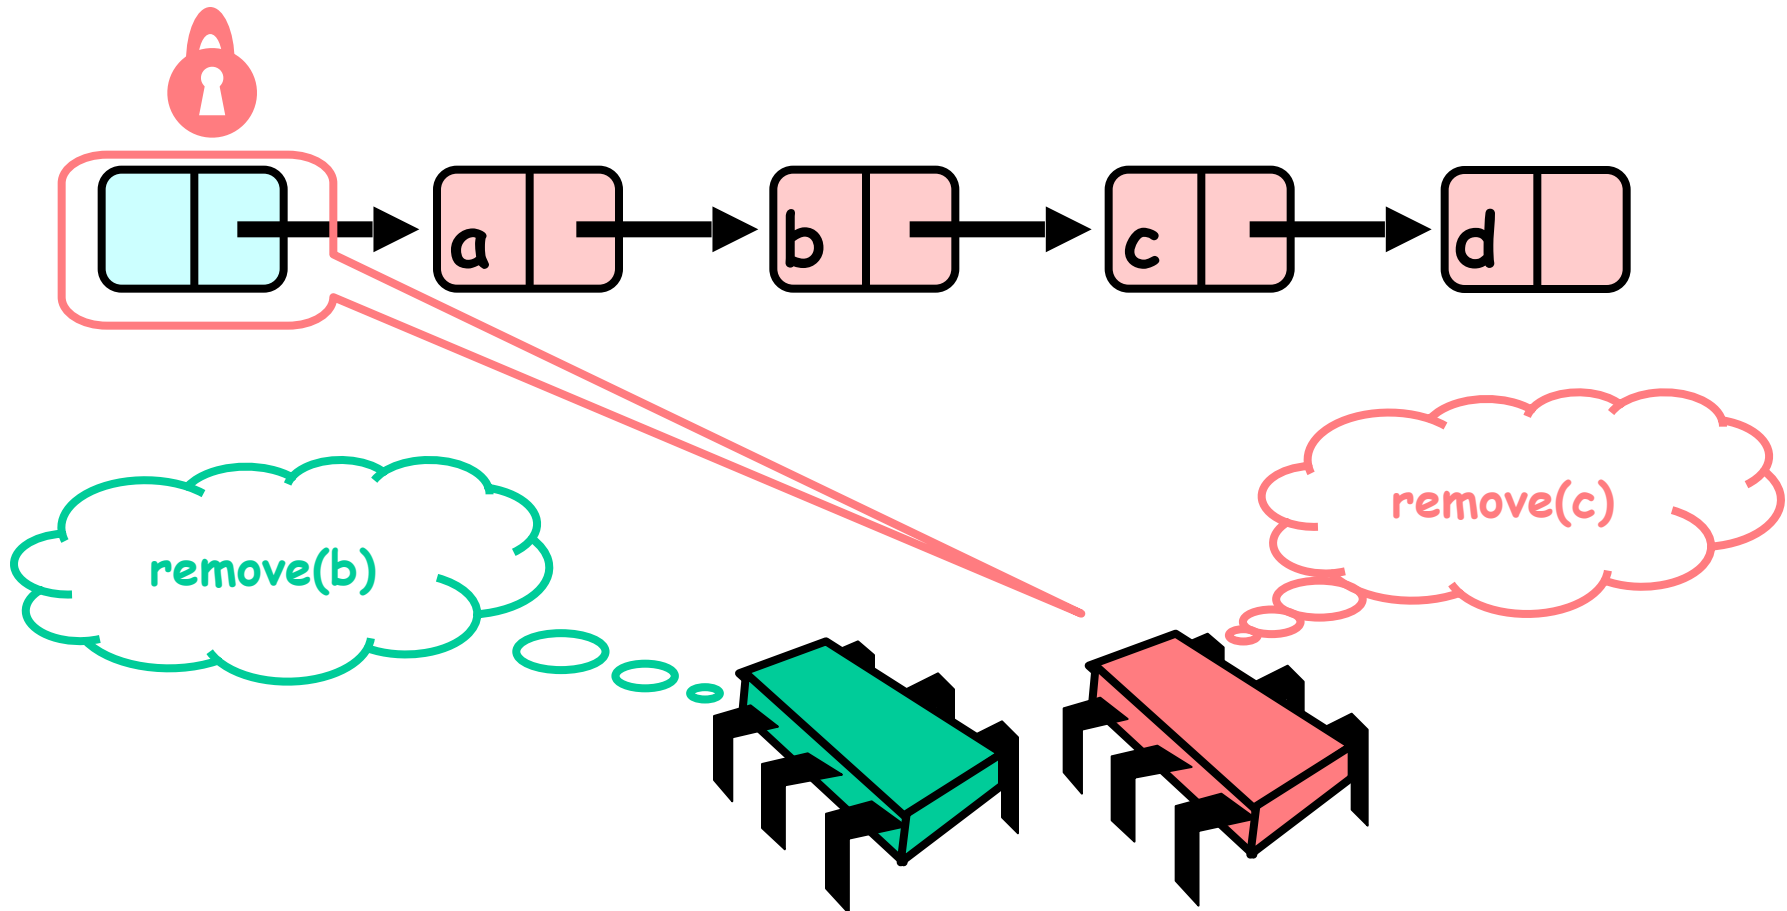
Hand-Over-Hand Again



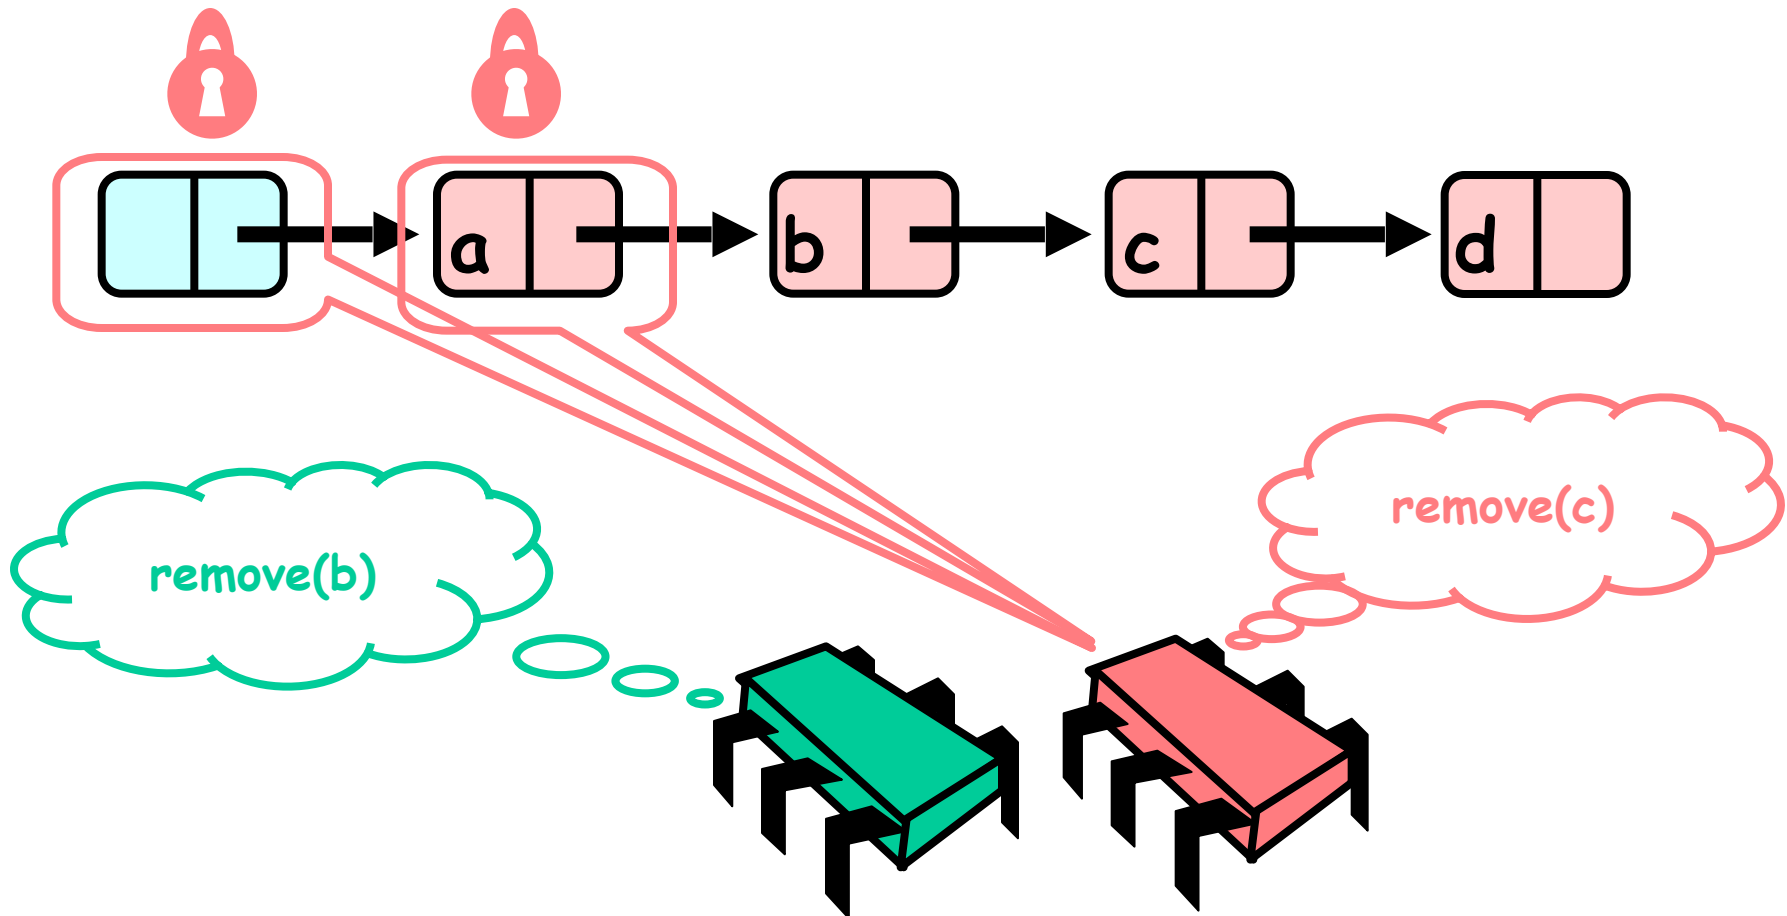
Removing a Node



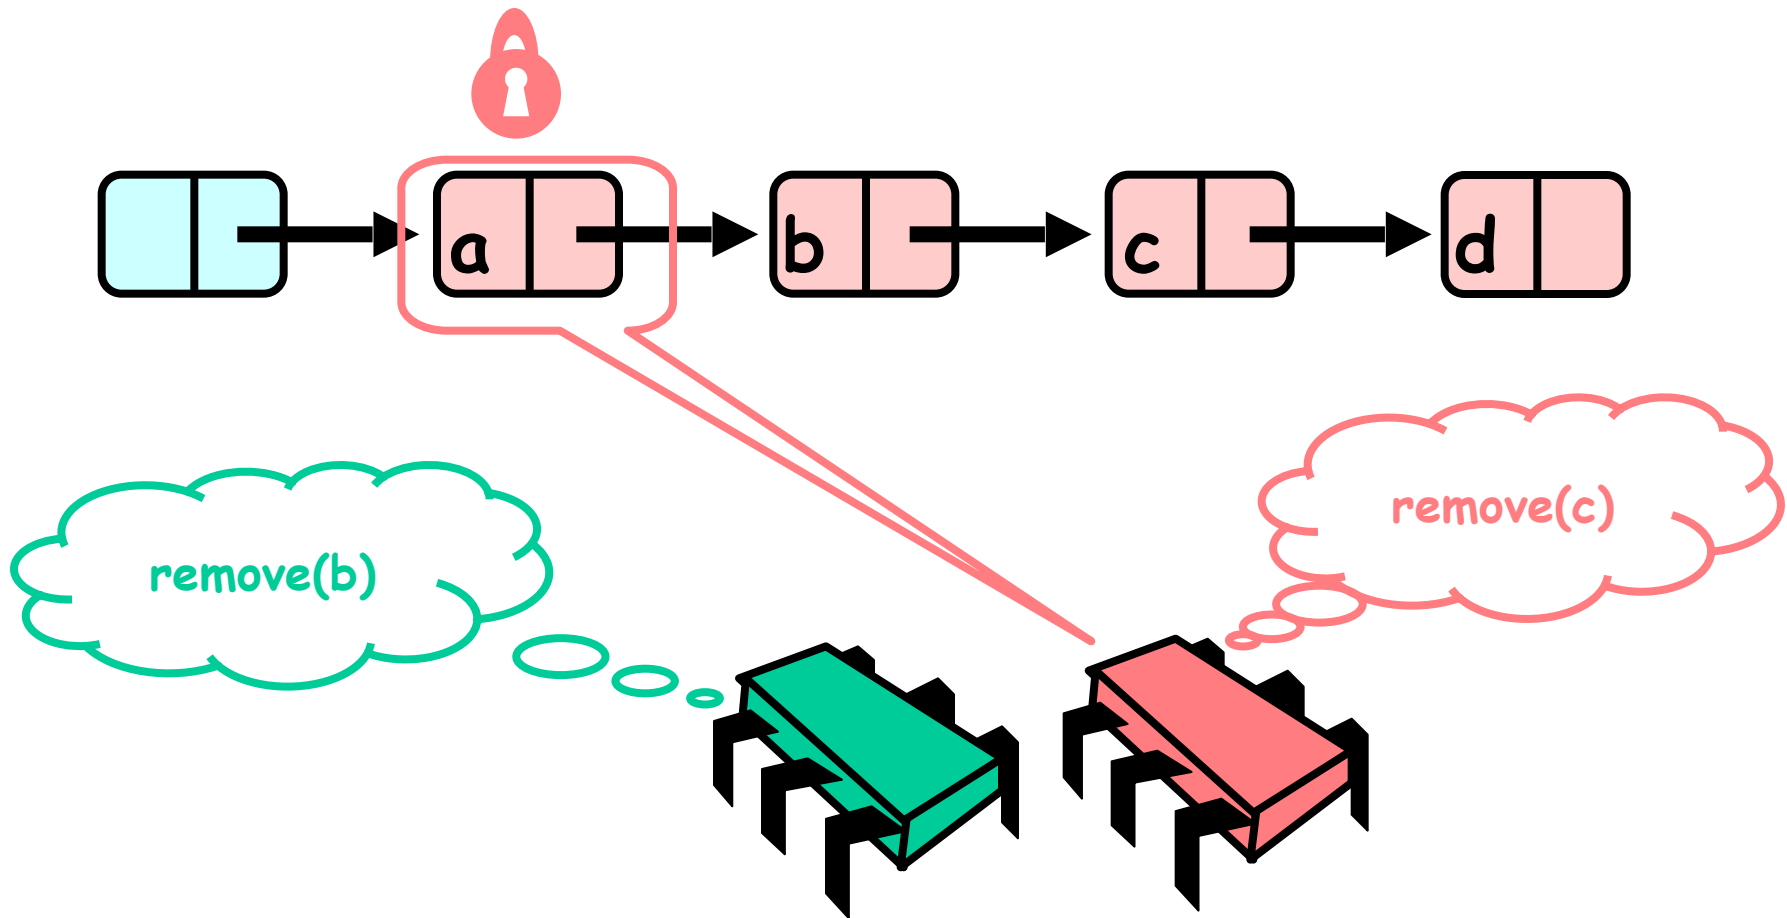
Removing a Node



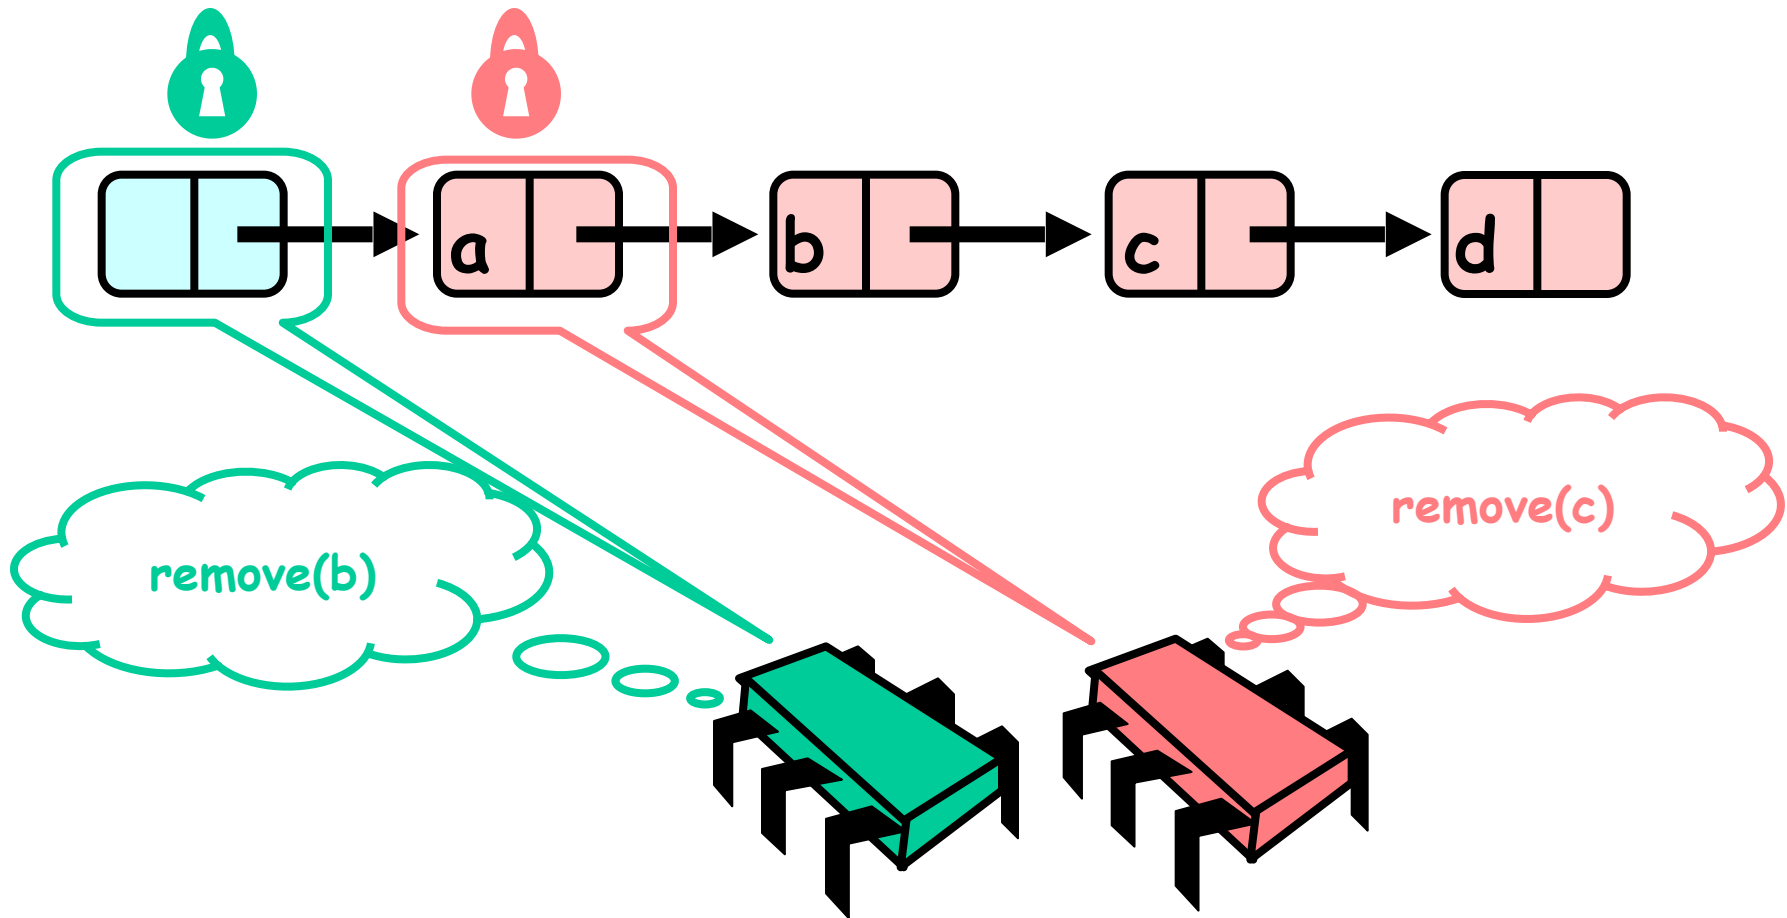
Removing a Node




Removing a Node



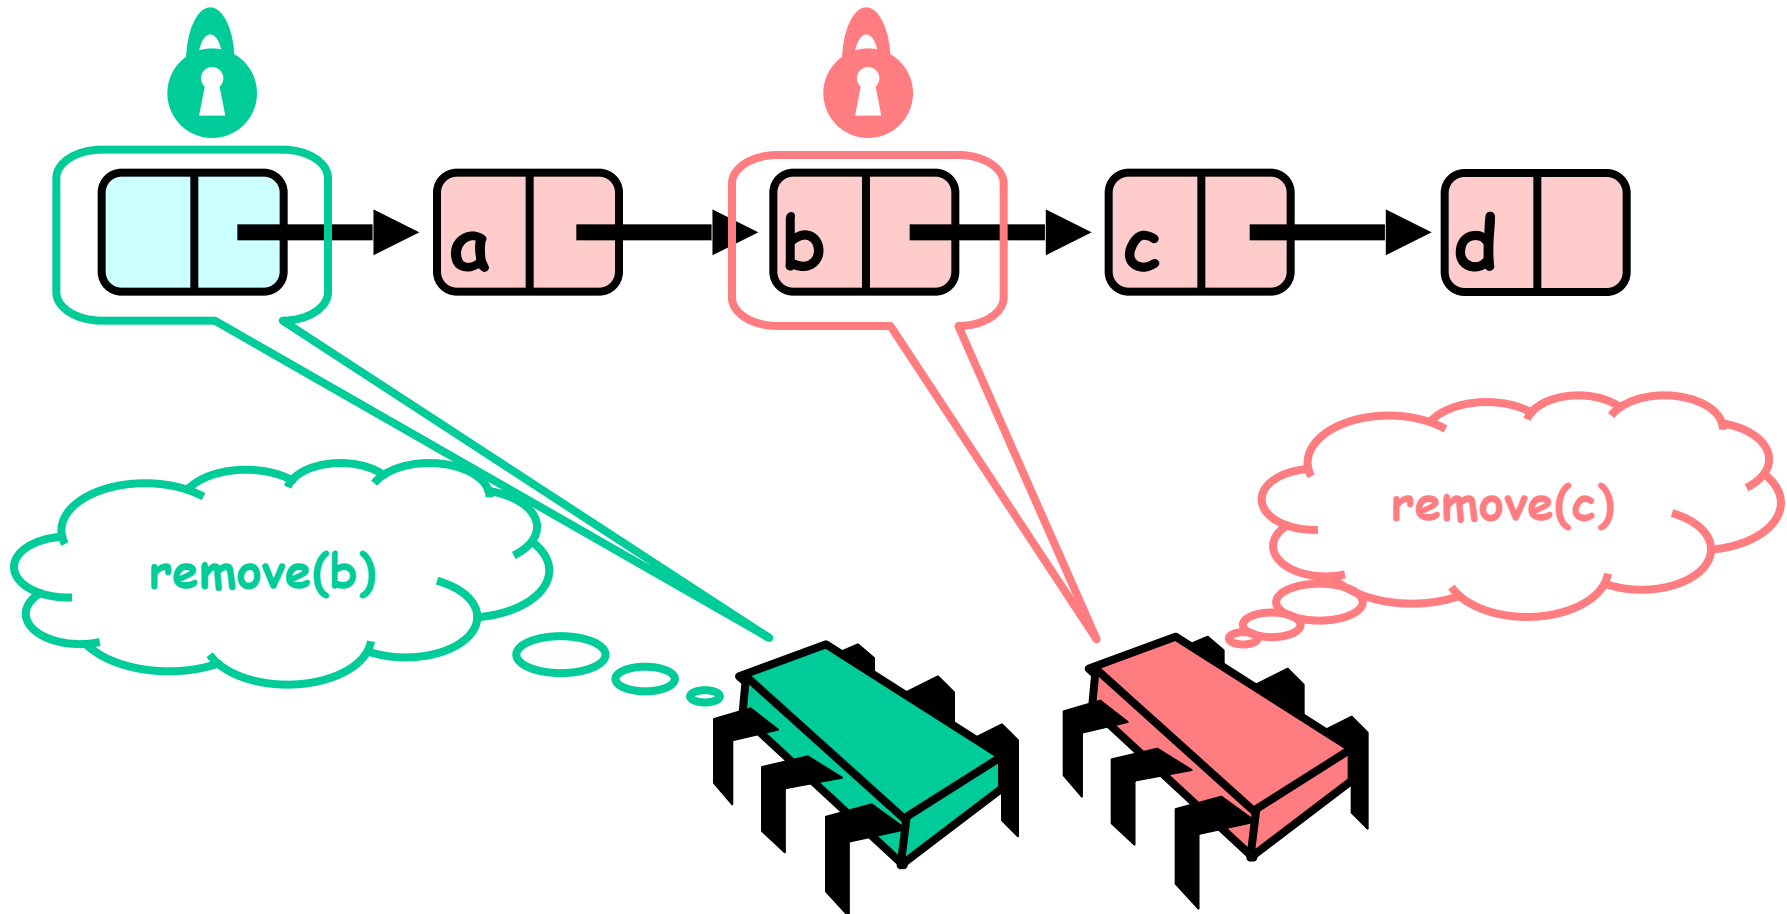
Removing a Node



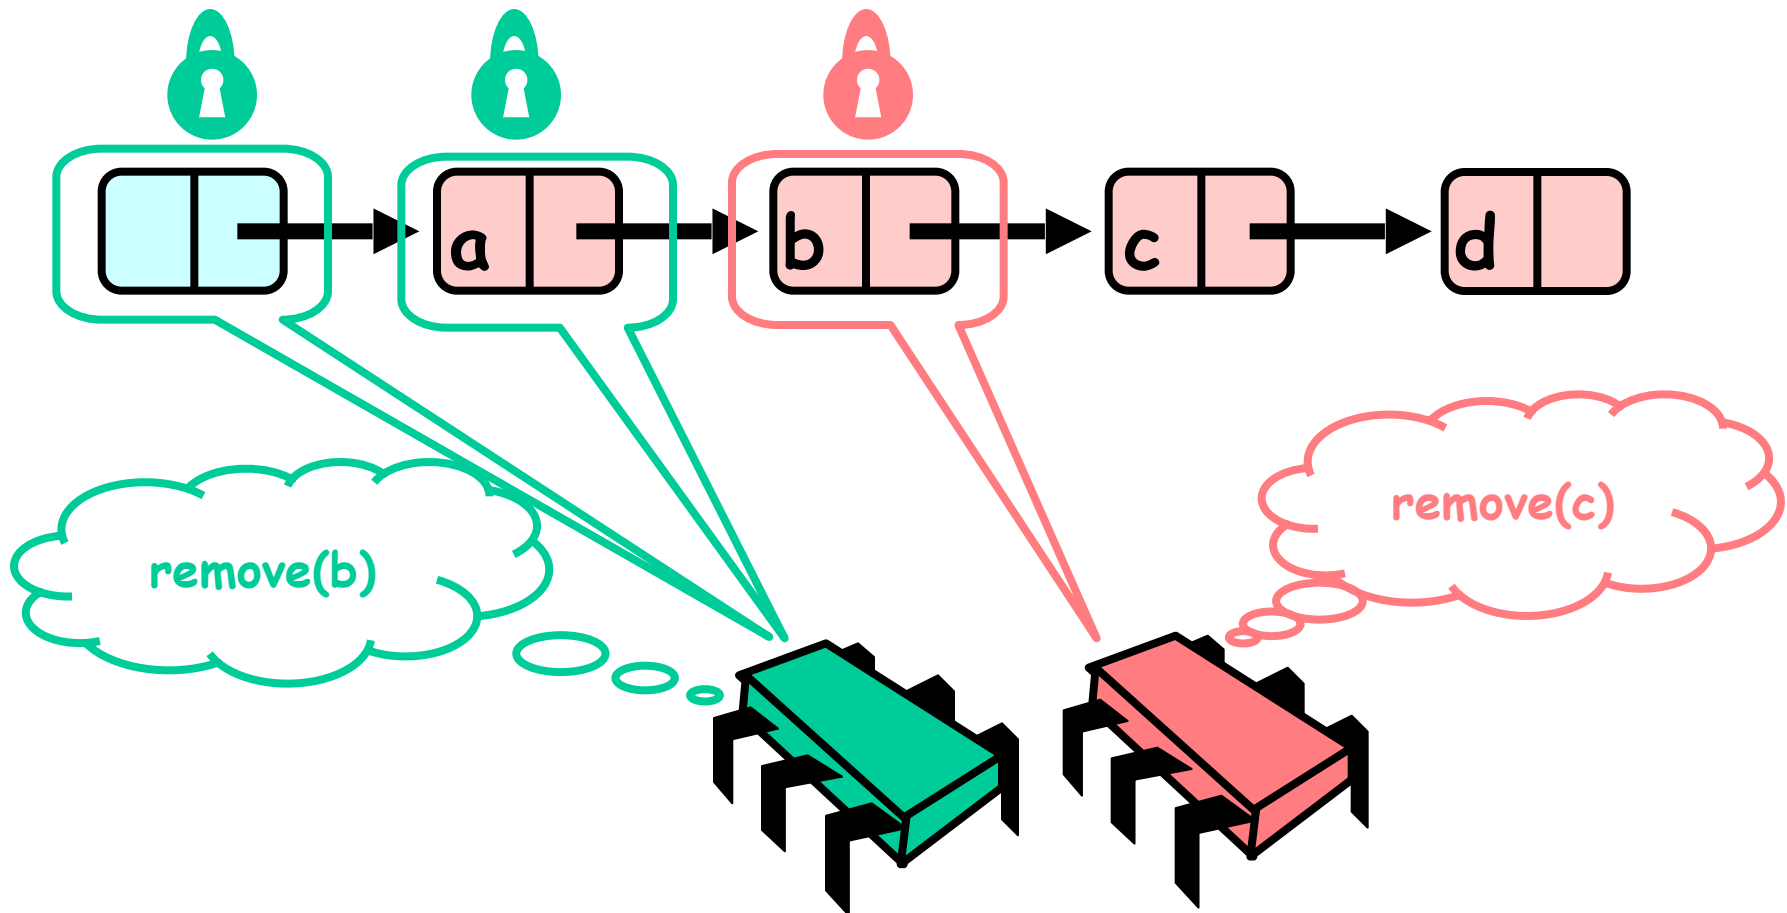
Removing a Node



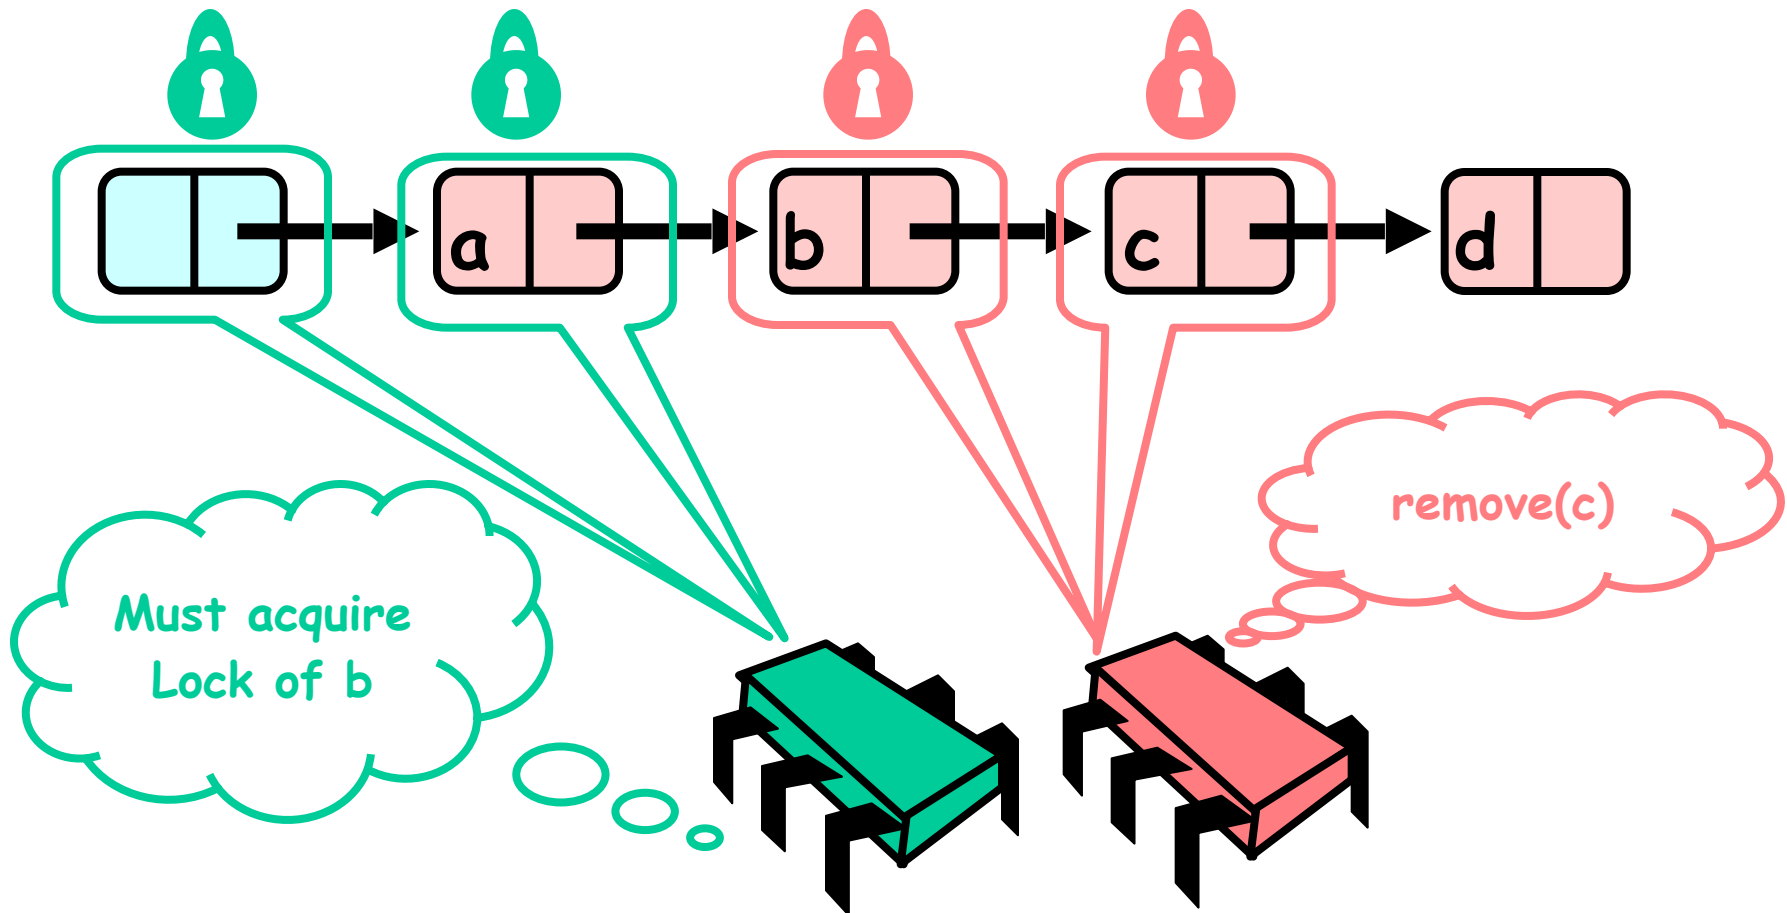
Removing a Node




Removing a Node



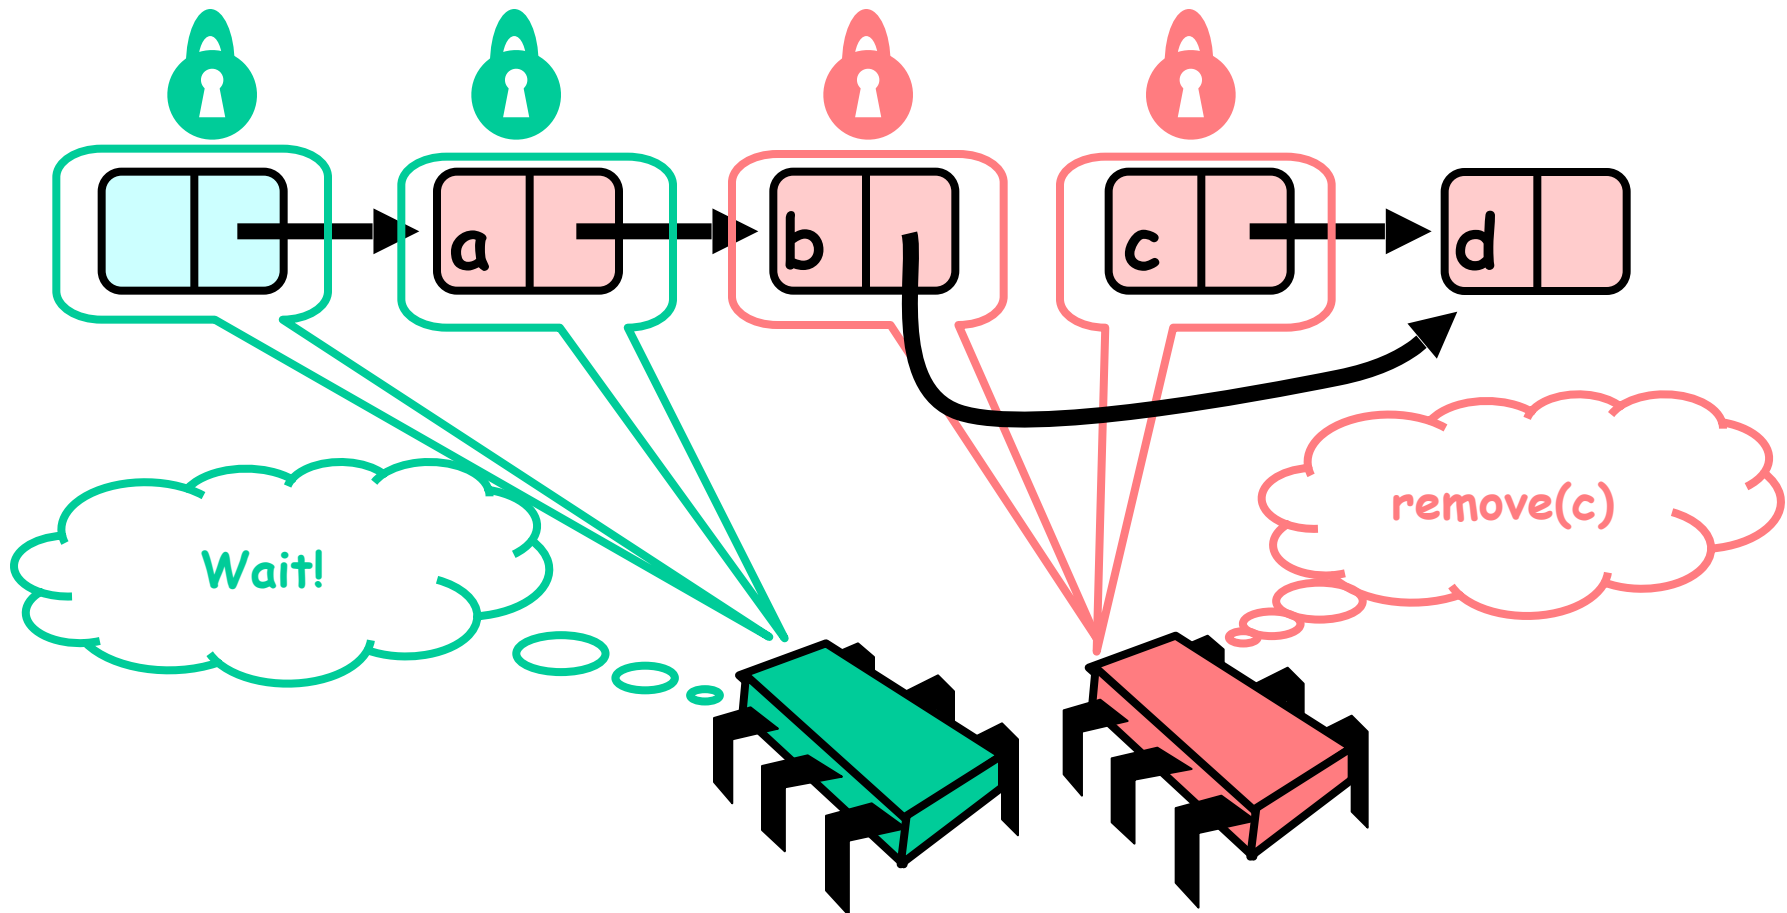
Removing a Node




Removing a Node



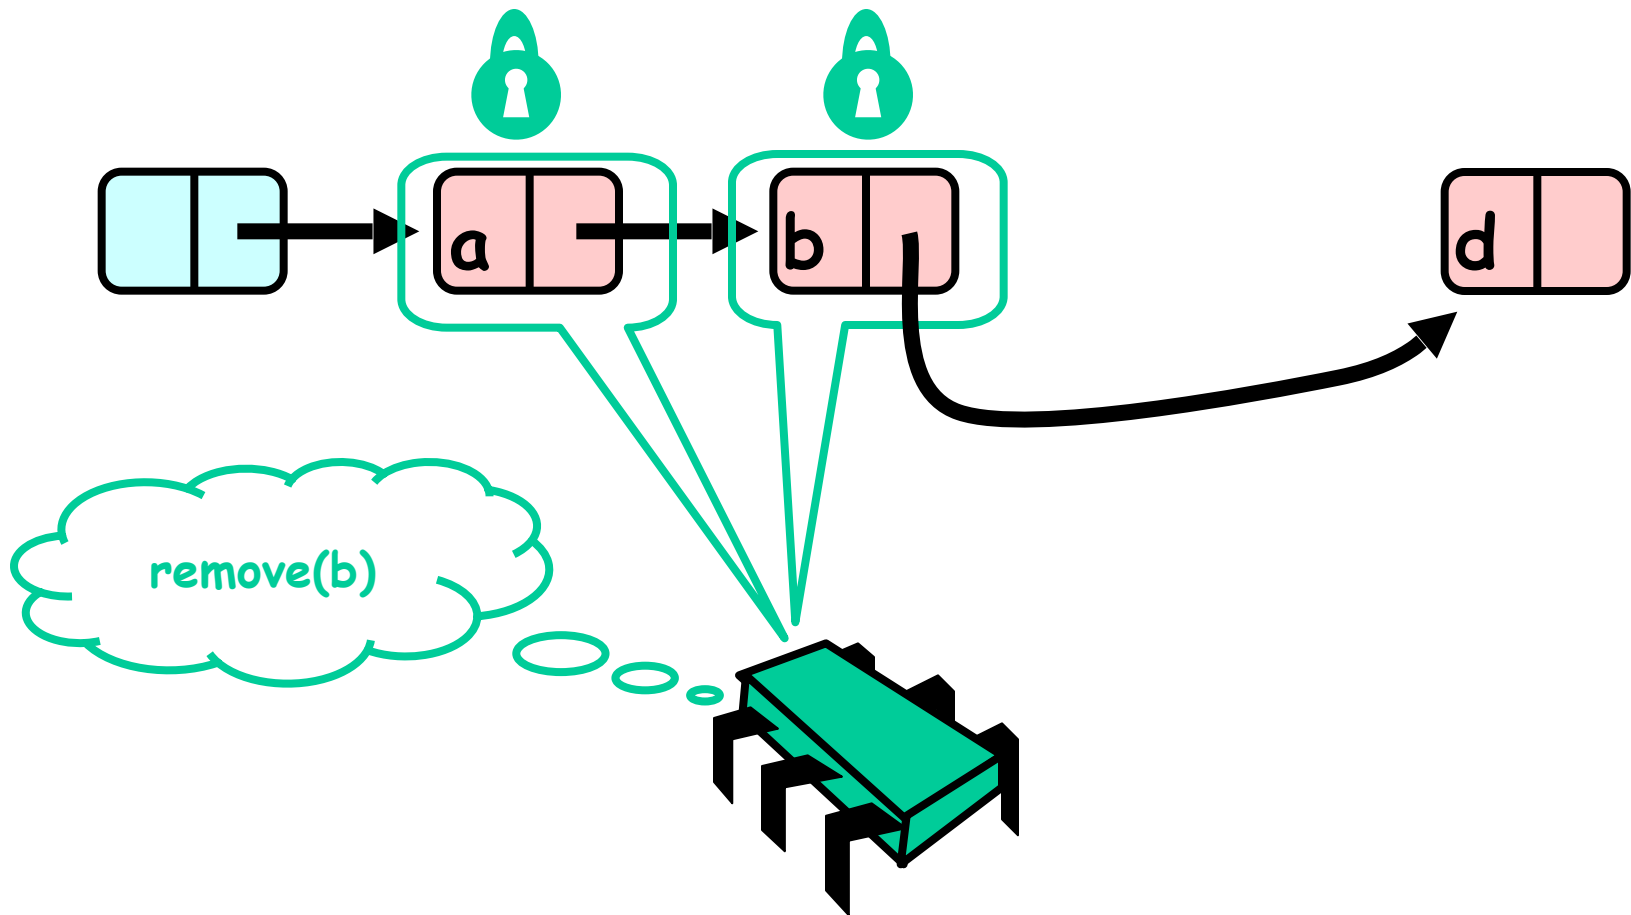
Removing a Node




Removing a Node



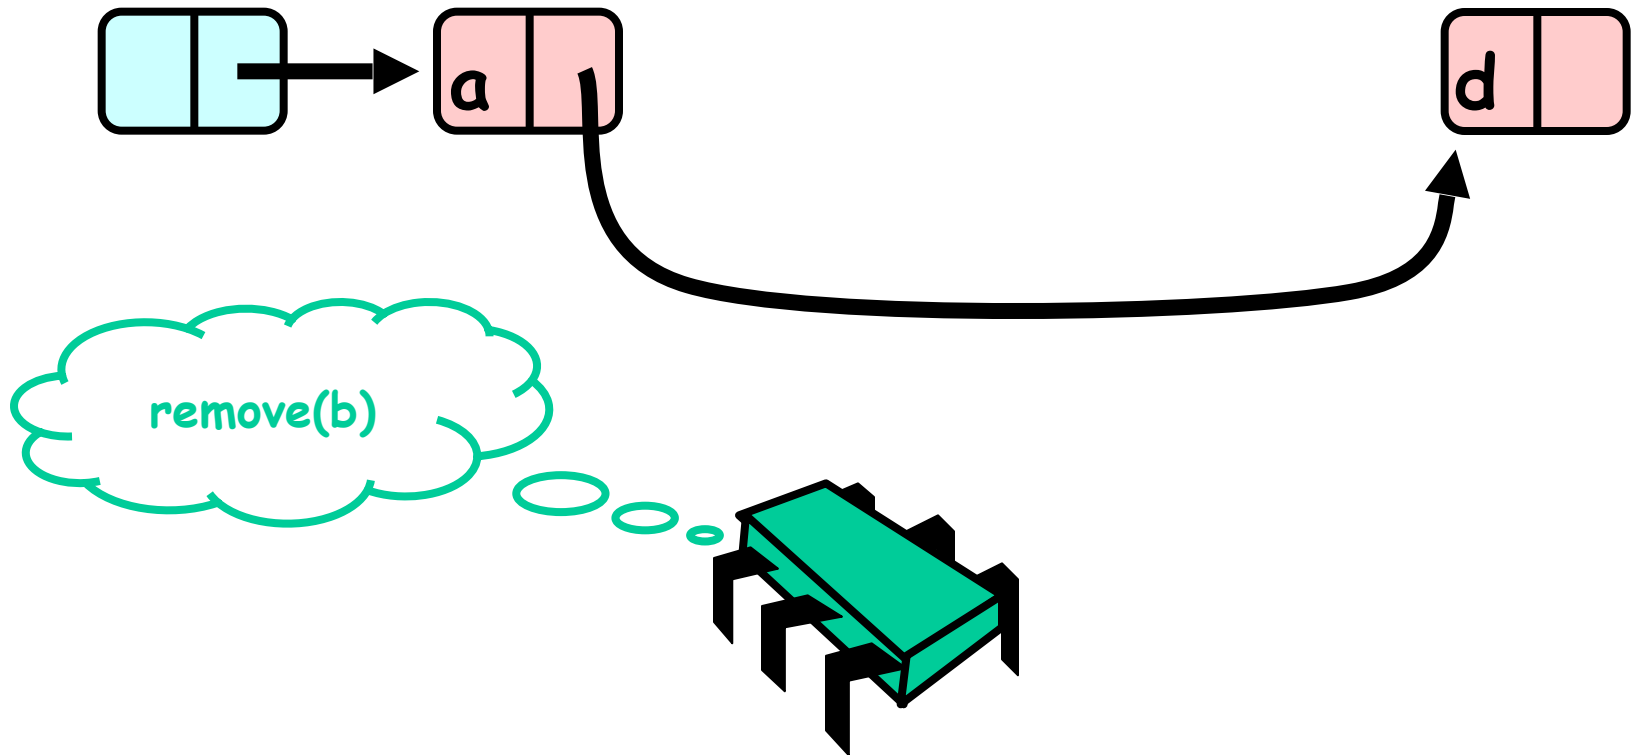
Removing a Node



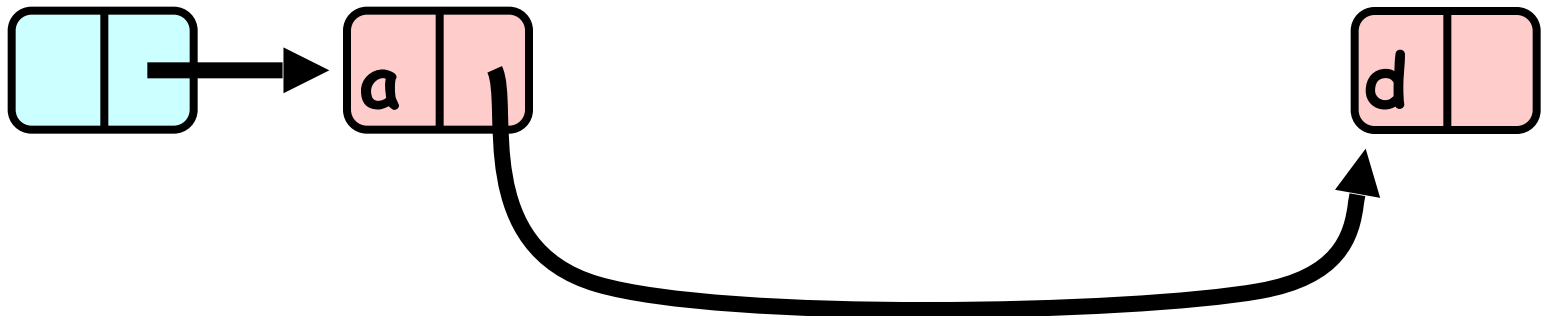
Removing a Node



Removing a Node



Removing a Node



Fine-Grained Synchronization: hand-over-hand locking Linked List

```
public boolean add(T item) {
    int key = item.hashCode();
    head.lock();
    Node pred = head;
    try {
        Node curr = pred.next;
        curr.lock();
        try {
            while (curr.key < key) {
                pred.unlock();
                pred = curr;
                curr = curr.next;
                curr.lock();
            }
            if (curr.key == key) return false;
            Node newNode = new Node(item);
            newNode.next = curr;
            pred.next = newNode;
            return true;
        } finally {
            curr.unlock();
        }
    } finally {
        pred.unlock();
    }
}
```

```

public boolean remove(T item) {
    Node pred = null, curr = null;
    int key = item.hashCode();
    head.lock();
    try {
        pred = head;
        curr = pred.next;
        curr.lock();
        try {
            while (curr.key < key) {
                pred.unlock();
                pred = curr;
                curr = curr.next;
                curr.lock();
            }
            if (curr.key == key) {
                pred.next = curr.next;
                return true;
            }
            return false;
        } finally {
            curr.unlock();
        }
    } finally {
        pred.unlock();
    }
}

```

```

public boolean contains(T item) {
    Node last = null, pred = null, curr
= null;
    int key = item.hashCode();
    head.lock();
    try {
        pred = head;
        curr = pred.next;
        curr.lock();
        try {
            while (curr.key < key) {
                pred.unlock();
                pred = curr;
                curr = curr.next;
                curr.lock();
            }
            return (curr.key == key);
        } finally {
            curr.unlock();
        }
    } finally {
        pred.unlock();
    }
}
}

```

Adding Nodes

- To add node e
 - Must lock predecessor
 - Must lock successor
- Neither can be deleted


Drawbacks

- Better than coarse-grained lock
 - Threads can traverse in parallel
- Still not ideal
 - Long chain of acquire/release
 - Inefficient

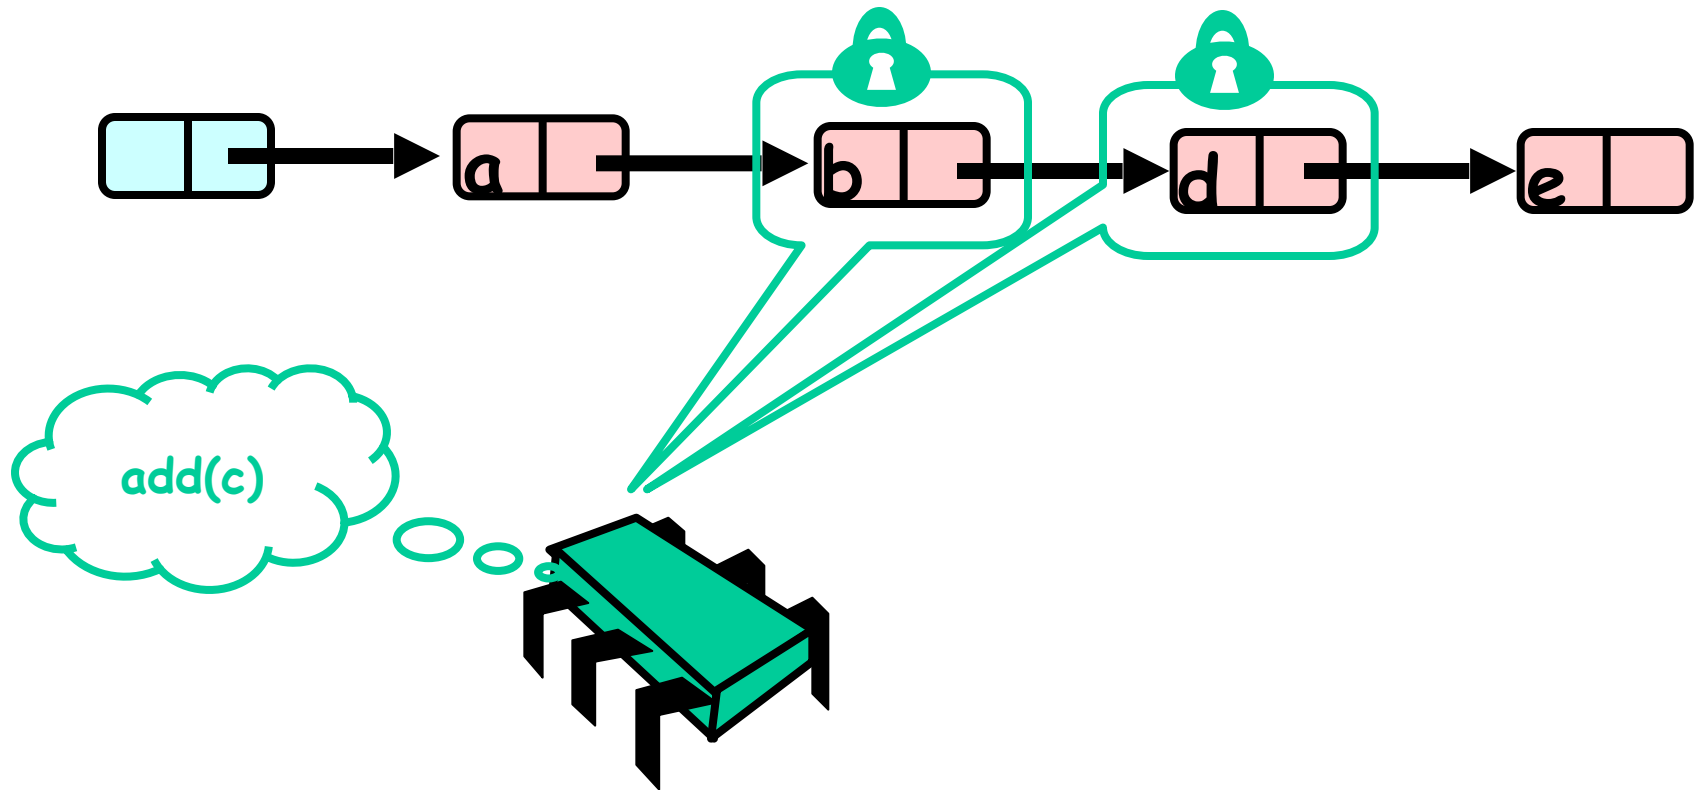
Optimistic Synchronization

- Find nodes without locking
- Lock nodes
- Check that everything is OK


Optimistic: Traverse without Locking




Optimistic: Lock and Load



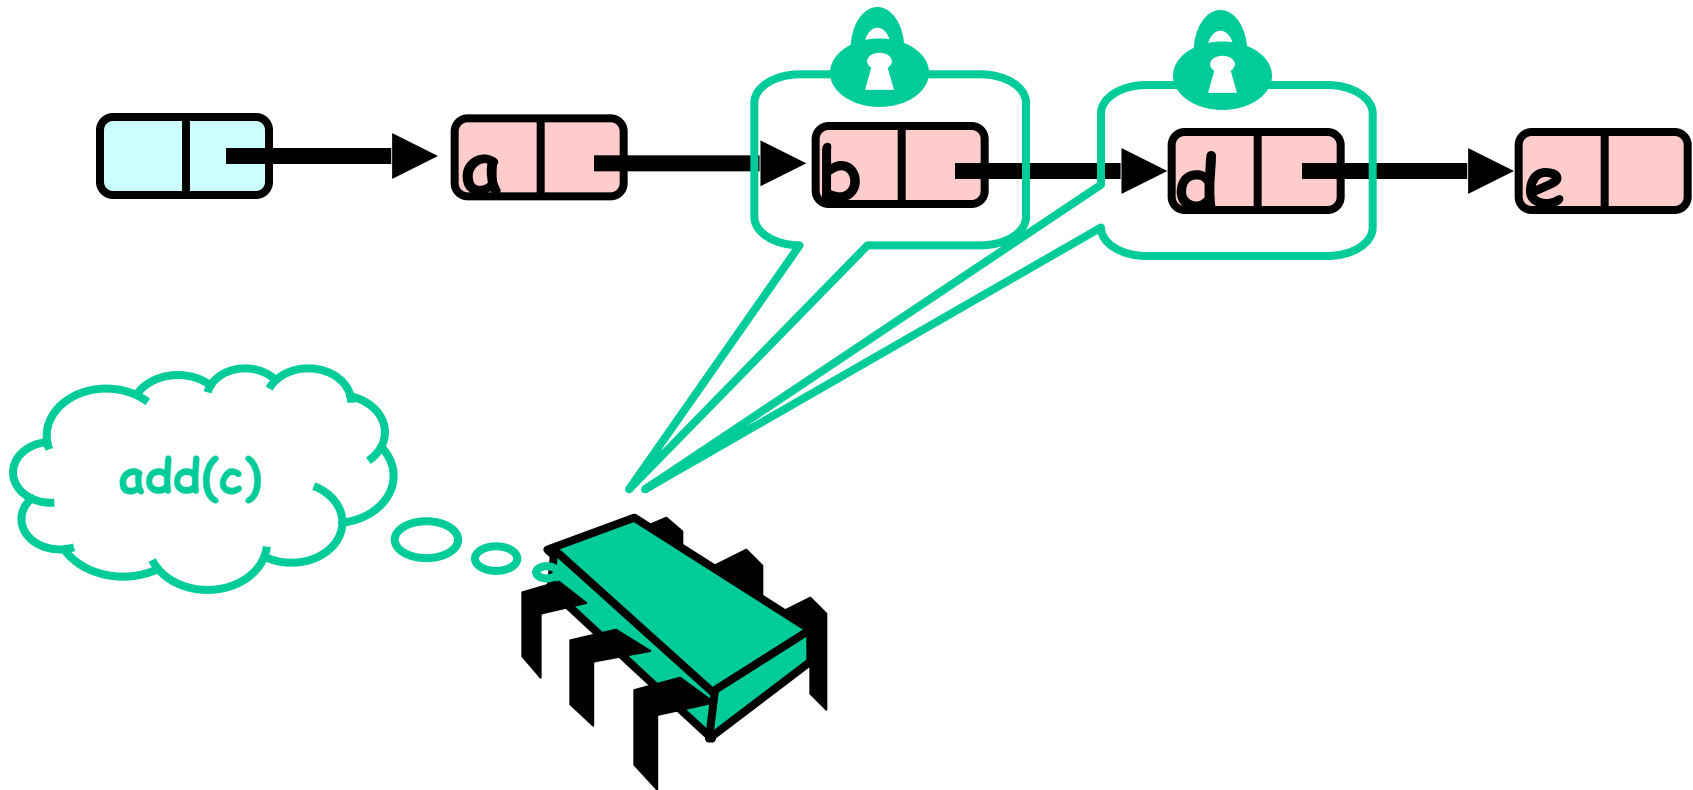
What could go wrong?



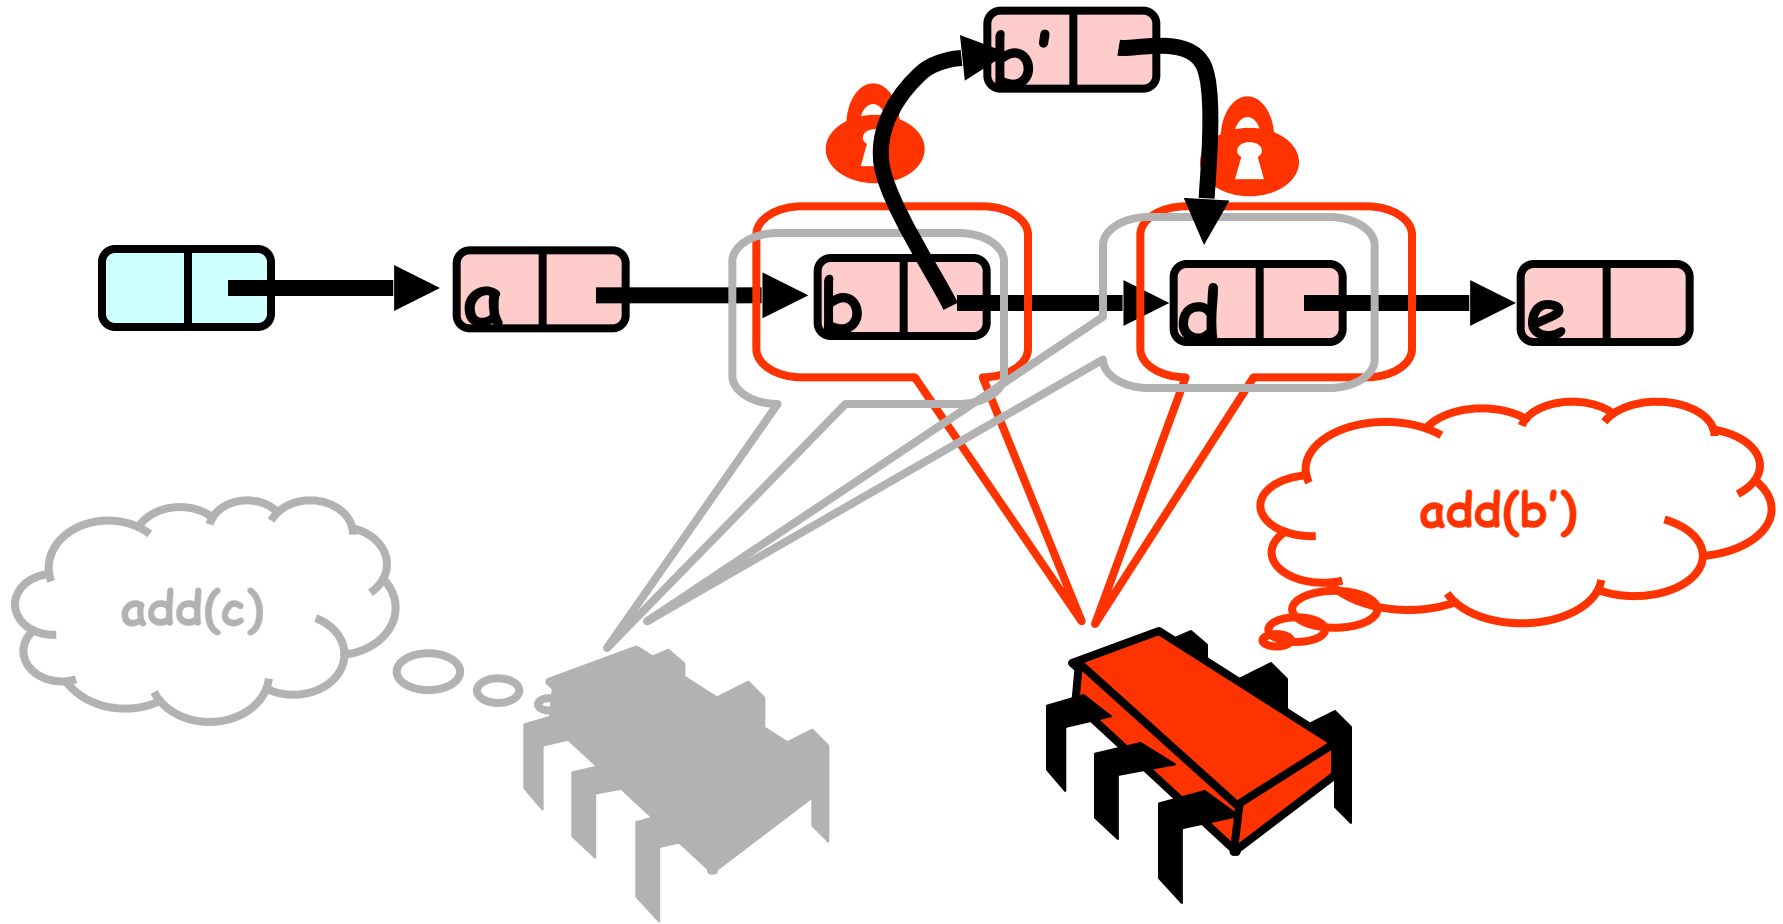
Validate - Part 1 (while holding locks)




What Else Can Go Wrong?



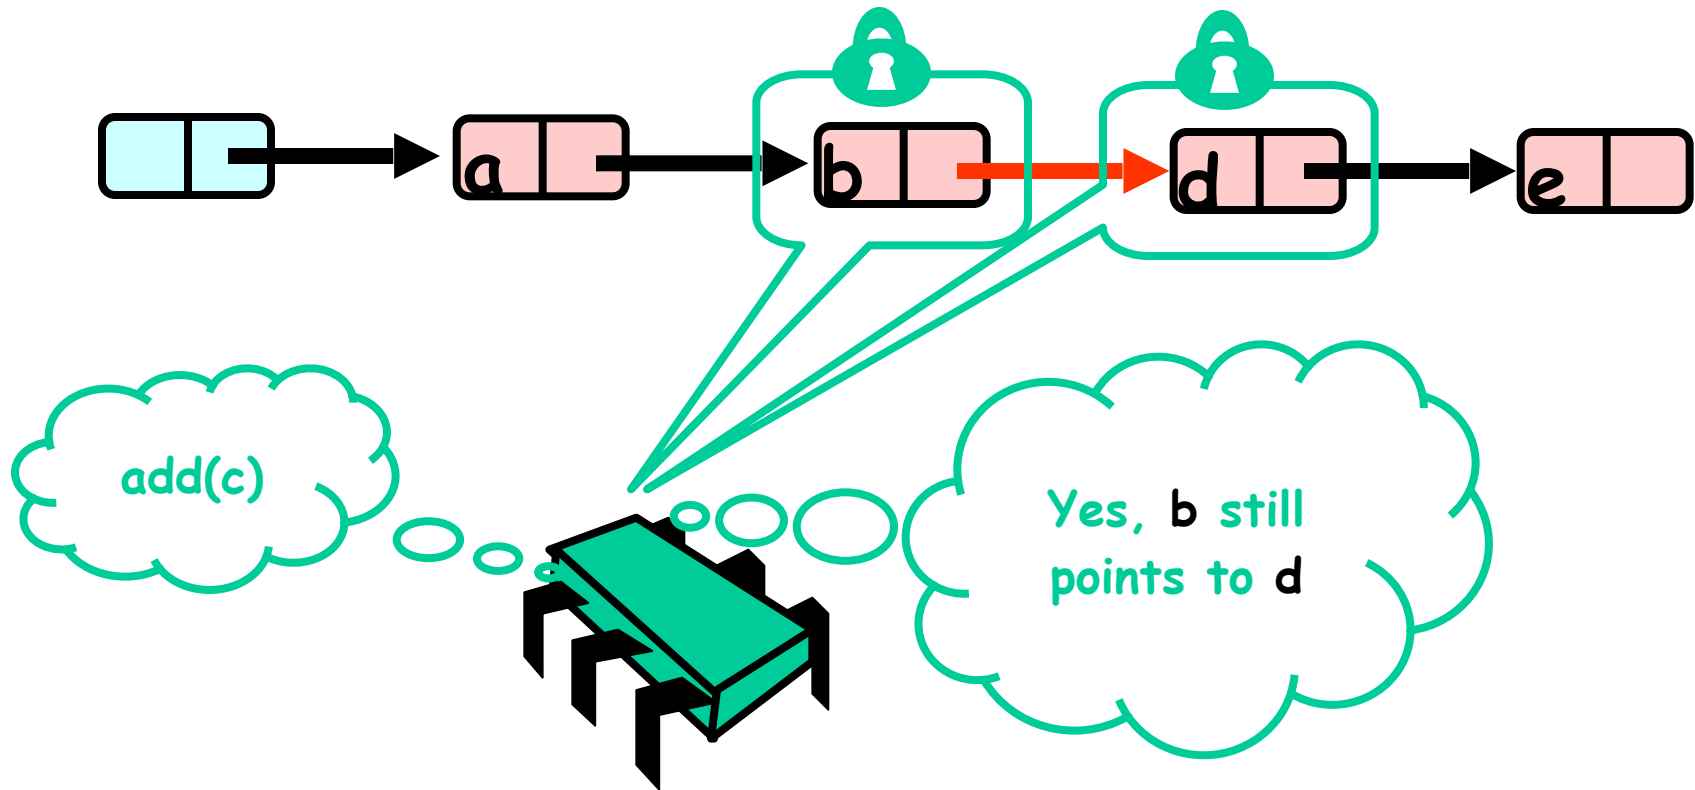
What Else Can Go Wrong?




What Else Can Go Wrong?



Validate Part 2 (while holding locks)



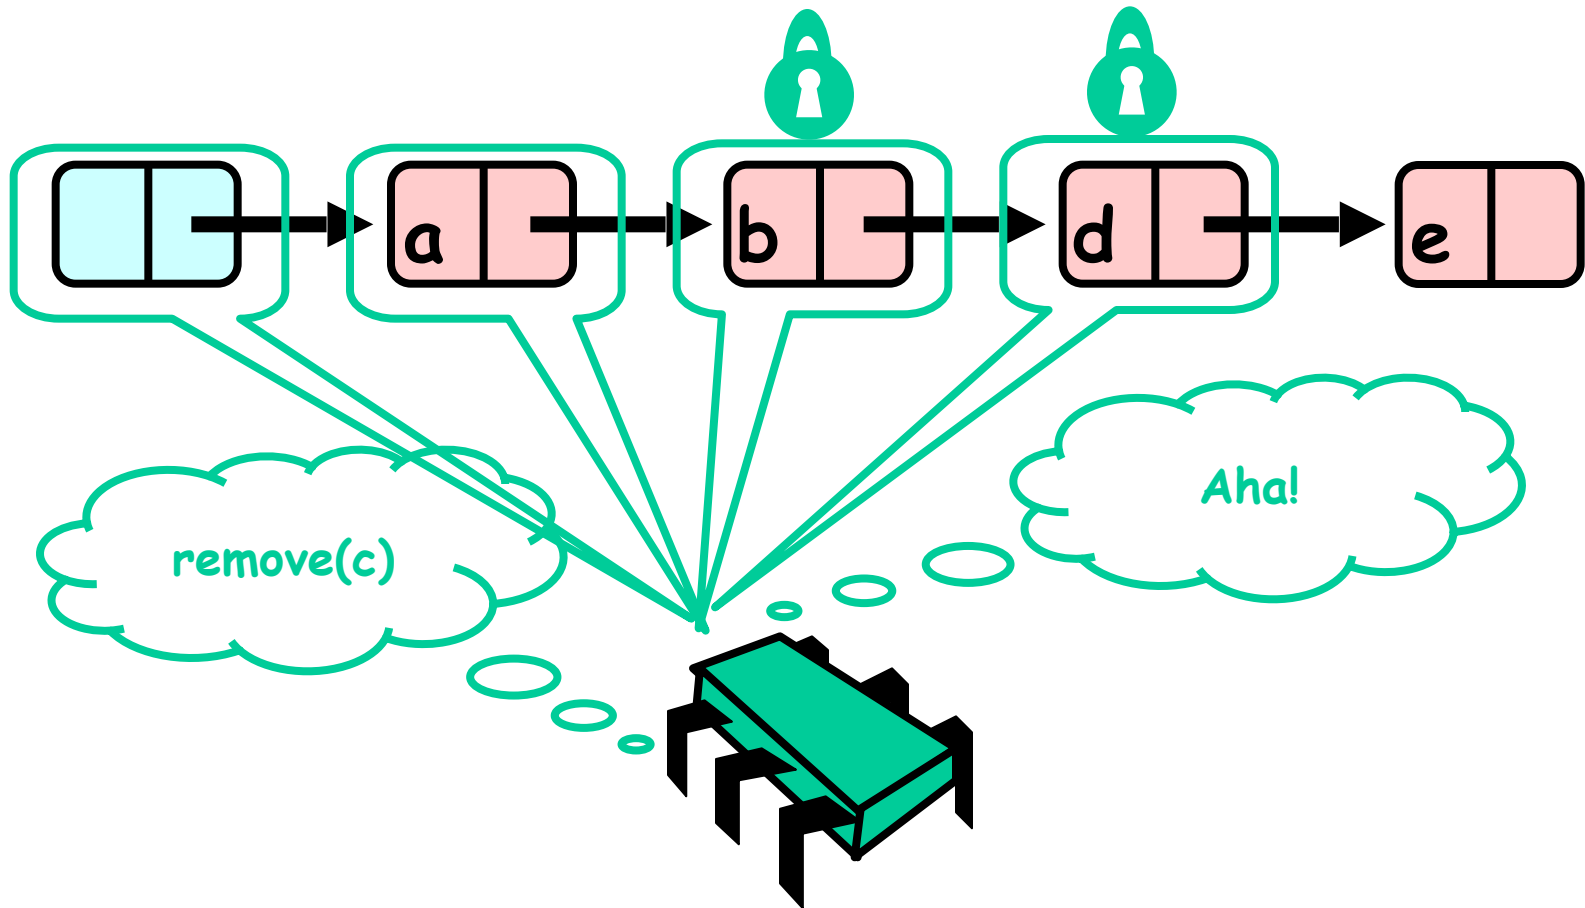
Optimistic: Critical Point




Correctness

- If
 - Nodes b and c both locked
 - Node b still accessible
 - Node c still successor to b
- Then
 - Neither will be deleted
 - OK to delete and return true

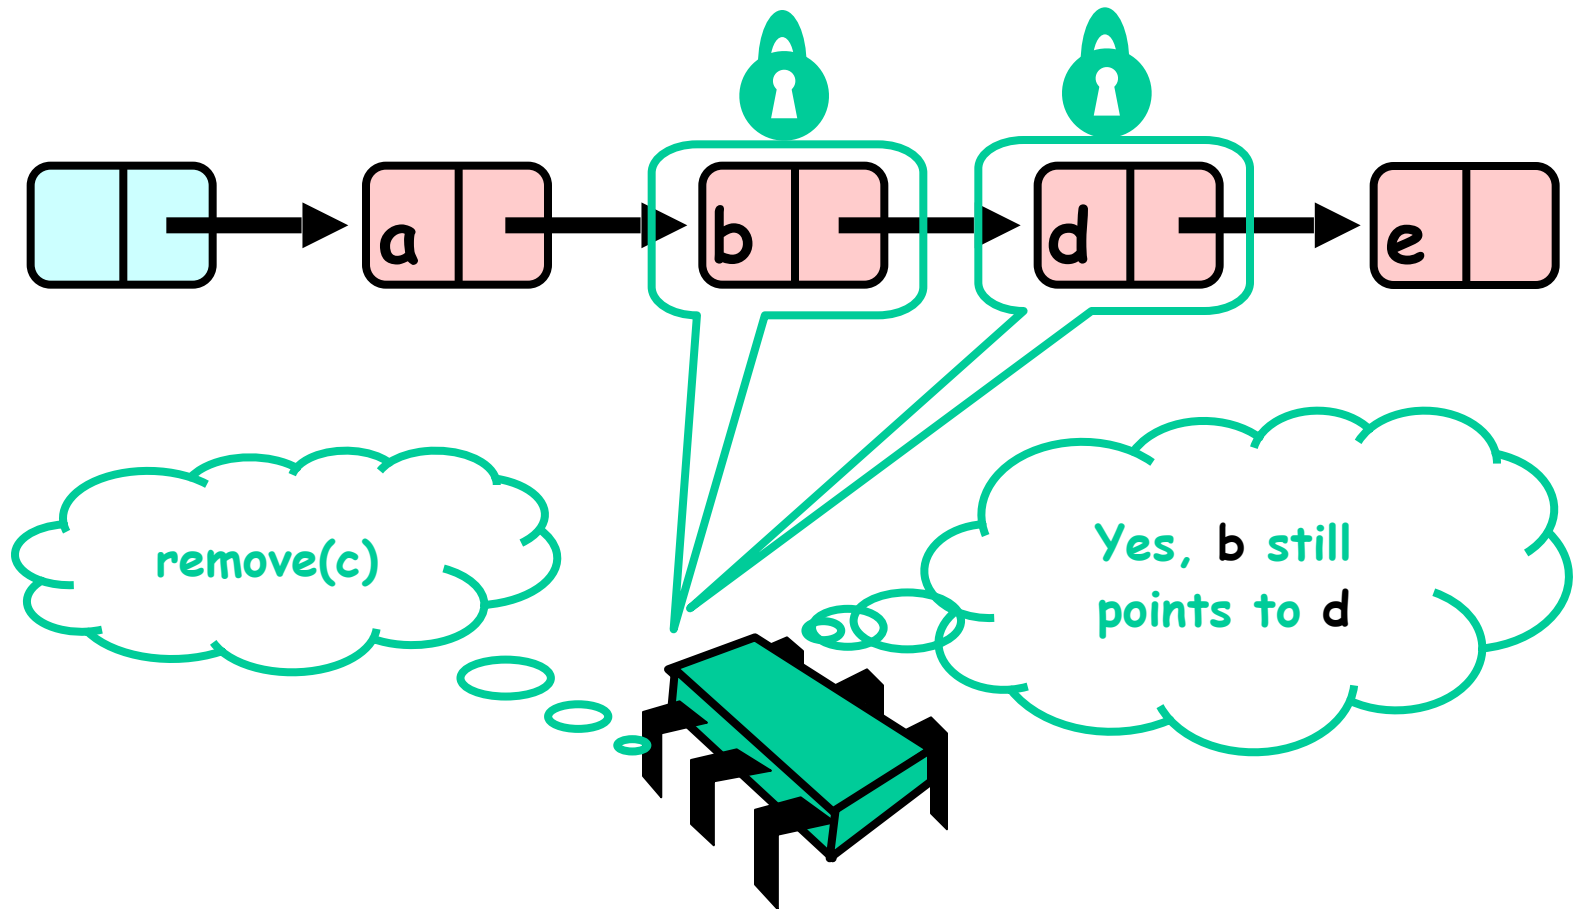
Unsuccessful Remove



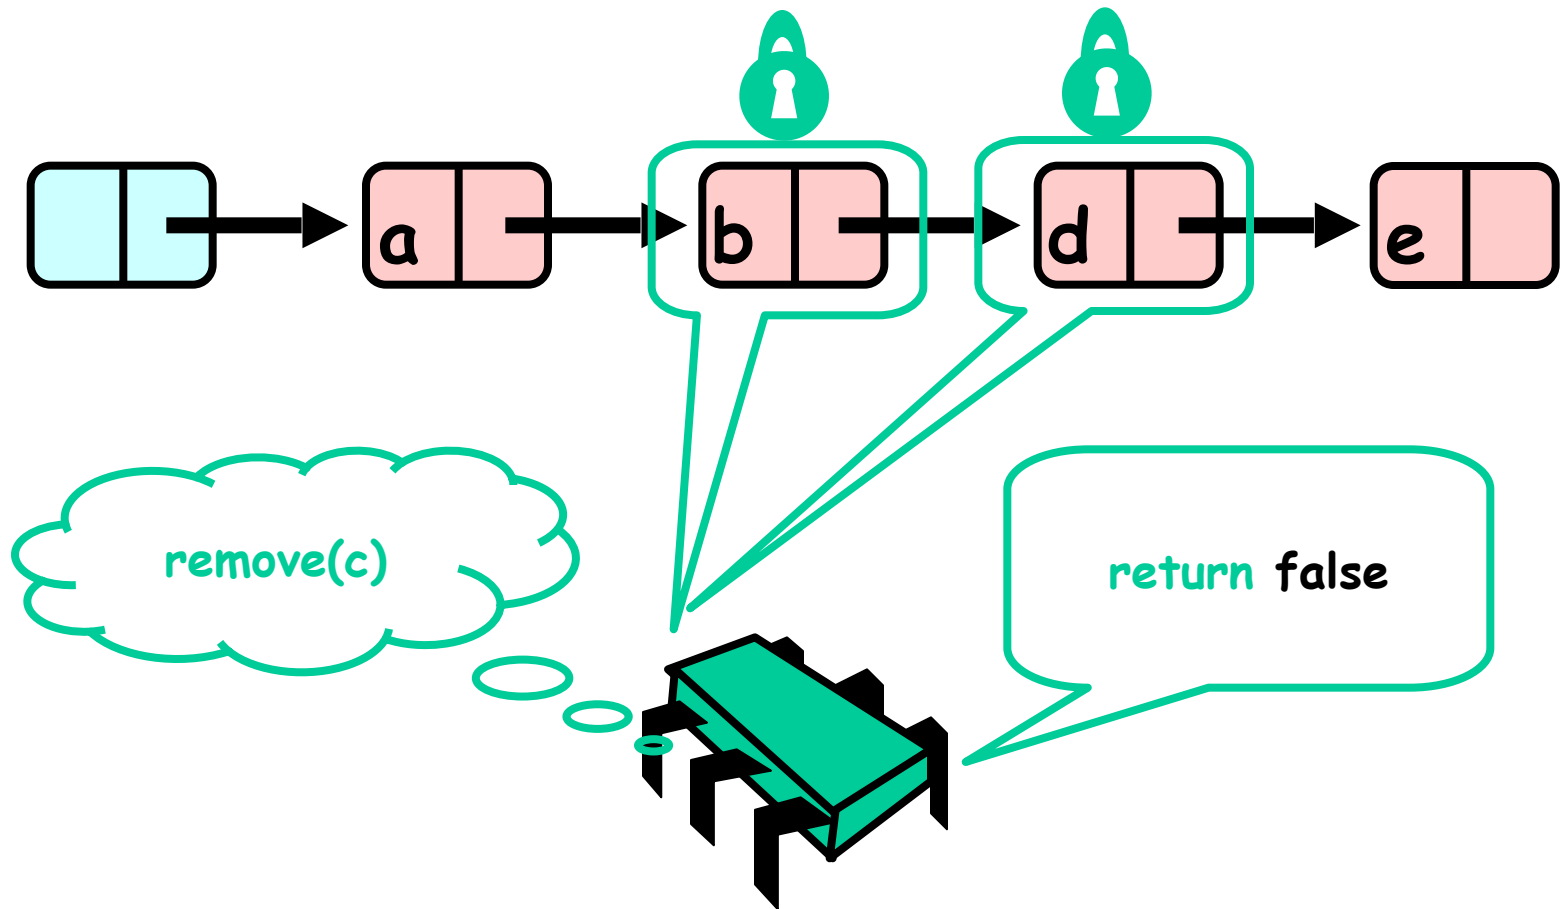
Validate (1)



Validate (2)



OK Computer



Correctness

- If
 - Nodes b and d both locked
 - Node b still accessible
 - Node d still successor to b
- Then
 - Neither will be deleted
 - No thread can add c after b
 - OK to return false


Validation

```
private boolean
validate(Node pred,
         Node curr) {
    Node node = head;
    while (node.key <= pred.key) {
        if (node == pred)
            return pred.next == curr;
        node = node.next;
    }
    return false;
}
```

Validation

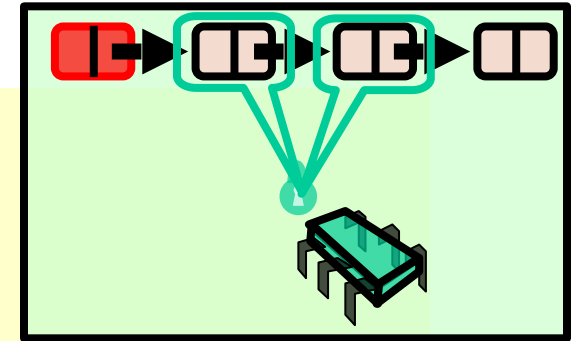
```
private boolean  
validate(Node pred,  
         Node curr) {  
    Node node = head;  
    while (node.key <= pred.key) {  
        if (node == pred)  
            return pred.next == curr;  
        node = node.next;  
    }  
    return false;  
}
```

**Predecessor &
current nodes**



Validation

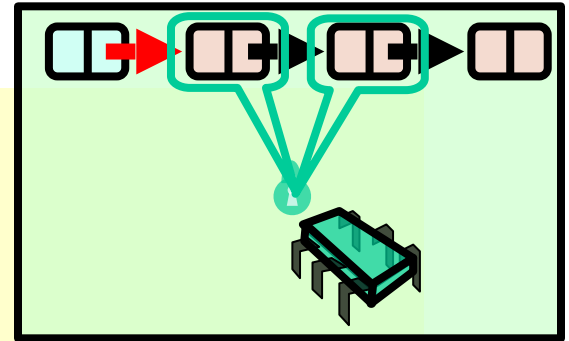
```
private boolean  
validate(Node pred,  
         Node curr) {  
Node node = head;  
while (node.key <= pred.key) {  
    if (node == pred)  
        return pred.next == curr;  
    node = node.next;  
}  
return false;  
}
```



**Begin at the
beginning**

Validation

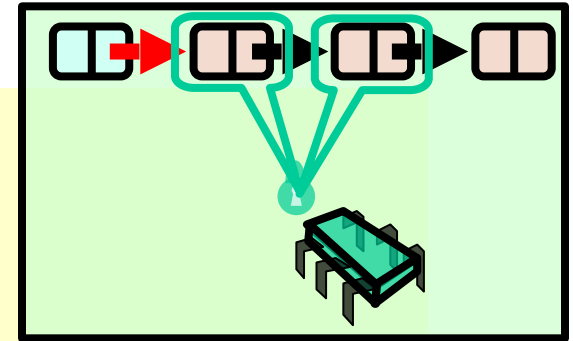
```
private boolean  
validate(Node pred,  
         Node curr) {  
    Node node = head;  
    while (node.key <= pred.key) {  
        if (node == pred)  
            return pred.next == curr;  
        node = node.next;  
    }  
    return false;  
}
```



Search range of keys

Validation

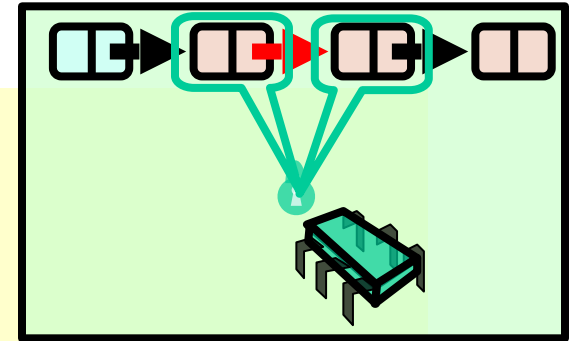
```
private boolean  
validate(Node pred,  
         Node curr) {  
    Node node = head;  
    while (node.key <= pred.key) {  
        if (node == pred)  
            return pred.next == curr;  
        node = node.next;  
    }  
    return false;  
}
```



Predecessor reachable

Validation

```
private boolean  
validate(Node pred,  
         Node curr) {  
    Node node = head;  
    while (node.key <= pred.key) {  
        if (node == pred)  
            return pred.next == curr;  
        node = node.next;  
    }  
    return false;  
}
```

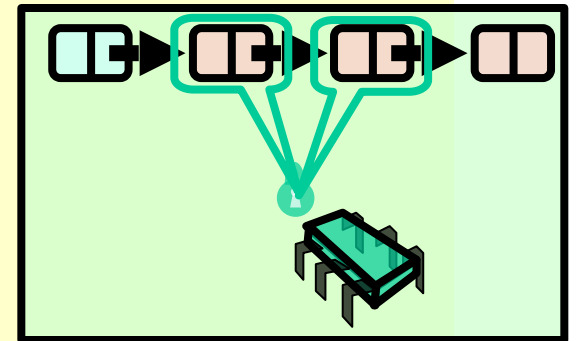


Is current node next?

Validation

```
private boolean
validate(Node pred,
        Node curr) {
    Node node = head;
    while (node.key <= pred.key) {
        if (node == pred)
            return pred.next == curr;
        node = node.next;
    }
    return false;
}
```

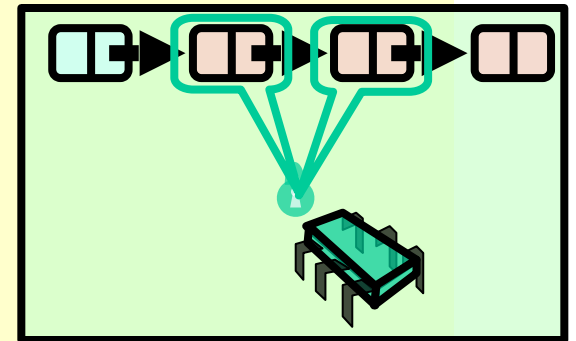
Otherwise move on



Validation

```
private boolean  
validate(Node pred,  
         Node curr) {  
    Node node = head;  
    while (node.key <= pred.key) {  
        if (node == pred)  
            return pred.next == curr;  
        node = node.next;  
    }  
    return false;  
}
```

Predecessor not reachable



Optimistic Synchronization

```
public boolean add(T item) {
    int key = item.hashCode();
    while (true) {
        Node pred = this.head;
        Node curr = pred.next;
        while (curr.key < key) {
            pred = curr; curr = curr.next;
        }
        pred.lock(); curr.lock();
        try {
            if (validate(pred, curr)) {
                if (curr.key == key) {
                    return false;
                } else {
                    Node node = new Entry(item);
                    entry.next = curr;
                    pred.next = node;
                    return true;
                }
            }
        } finally {
            pred.unlock(); curr.unlock();
        }
    }
}
```

```

public boolean remove(T item) {
    int key = item.hashCode();
    while (true) {
        Node pred = this.head;
        Node curr = pred.next;
        while (curr.key < key) {
            pred = curr; curr = curr.next;
        }
        pred.lock(); curr.lock();
        try {
            if (validate(pred, curr)) {
                if (curr.key == key) {
                    pred.next = curr.next;
                    return true;
                } else {
                    return false;
                }
            }
        } finally {
            pred.unlock(); curr.unlock();
        }
    }
}

```

```

public boolean contains(T item) {
    int key = item.hashCode();
    while (true) {
        Node pred = this.head;
        Node curr = pred.next;
        while (curr.key < key) {
            pred = curr; curr = curr.next;
        }
        try {
            pred.lock(); curr.lock();
            if (validate(pred, curr)) {
                return (curr.key == key);
            }
        } finally {
            pred.unlock(); curr.unlock();
        }
    }
}

```

```

private boolean validate(Node pred, Node
curr) {
    Node node = head;
    while (node.key <= pred.key) {
        if (node == pred)
            return pred.next == curr;
        Node = node.next;
    }
    return false;
}

```


Optimistic List

- Limited hot-spots
 - Targets of `add()`, `remove()`, `contains()`
 - No contention on traversals
- Moreover
 - Traversals are wait-free
 - Food for thought ...

So Far, So Good

- Much less lock acquisition/release
 - Performance
 - Concurrency
- Problems
 - Need to traverse list twice
 - contains() method acquires locks

Evaluation

- Optimistic is effective if
 - cost of scanning twice without locks is less than
 - cost of scanning once with locks
- Drawback
 - contains() acquires locks
 - 90% of calls in many apps

Lazy List

- Like optimistic, except
 - Scan once
 - contains(x) never locks ...
- Key insight
 - Removing nodes causes trouble
 - Do it "lazily"


Lazy List

- remove()
 - Scans list (as before)
 - Locks predecessor & current (as before)
- Logical delete
 - Marks current node as removed (new!)
- Physical delete
 - Redirects predecessor's next (as before)

Lazy Removal

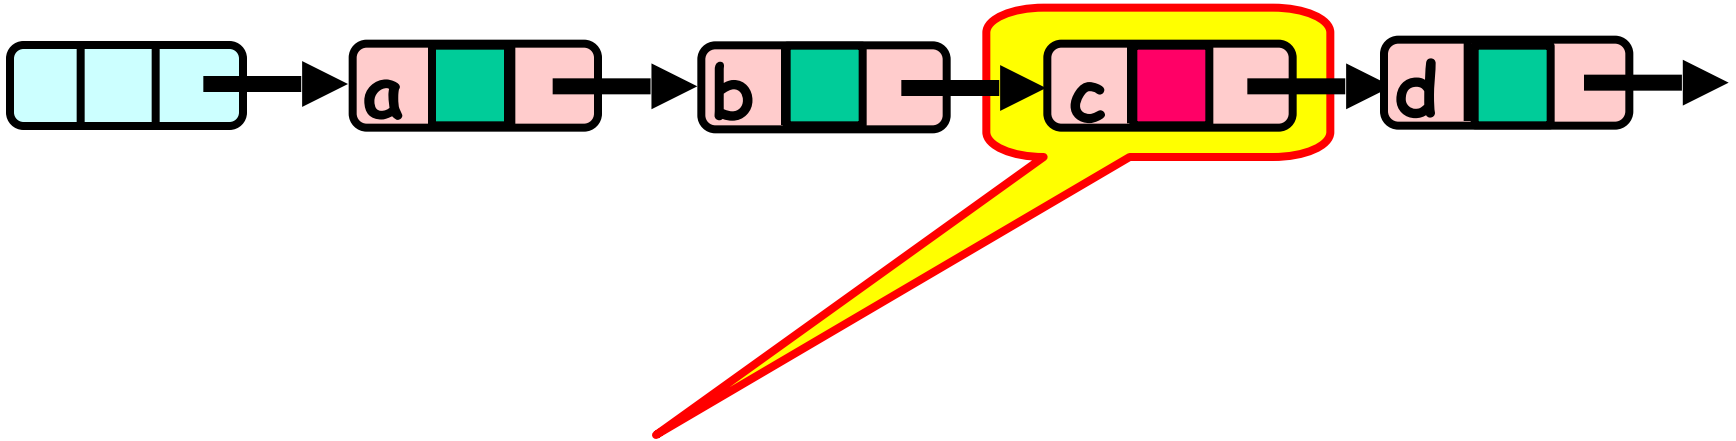


Lazy Removal



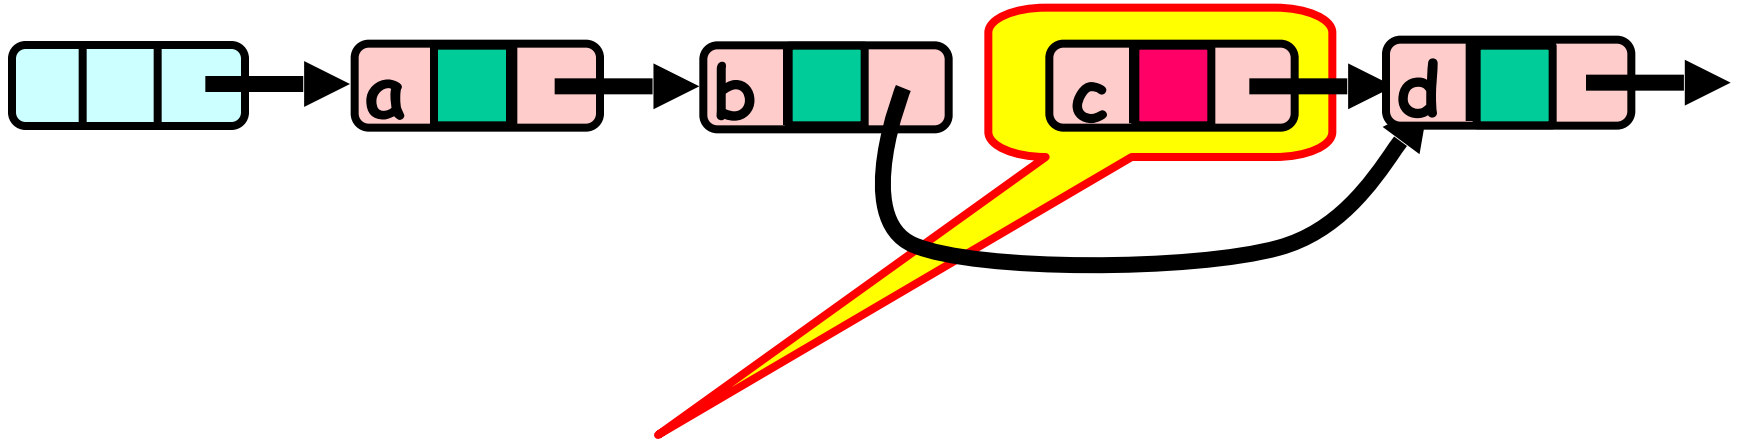
Present in list

Lazy Removal



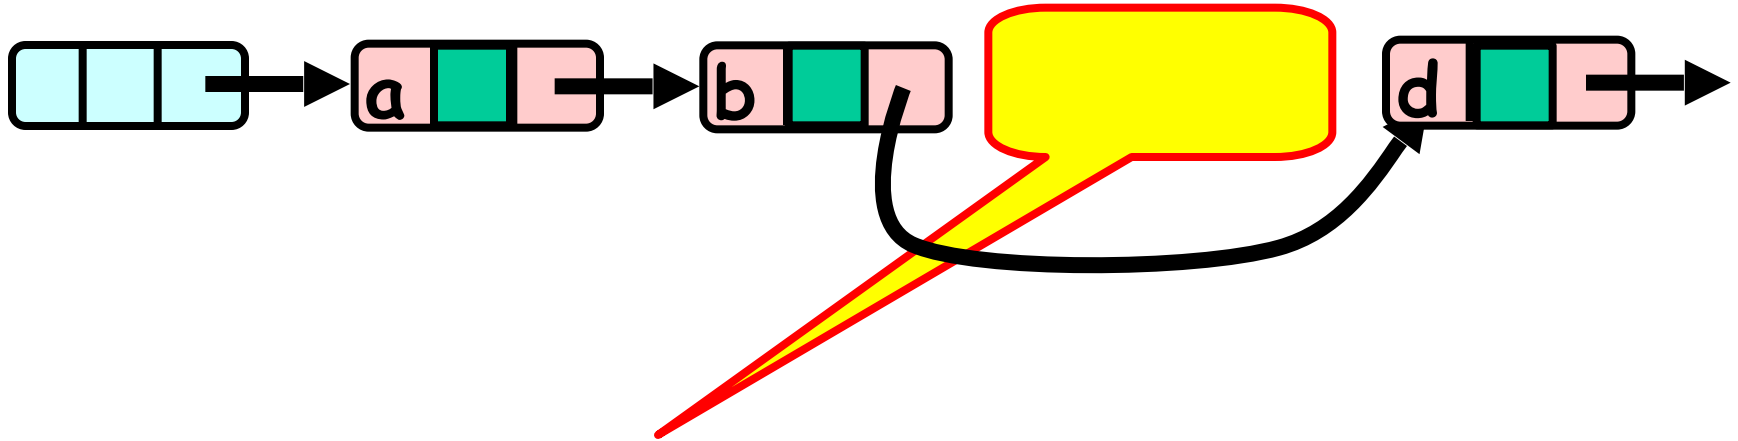
Logically deleted

Lazy Removal



Physically deleted

Lazy Removal



Physically deleted

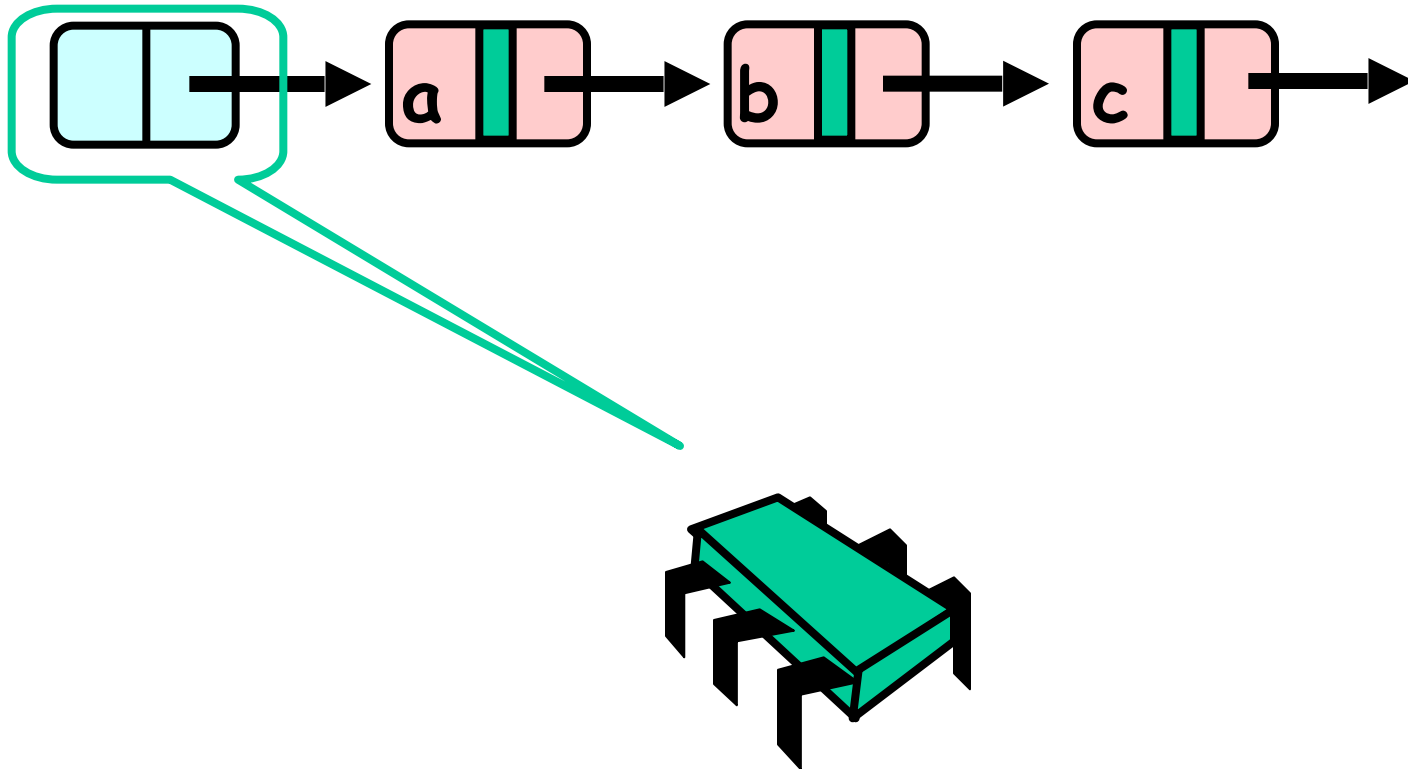
Lazy List

- All Methods
 - Scan through locked and marked nodes
 - Removing a node doesn't slow down other method calls ...
- Must still lock pred and curr nodes.

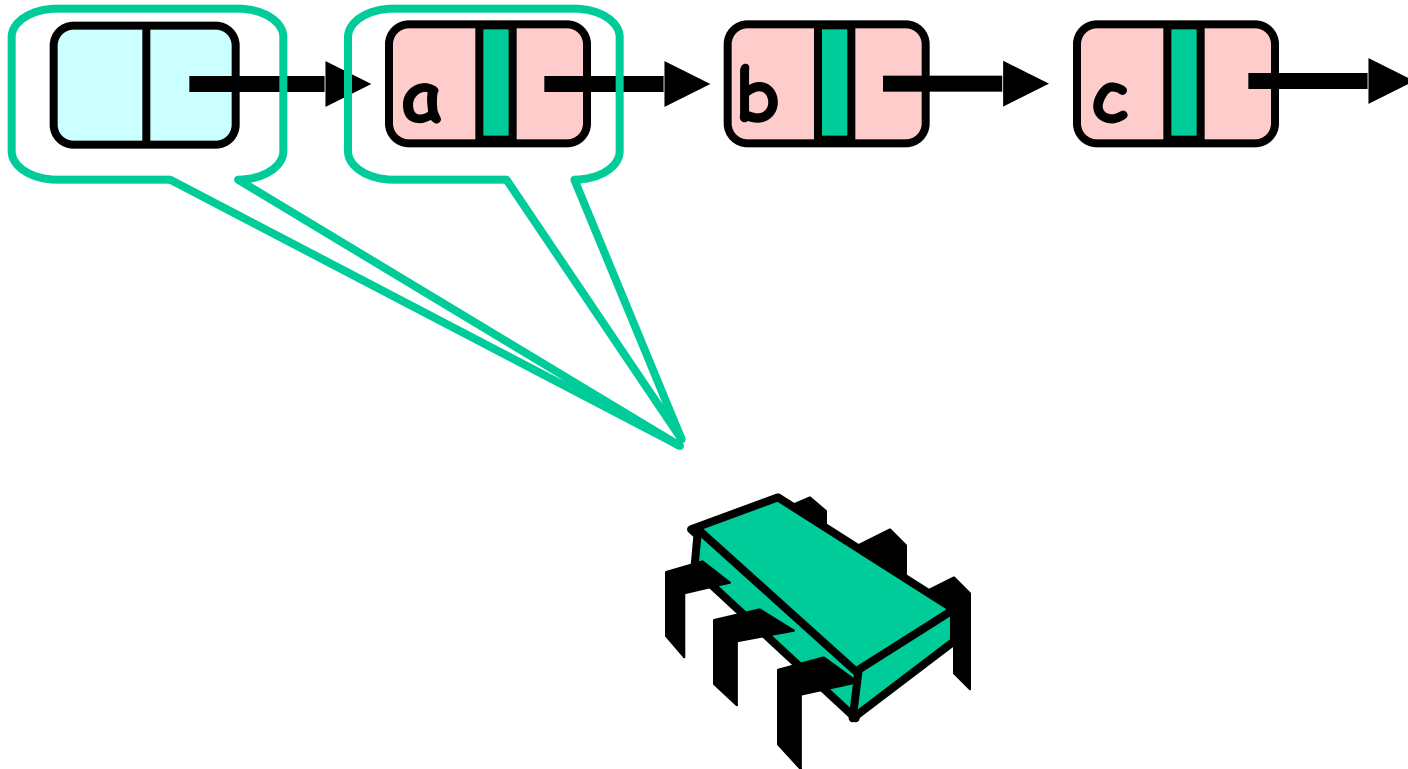
Validation

- No need to rescan list!
- Check that pred is not marked
- Check that curr is not marked
- Check that pred points to curr

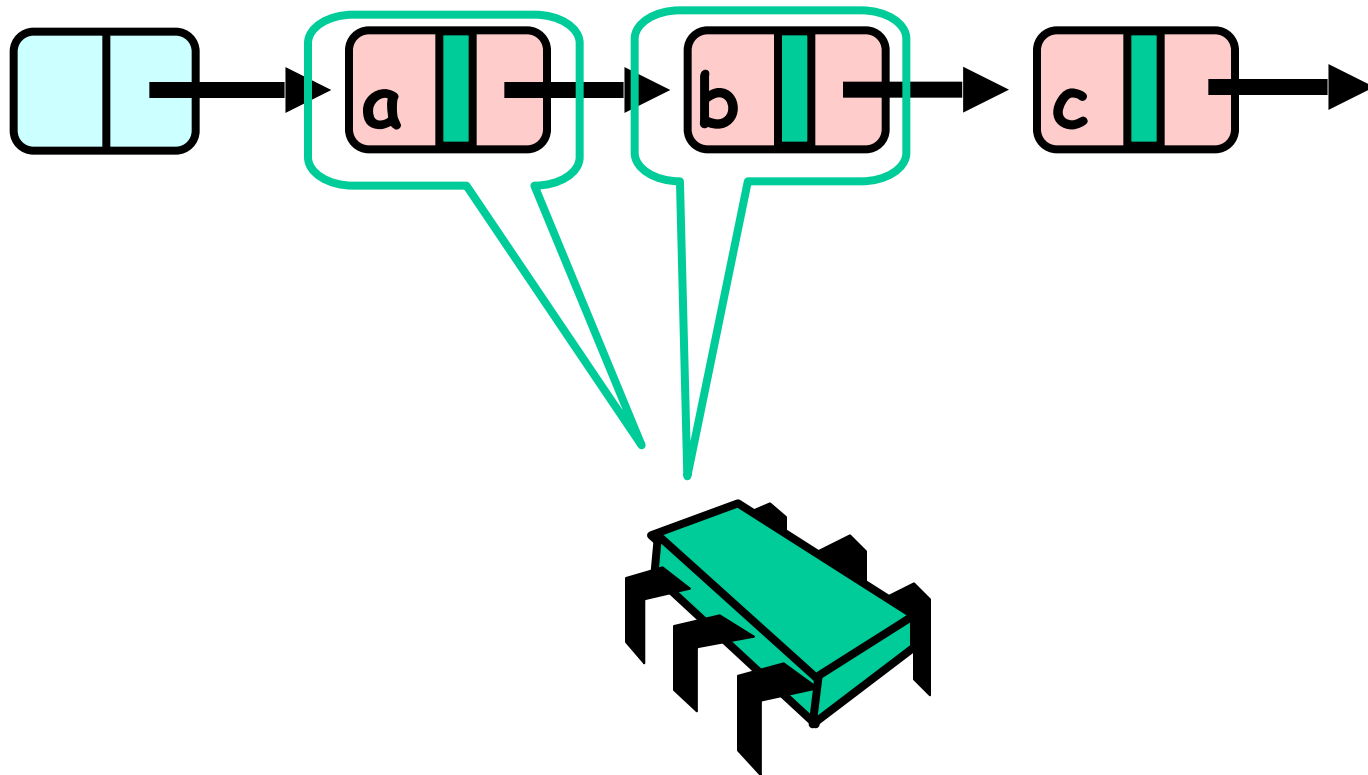
Business as Usual




Business as Usual



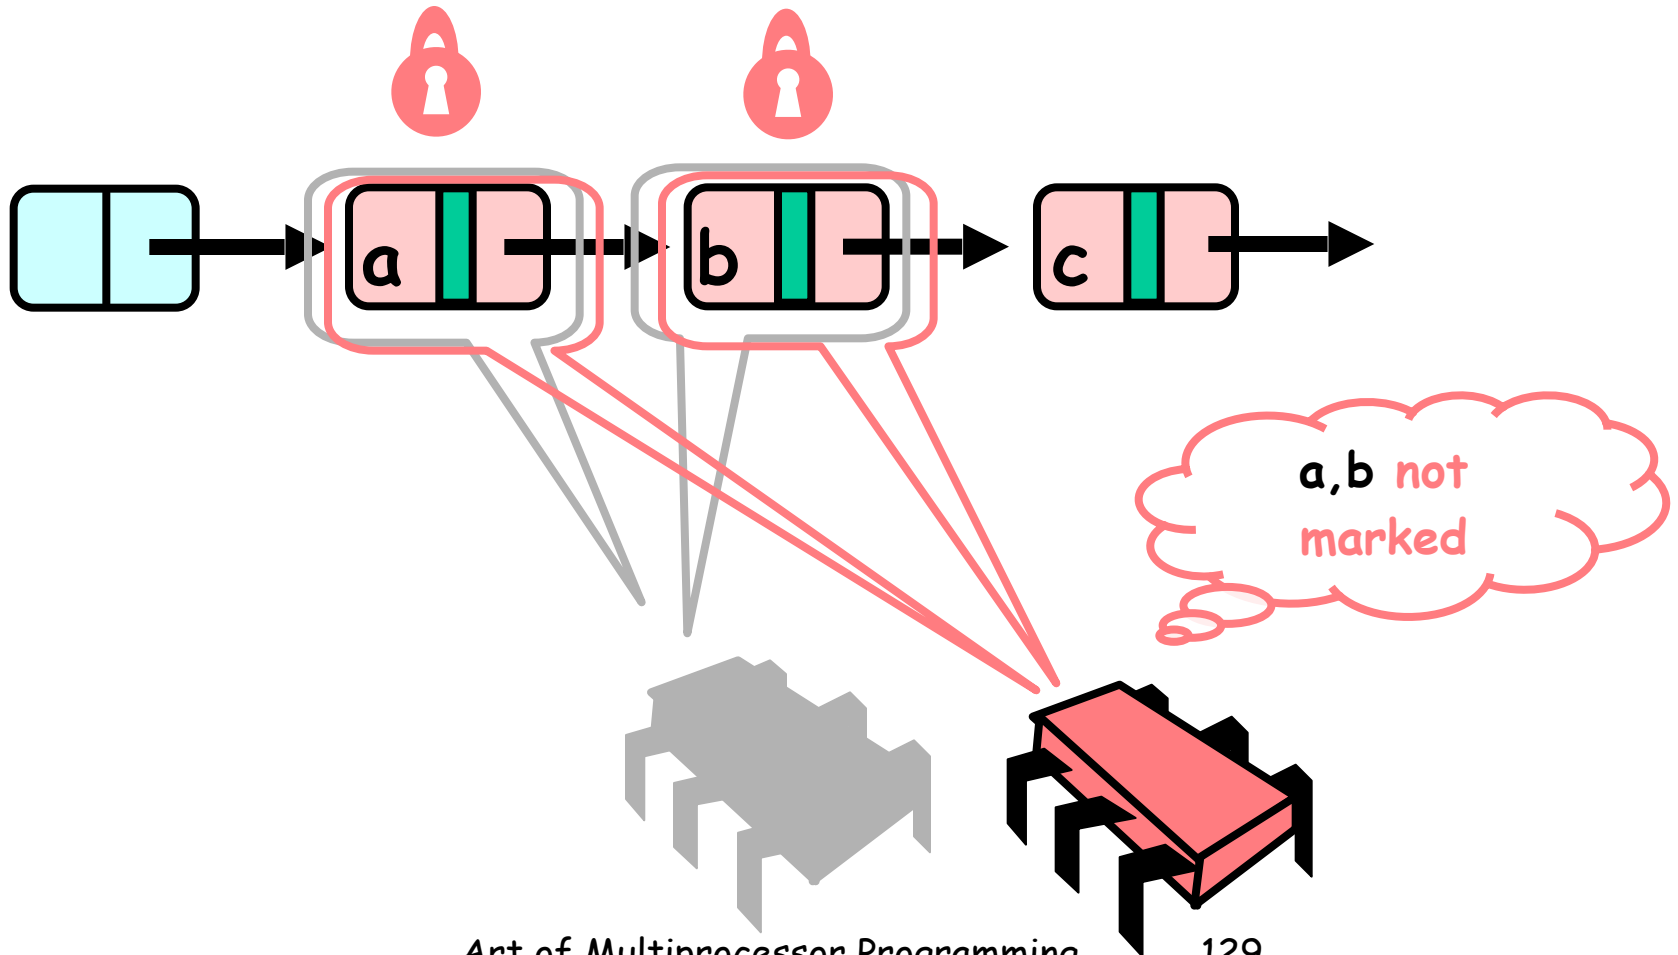
Business as Usual



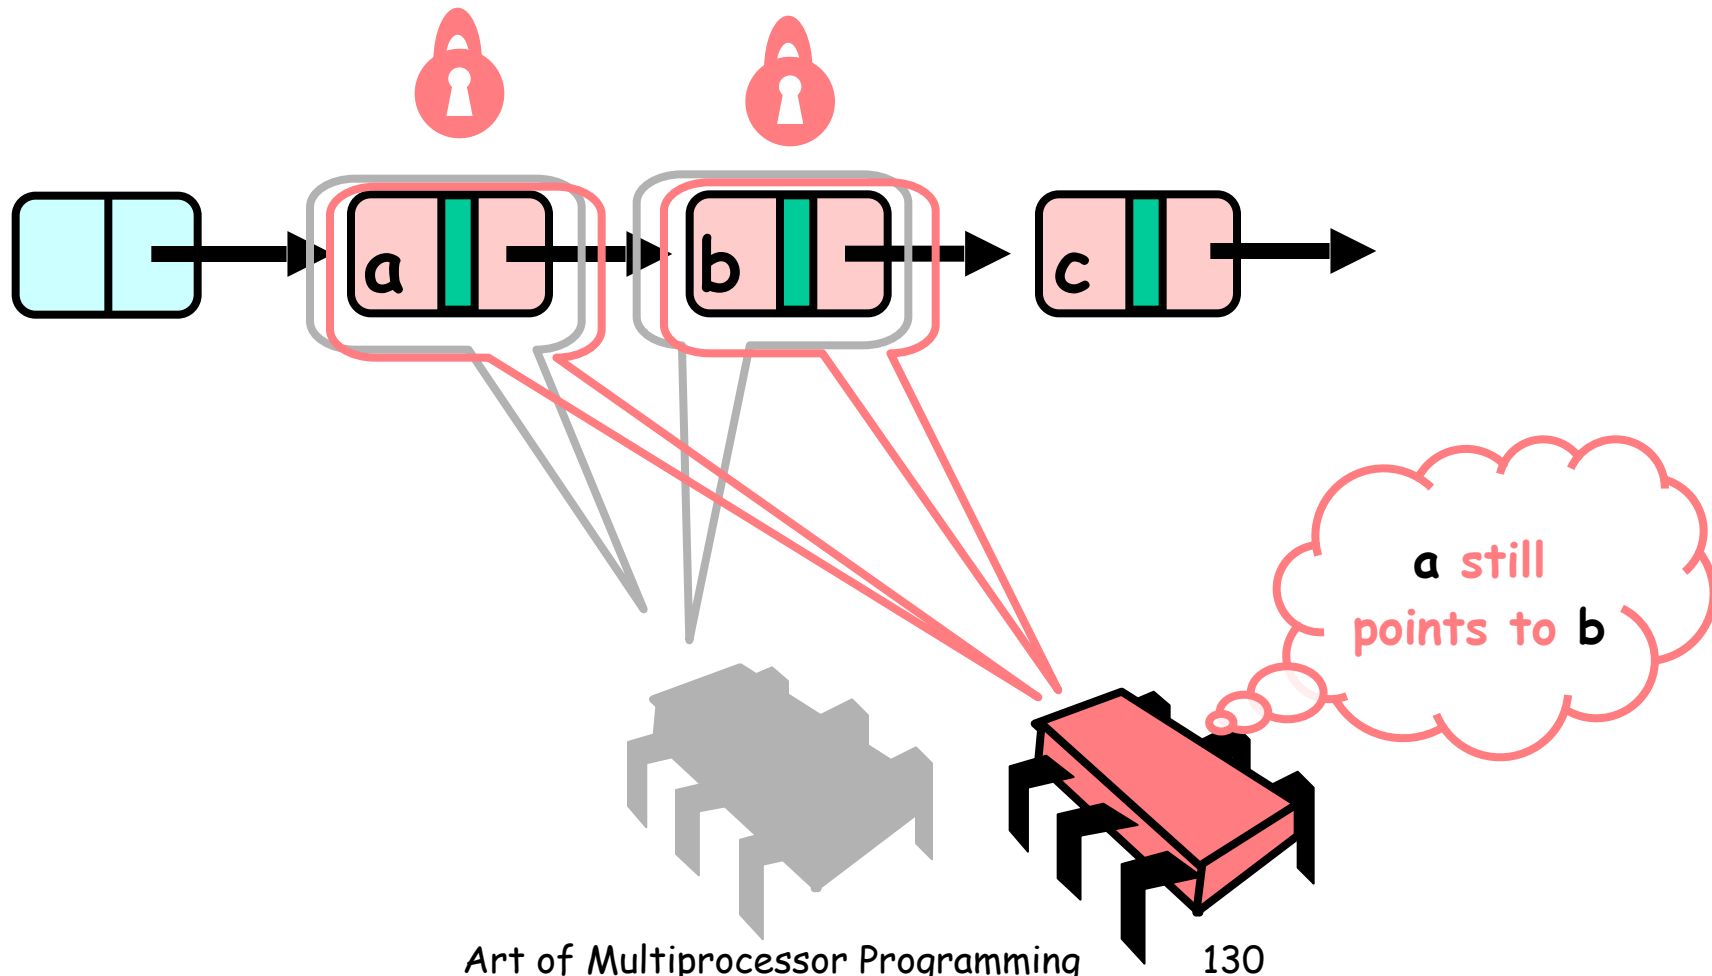
Business as Usual



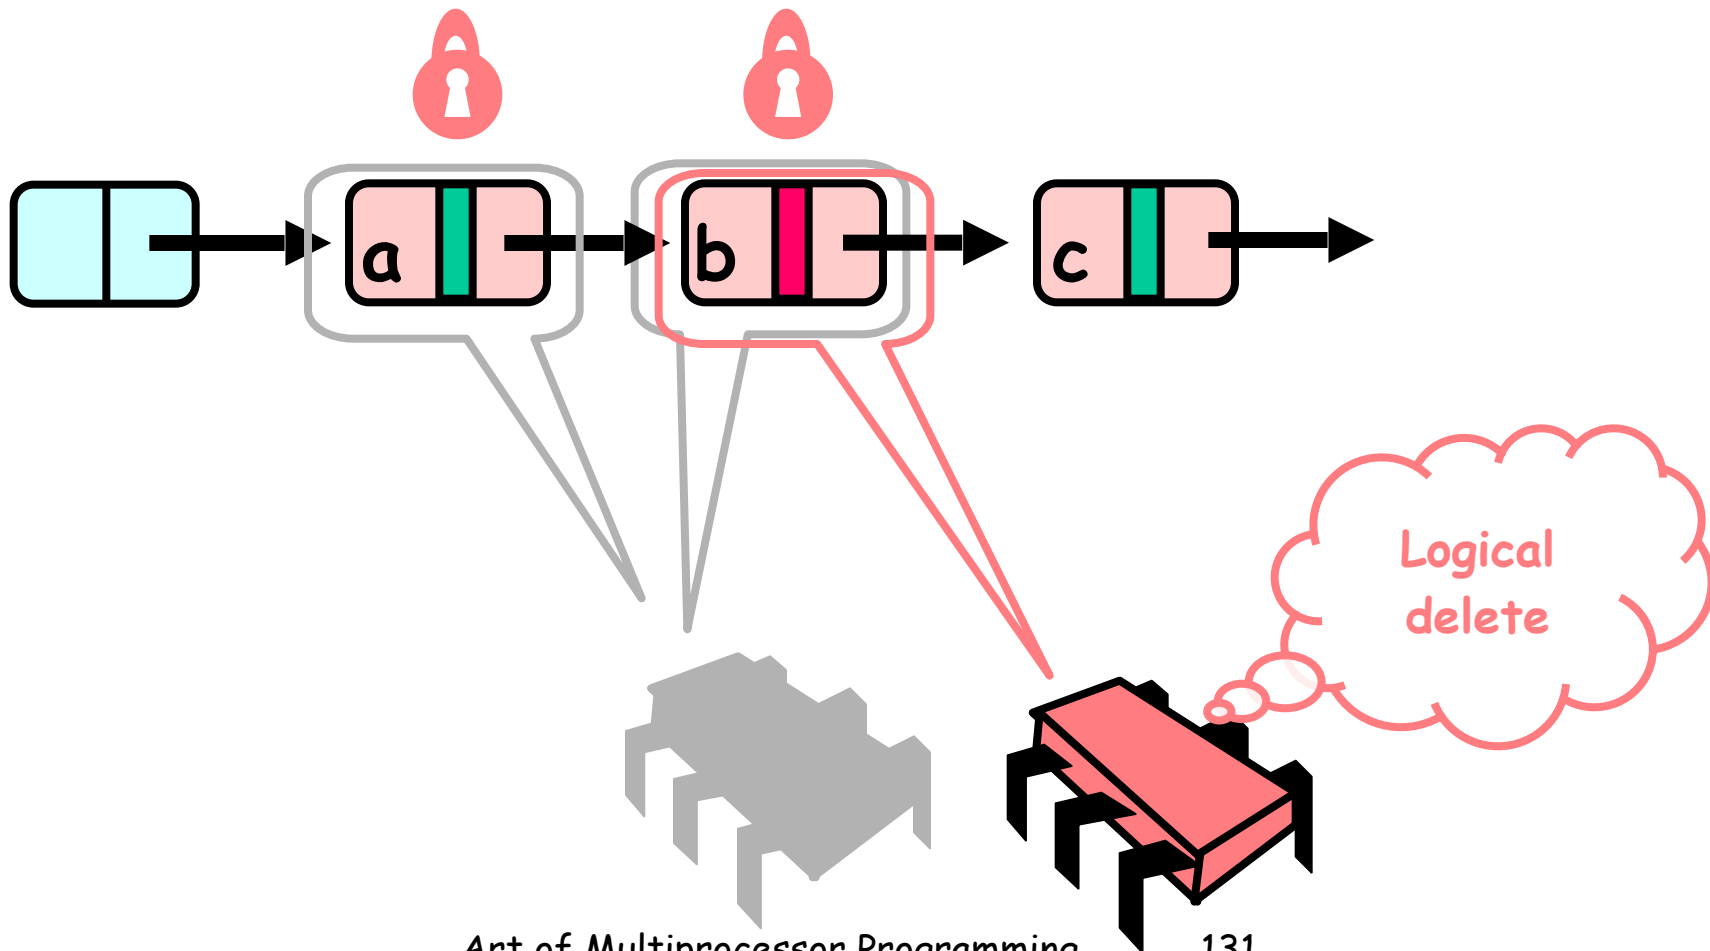
Business as Usual




Business as Usual



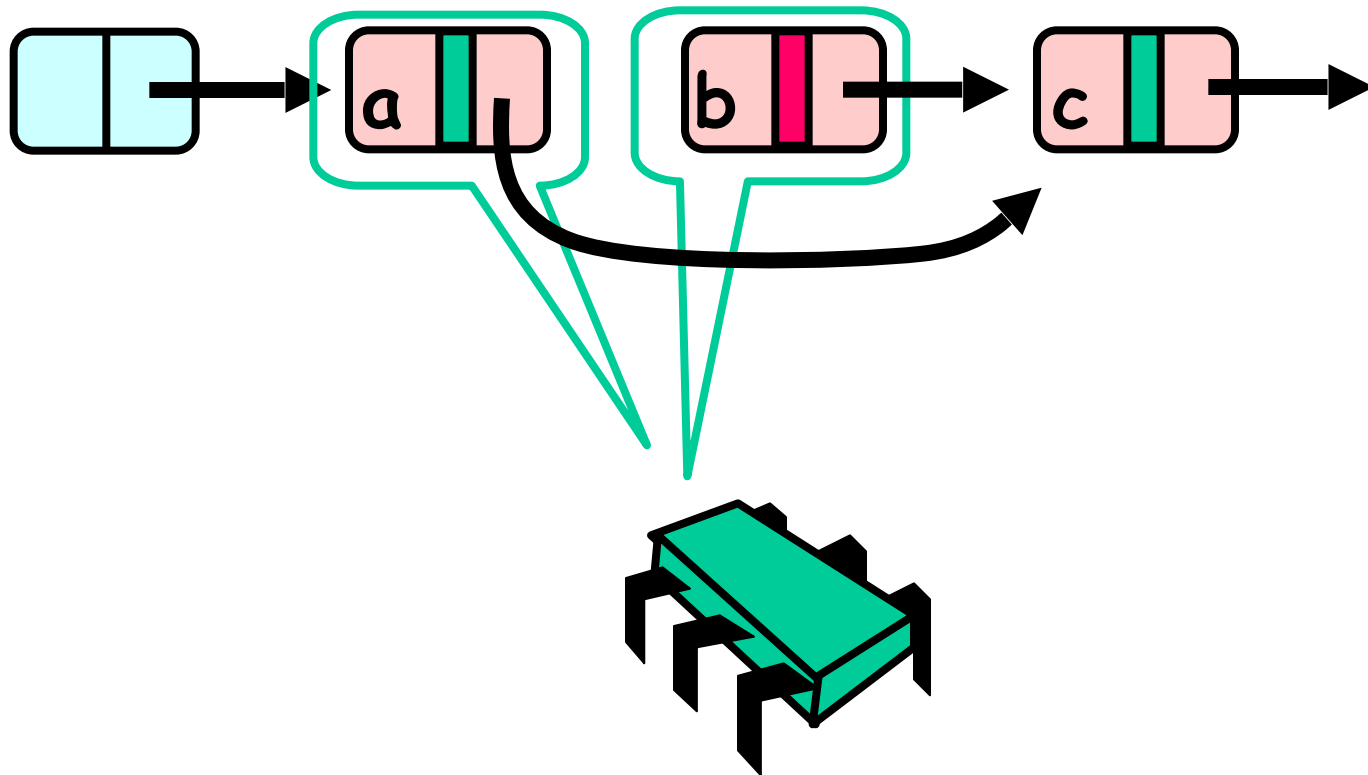
Business as Usual



Business as Usual



Business as Usual



Invariant

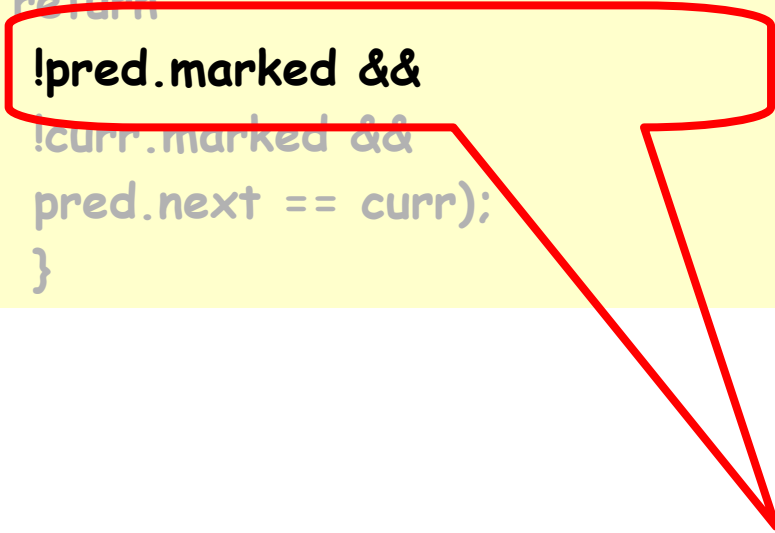
- If not marked then item in the set
- and reachable from head
- and if not yet traversed it is reachable from pred

Validation

```
private boolean  
    validate(Node pred, Node curr) {  
return  
    !pred.marked &&  
    !curr.marked &&  
    pred.next == curr);  
}
```

List Validate Method

```
private boolean  
    validate(Node pred, Node curr) {  
return  
    !pred.marked &&  
    !curr.marked &&  
    pred.next == curr);  
}
```



**Predecessor not
Logically removed**

List Validate Method

```
private boolean  
    validate(Node pred, Node curr) {  
    return  
        !pred.marked &&  
        !curr.marked &&  
        pred.next == curr);  
}
```



**Current not
Logically removed**

List Validate Method

```
private boolean  
    validate(Node pred, Node curr) {  
    return  
        !pred.marked &&  
        !curr.marked &&  
        pred.next == curr);  
}
```

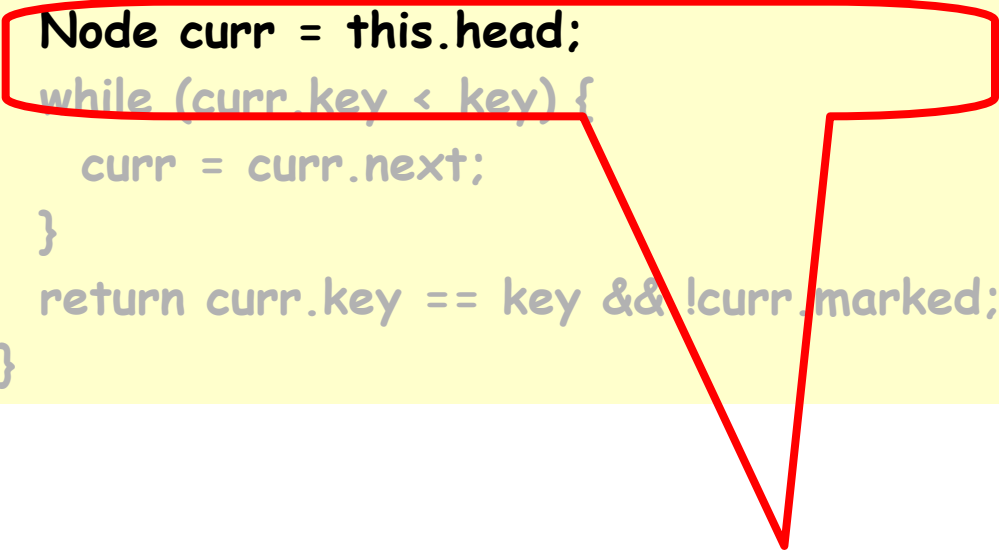
**Predecessor still
Points to current**

Contains

```
public boolean contains(Item item) {  
    int key = item.hashCode();  
    Node curr = this.head;  
    while (curr.key < key) {  
        curr = curr.next;  
    }  
    return curr.key == key && !curr.marked;  
}
```

Contains

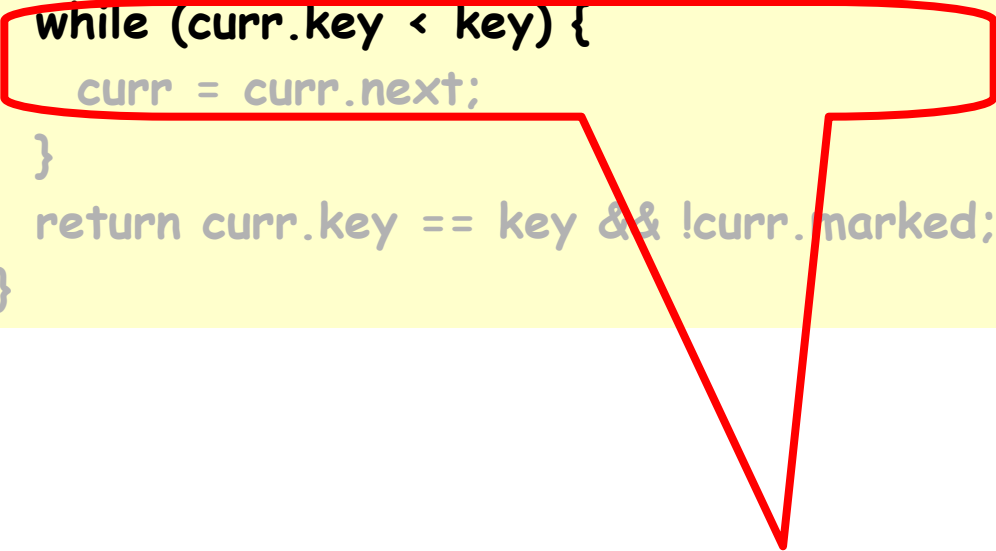
```
public boolean contains(Item item) {  
    int key = item.hashCode();  
    Node curr = this.head;  
    while (curr.key < key) {  
        curr = curr.next;  
    }  
    return curr.key == key && !curr.marked;  
}
```



Start at the head

Contains

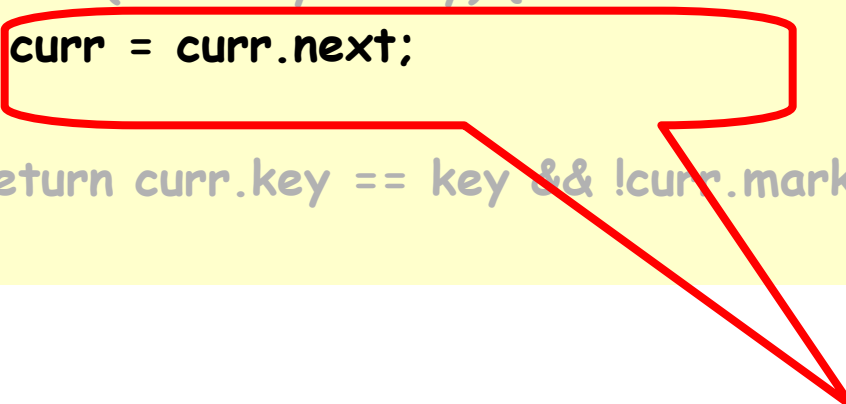
```
public boolean contains(Item item) {  
    int key = item.hashCode();  
    Node curr = this.head;  
    while (curr.key < key) {  
        curr = curr.next;  
    }  
    return curr.key == key && !curr.marked;  
}
```



Search key range

Contains


```
public boolean contains(Item item) {  
    int key = item.hashCode();  
    Node curr = this.head;  
    while (curr.key < key) {  
        curr = curr.next;  
    }  
    return curr.key == key && !curr.marked;  
}
```



**Traverse without locking
(nodes may have been removed)**

Contains

```
public boolean contains(Item item) {  
    int key = item.hashCode();  
    Node curr = this.head;  
    while (curr.key < key) {  
        curr = curr.next;  
    }  
    return curr.key == key && !curr.marked;  
}
```



Present and undeleted?

```

public boolean add(T item) {
    int key = item.hashCode();
    while (true) {
        Node pred = this.head;
        Node curr = head.next;
        while (curr.key < key) {
            pred = curr; curr = curr.next;
        }
        pred.lock();
        try {
            curr.lock();
            try {
                if (validate(pred, curr)) {
                    if (curr.key == key) {
                        return false;
                    } else {
                        Node Node = new Node(item);
                        Node.next = curr;
                        pred.next = Node;
                        return true;
                    }
                }
            }
        } finally { // always unlock
            curr.unlock();
        }
    } finally { // always unlock
        pred.unlock();
    }
}
}

```

Lazy Synchronization


```

public boolean remove(T item) {
    int key = item.hashCode();
    while (true) {
        Node pred = this.head;
        Node curr = head.next;
        while (curr.key < key) {
            pred = curr; curr = curr.next;
        }
        pred.lock();
        try {
            curr.lock();
            try {
                if (validate(pred, curr)) {
                    if (curr.key != key) {
                        return false;
                    } else {
                        curr.marked = true;
                        pred.next = curr.next;
                        return true;
                    }
                }
            } finally {
                curr.unlock();
            }
        } finally {
            pred.unlock();
        }
    }
}

```

```

public boolean contains(T item) {
    int key = item.hashCode();
    Node curr = this.head;
    while (curr.key < key)
        curr = curr.next;
    return curr.key == key && !curr.marked;
}

```

```

private boolean validate(Node pred, Node
curr) {
    return !pred.marked && !curr.marked &&
pred.next == curr;
}

```

Evaluation

- Good:
 - `contains()` doesn't lock
 - Good because typically high % `contains()`
 - Uncontended calls don't re-traverse
- Bad
 - Contended `add()` and `remove()` calls do re-traverse
 - Traffic jam if one thread delays

Traffic Jam

- Any concurrent data structure based on mutual exclusion has a weakness
- If one thread
 - Enters critical section
 - And "eats the big muffin"
 - Cache miss, page fault, descheduled ...
 - Everyone else using that lock is stuck!
 - Need to trust the scheduler....

Reminder: Lock-Free Data Structures



- No matter what ...
 - Guarantees minimal progress in any execution
 - i.e. Some thread will always complete a method call, even if others halt at malicious times
 - Implies that implementation can't use locks