


# MPCS 52060 - Parallel Programming

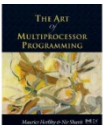
## M6: Advanced Parallel Patterns and Techniques (Part 1)



Original slides from “The Art of Multiprocessor Programming” by Maurice Herlihy & Nir Shavit with modifications by Lamont Samuels

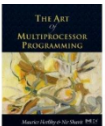
# How to write Parallel Apps?

- Split a program into parallel parts
- In an effective way
- Thread management
- We've seen a few ways so far to do this
  - Fork-Join (Static, Map)
  - Producer/Consumer model with global queue
  - Pipeline (M6)



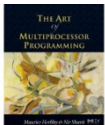
# Matrix Multiplication

$$(C) = (A) \cdot (B)$$



# Matrix Multiplication

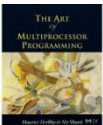
```
class Worker extends Thread {  
    int row, col;  
    Worker(int row, int col) {  
        this.row = row; this.col = col;  
    }  
    public void run() {  
        double dotProduct = 0.0;  
        for (int i = 0; i < n; i++)  
            dotProduct += a[row][i] * b[i][col];  
        c[row][col] = dotProduct;  
    }  
}
```



# Matrix Multiplication

```
class Worker extends Thread {  
    int row, col;  
    Worker(int row, int col) {  
        this.row = row; this.col = col;  
    }  
    public void run() {  
        double dotProduct = 0.0;  
        for (int i = 0; i < n; i++)  
            dotProduct += a[row][i] * b[i][col];  
        c[row][col] = dotProduct;  
    }  
}
```

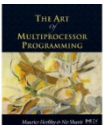
**a thread**



# Matrix Multiplication

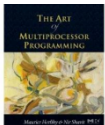
```
class Worker extends Thread {  
    int row, col;  
    Worker(int row, int col) {  
        this.row = row; this.col = col;  
    }  
    public void run() {  
        double dotProduct = 0.0;  
        for (int i = 0; i < n; i++)  
            dotProduct += a[i][row] * b[i][col];  
        c[row][col] = dotProduct;  
    }  
}
```

**Which matrix entry to compute**



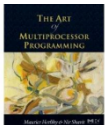
# Matrix Multiplication

```
class Worker extends Thread {  
    int row, col;  
    Worker(int row, int col) {  
        this.row = row; this.col = col;  
    }  
    public void run() { Actual computation  
        double dotProduct = 0.0;  
        for (int i = 0; i < n; i++)  
            dotProduct += a[row][i] * b[i][col];  
        c[row][col] = dotProduct;  
    }  
}
```



# Matrix Multiplication

```
void multiply() {  
    Worker[][] worker = new Worker[n][n];  
    for (int row ...)  
        for (int col ...)  
            worker[row][col] = new Worker(row,col);  
    for (int row ...)  
        for (int col ...)  
            worker[row][col].start();  
    for (int row ...)  
        for (int col ...)  
            worker[row][col].join();  
}
```

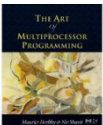




# Matrix Multiplication

```
void multiply() {  
    Worker[][] worker = new Worker[n][n];  
    for (int row ...)   
        for (int col ...)   
            worker[row][col] = new Worker(row,col);  
    for (int row ...)   
        for (int col ...)   
            worker[row][col].start();  
    for (int row ...)   
        for (int col ...)   
            worker[row][col].join();  
}
```

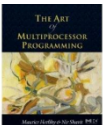
**Create n x n  
threads**



# Matrix Multiplication

```
void multiply() {  
    Worker[][] worker = new Worker[n][n];  
    for (int row ...)  
        for (int col ...)  
            worker[row][col] = new Worker(row,col);  
    for (int row ...)  
        for (int col ...)  
            worker[row][col].start();  
    for (int row ...)  
        for (int col ...)  
            worker[row][col].join();  
}
```

**Start them**




# Matrix Multiplication

```
void multiply() {  
    Worker[][] worker = new Worker[n][n];  
    for (int row ...)  
        for (int col ...)  
            worker[row][col] = new Worker(row,col);  
    for (int row ...)  
        for (int col ...)  
            worker[row][col].start();  
    for (int row ...)  
        for (int col ...)  
            worker[row][col].join();  
}
```

**Start them**

**Wait for  
them to  
finish**




# Matrix Multiplication

```
void multiply() {  
    Worker[][] worker = new Worker[n][n];  
    for (int row ...)  
        for (int col ...)  
            worker[row][col] = new Worker(row,col);  
    for (int row ...)  
        for (int col ...)  
            worker[row][col].start();  
    for (int row ...)  
        for (int col ...)  
            worker[row][col].join();  
}
```

Start them

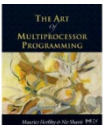
What's wrong with this picture?

Wait for  
them to  
finish




# Thread Overhead

- Threads Require resources
  - Memory for stacks
  - Setup, teardown
  - Scheduler overhead
- Short-lived threads
  - Ratio of work versus overhead bad



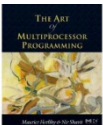
# Thread Pools

- More sensible to keep a pool of
  - long-lived threads
- Threads assigned short-lived tasks
  - Run the task
  - Rejoin pool
  - Wait for next assignment



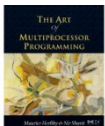
# Thread Pool = Abstraction

- Insulate programmer from platform
  - Big machine, big pool
  - Small machine, small pool
- Portable code
  - Works across platforms
  - Worry about algorithm, not platform



# ExecutorService Interface

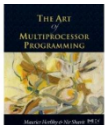
- In `java.util.concurrent`
  - Task = **Runnable** object
    - If no result value expected
    - Calls **run()** method.
  - Task = **Callable<T>** object
    - If result value of type **T** expected
    - Calls **T call()** method.





# Future<T>

```
Callable<T> task = ...;  
...  
Future<T> future = executor.submit(task);  
...  
T value = future.get();
```



# Future<T>

```
Callable<T> task = ...;
```


```
...
```

```
Future<T> future = executor.submit(task);
```

```
...
```

```
T value = future.get();
```

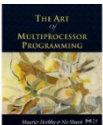
Submitting a **Callable<T>** task  
returns a **Future<T>** object



# Future<T>

```
Callable<T> task = ...;  
...  
Future<T> future = executor.submit(task);  
...  
T value = future.get();
```

The Future's **get()** method blocks  
until the value is available



# Future<?>


```
Runnable task = ...;
```

```
...
```

```
Future<?> future = executor.submit(task);
```

```
...
```

```
future.get();
```



# Future<?>

```
Runnable task = ...;
```


```
...
```

```
Future<?> future = executor.submit(task);
```

```
...
```

```
future.get();
```

Submitting a **Runnable** task  
returns a **Future<?>** object



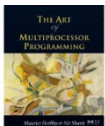
# Future<?>

```
Runnable task = ...;  
...  
Future<?> future = executor.submit(task);  
...  
future.get();
```

The Future's **get()** method blocks until the computation is complete

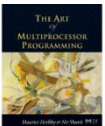
# Note

- Executor Service submissions
  - Like New England traffic signs
  - Are purely advisory in nature
- The executor
  - Like the Boston driver
  - Is free to ignore any such advice
  - And could execute tasks sequentially ...



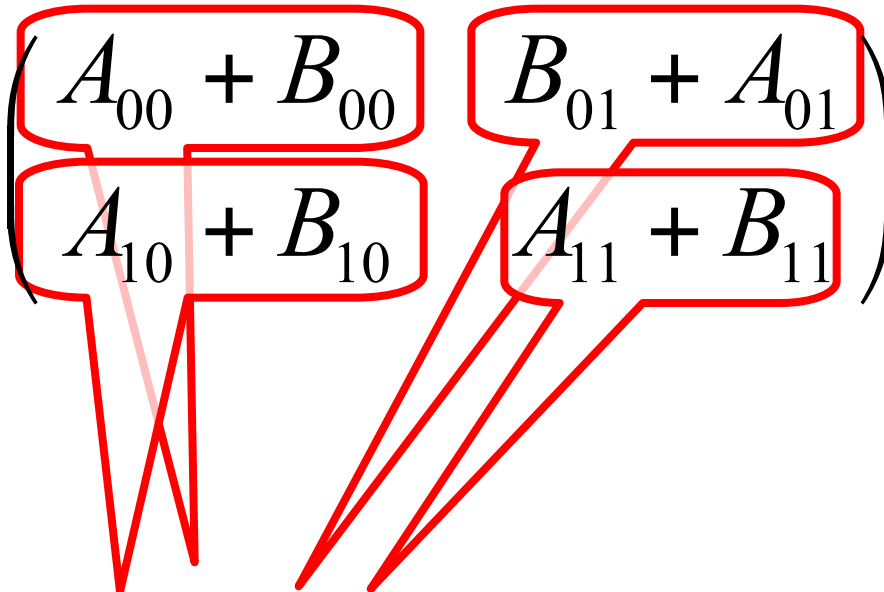
# Matrix Addition

$$\begin{pmatrix} C_{00} & C_{01} \\ C_{10} & C_{11} \end{pmatrix} = \begin{pmatrix} A_{00} + B_{00} & A_{01} + B_{01} \\ A_{10} + B_{10} & A_{11} + B_{11} \end{pmatrix}$$





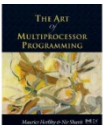
# Matrix Addition

$$\begin{pmatrix} C_{00} & C_{01} \\ C_{10} & C_{11} \end{pmatrix} = \begin{pmatrix} A_{00} + B_{00} & B_{01} + A_{01} \\ A_{10} + B_{10} & A_{11} + B_{11} \end{pmatrix}$$


4 parallel additions

# Matrix Addition Task

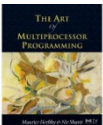
```
class AddTask implements Runnable {  
    Matrix a, b; // add this!  
    public void run() {  
        if (a.dim == 1) {  
            c[0][0] = a[0][0] + b[0][0]; // base case  
        } else {  
            (partition a, b into half-size matrices aij and bij)  
            Future<?> f00 = exec.submit(add(a00, b00));  
  
            ...  
            Future<?> f11 = exec.submit(add(a11, b11));  
            f00.get(); ...; f11.get();  
  
            ...  
        }  
    }  
}
```



# Matrix Addition Task

```
class AddTask implements Runnable {  
    Matrix a, b; // add this!  
    public void run() {  
        if (a.dim == 1) {  
            c[0][0] = a[0][0] + b[0][0]; // base case  
        } else {  
            (partition a, b into half-size matrices aij and bij)  
            Future<?> f00 = exec.submit(add(a00, b00));  
            ...  
            Future<?> f11 = exec.submit(add(a11, b11));  
            f00.get(); ...; f11.get();  
            ...  
        }  
    }  
}
```

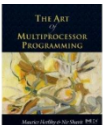
**Constant-time operation**



# Matrix Addition Task

```
class AddTask implements Runnable {  
    Matrix a, b; // add this!  
    public void run() {  
        if (a.dim == 1) {  
            c[0][0] = a[0][0] + b[0][0]; // base case  
        } else {  
            (partition a, b into half-size matrices aij and bij)  
            Future<?> f00 = exec.submit(add(a00, b00));  
            ...  
            Future<?> f11 = exec.submit(add(a11, b11));  
            f00.get(); ...; f11.get();  
            ...  
        }  
    }  
}
```

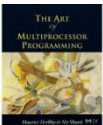
**Submit 4 tasks**



# Matrix Addition Task

```
class AddTask implements Runnable {  
    Matrix a, b; // add this!  
    public void run() {  
        if (a.dim == 1) {  
            c[0][0] = a[0][0] + b[0][0]; // base case  
        } else {  
            (partition a, b into half-size matrices aij and bij)  
            Future<?> f00 = exec.submit(add(a00, b00));  
            ...  
            Future<?> f11 = exec.submit(add(a11, b11));  
            f00.get(); ...; f11.get();  
            ...  
        }  
    }  
}
```

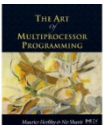
**Base case: add directly**



# Matrix Addition Task

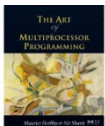
```
class AddTask implements Runnable {  
    Matrix a, b; // multiply this!  
    public void run() {  
        if (a.dim == 1) {  
            c[0][0] = a[0][0] + b[0][0]; // base case  
        } else {  
            (partition a, b into half-size matrices aij and bij)  
            Future<?> f00 = exec.submit(add(a00, b00));  
  
            ...  
            Future<?> f11 = exec.submit(add(a11, b11));  
            f00.get(); ...; f11.get();  
  
            ...  
        }  
    }  
}
```

**Let them finish**

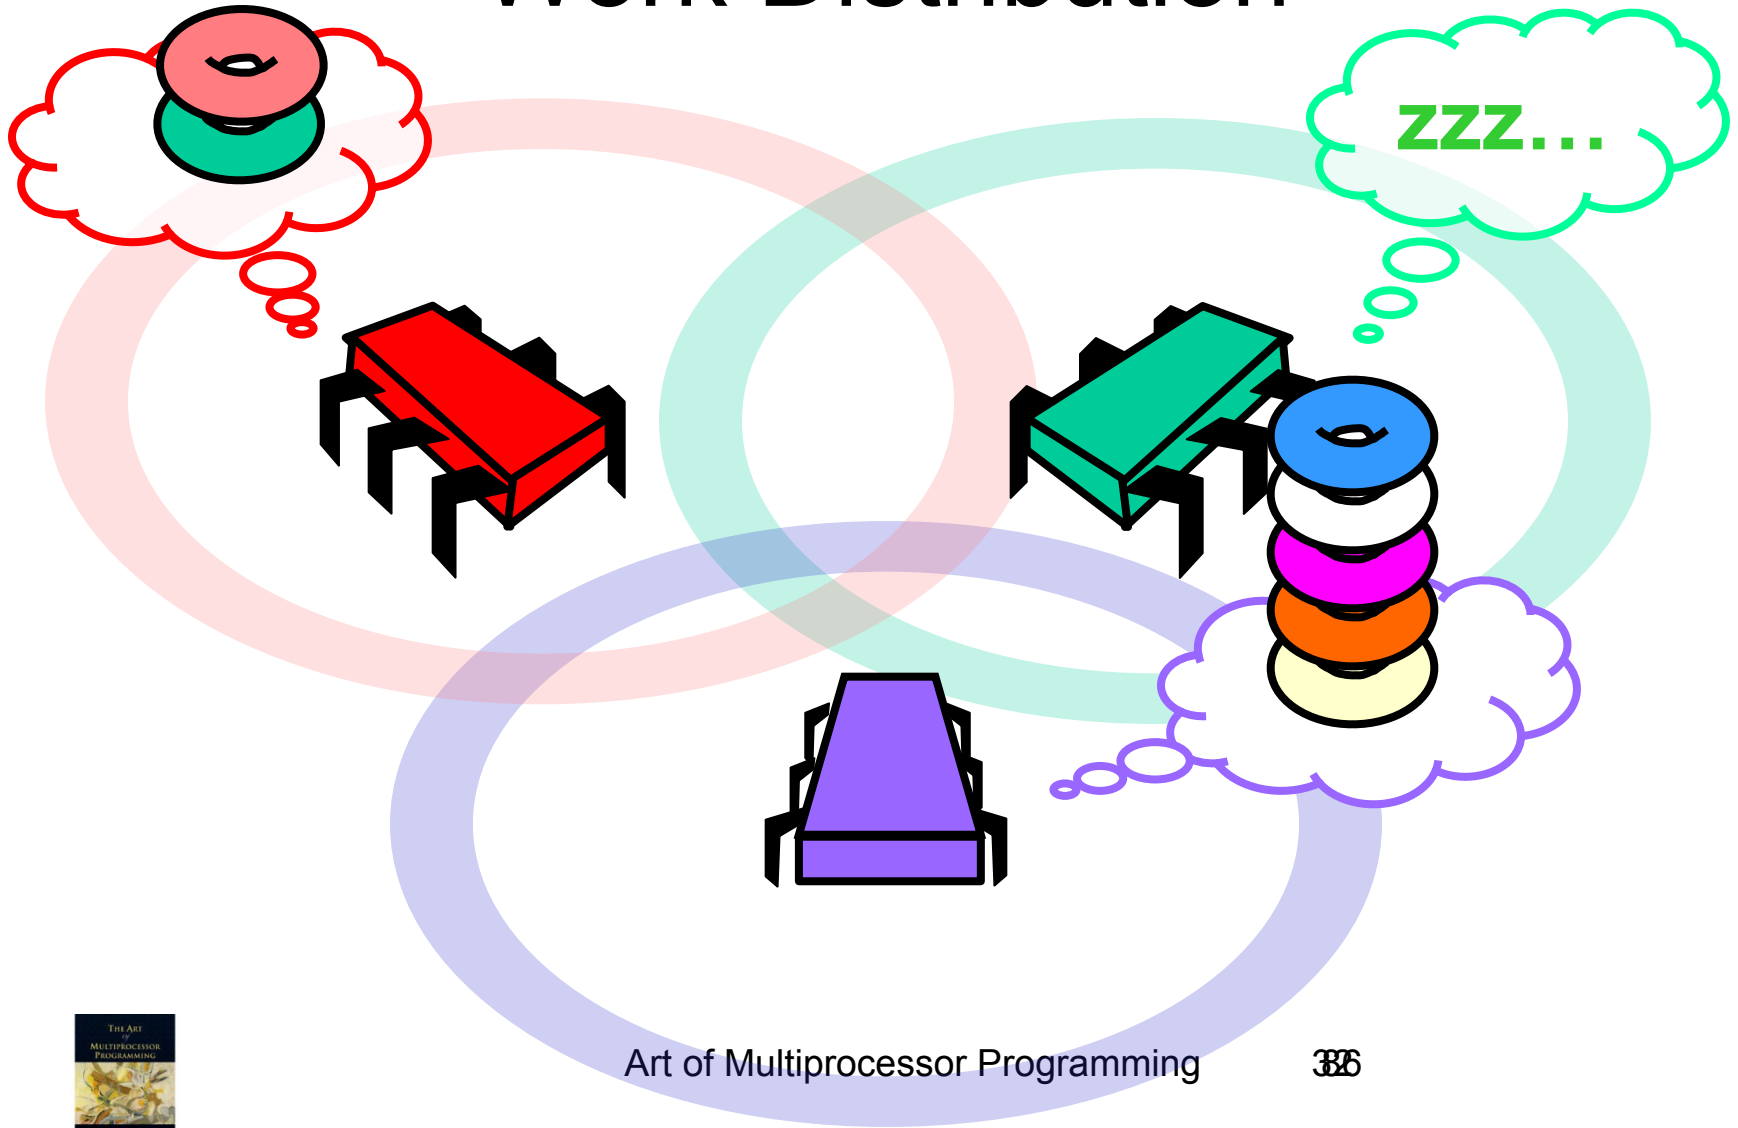


# Dependencies

- Matrix example is not typical
- Tasks are independent
  - Don't need results of one task ...
  - To complete another
- Often tasks are not independent
  - We can use other ways to handle these types of tasks
  - For example - pipelining, barriers, etc.

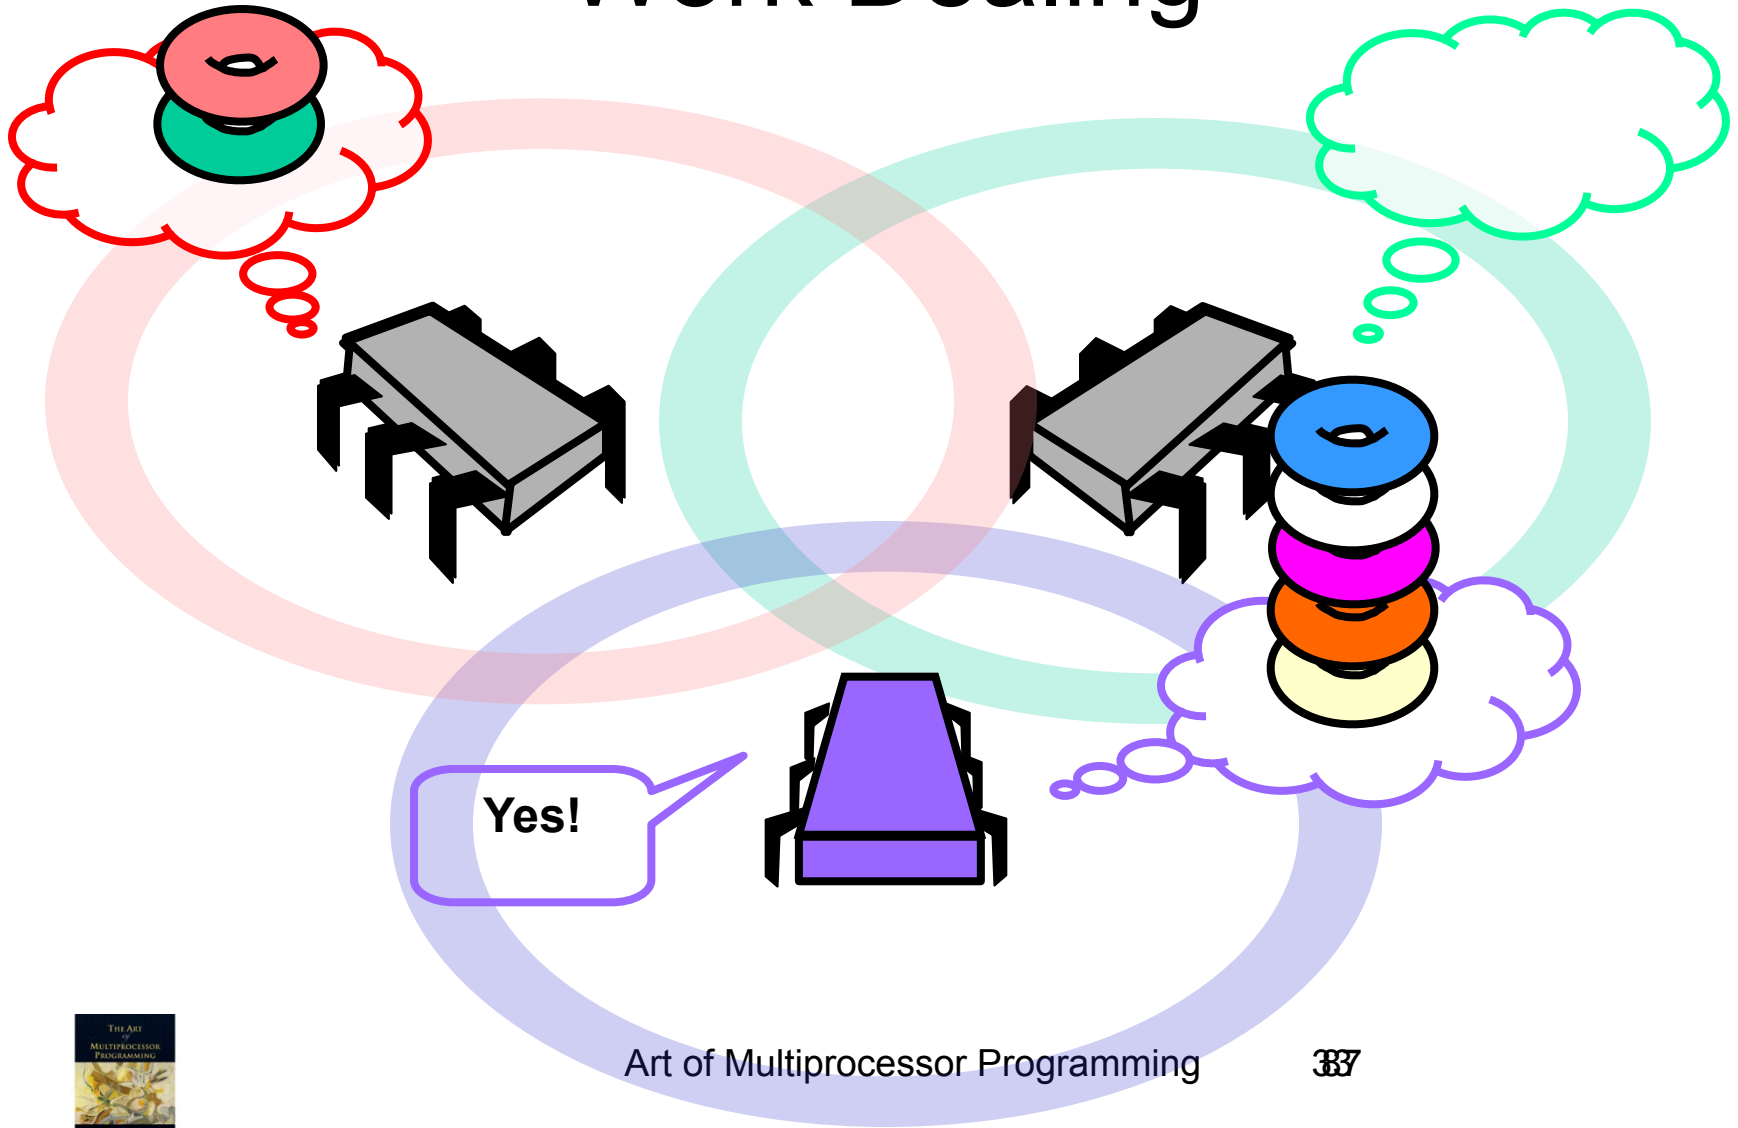


# Work Distribution


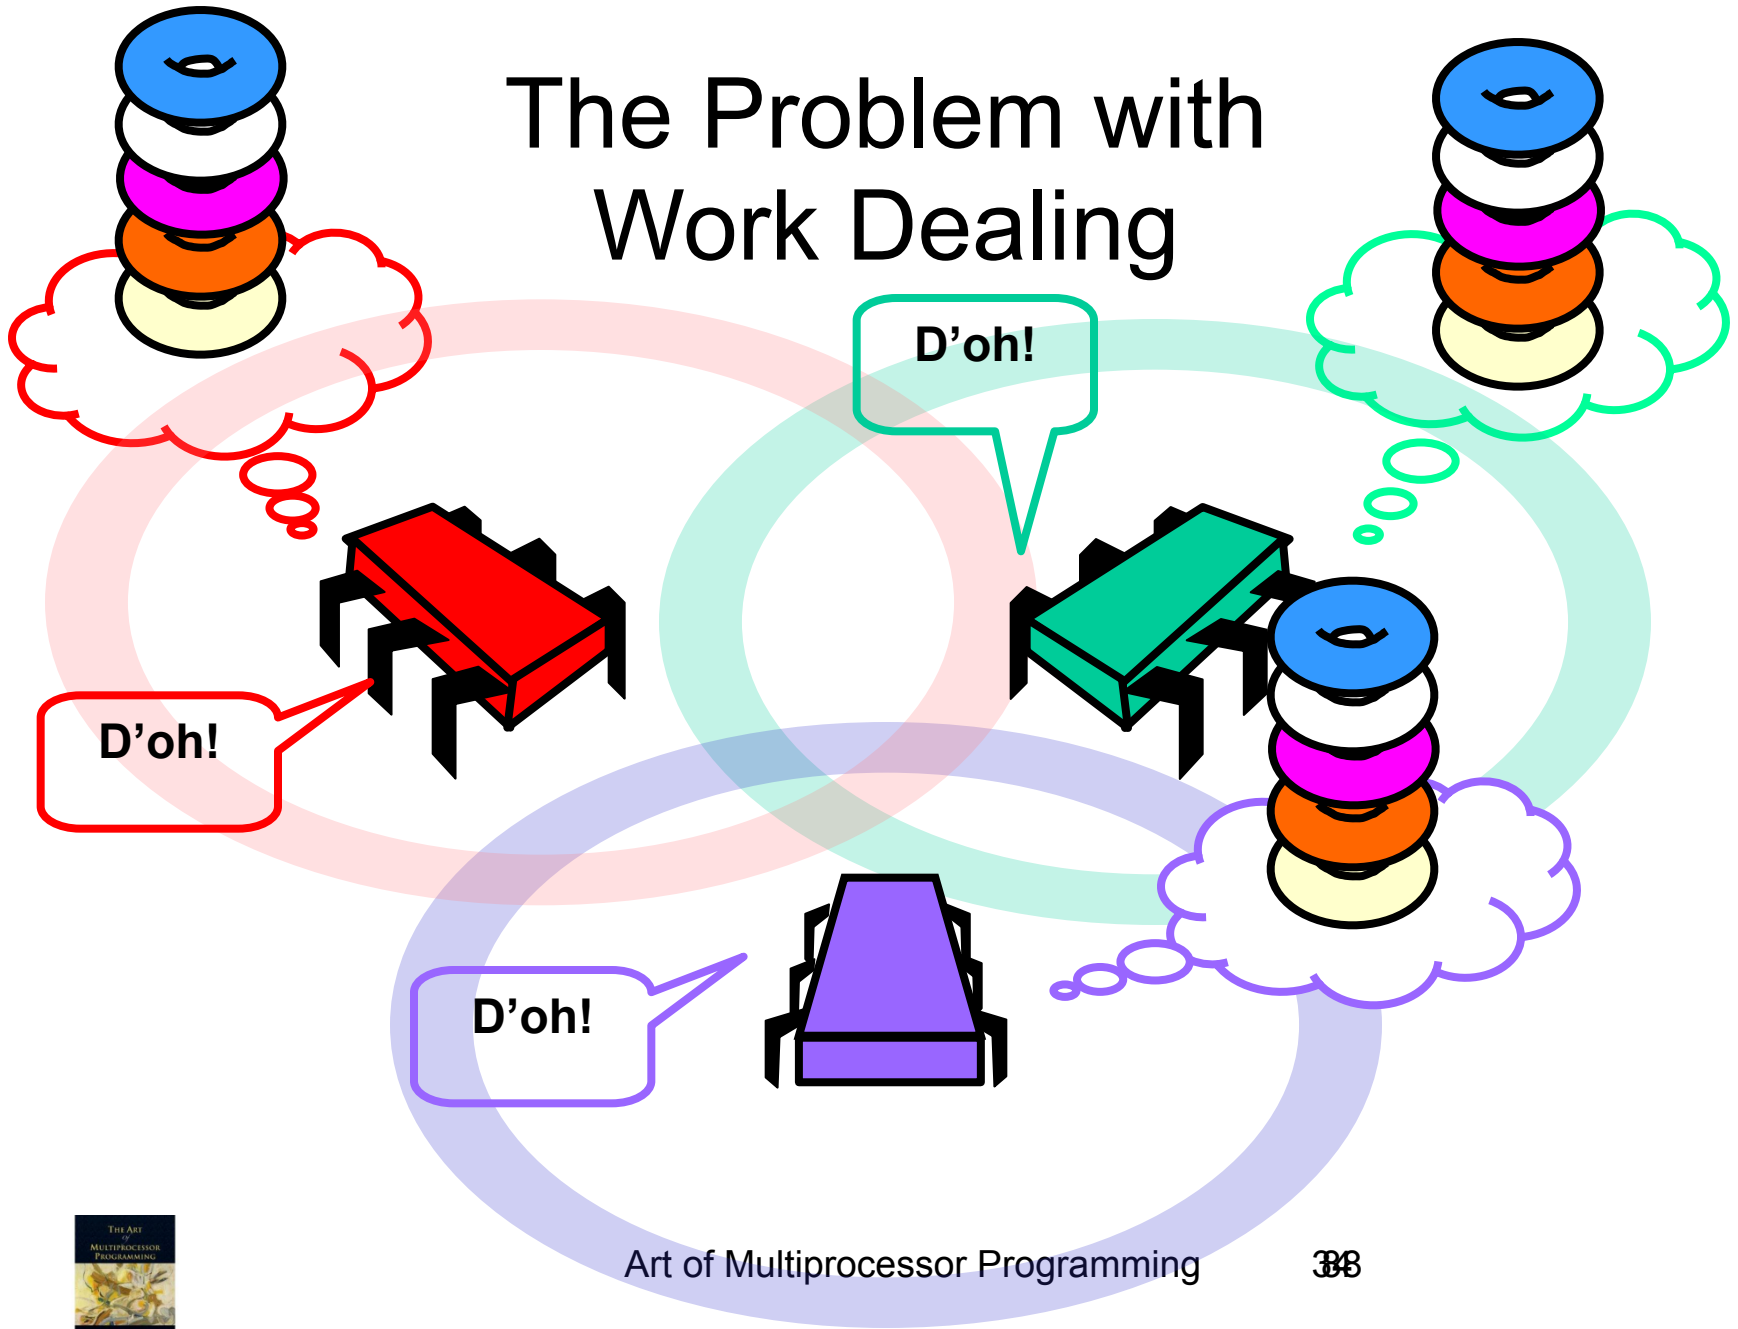




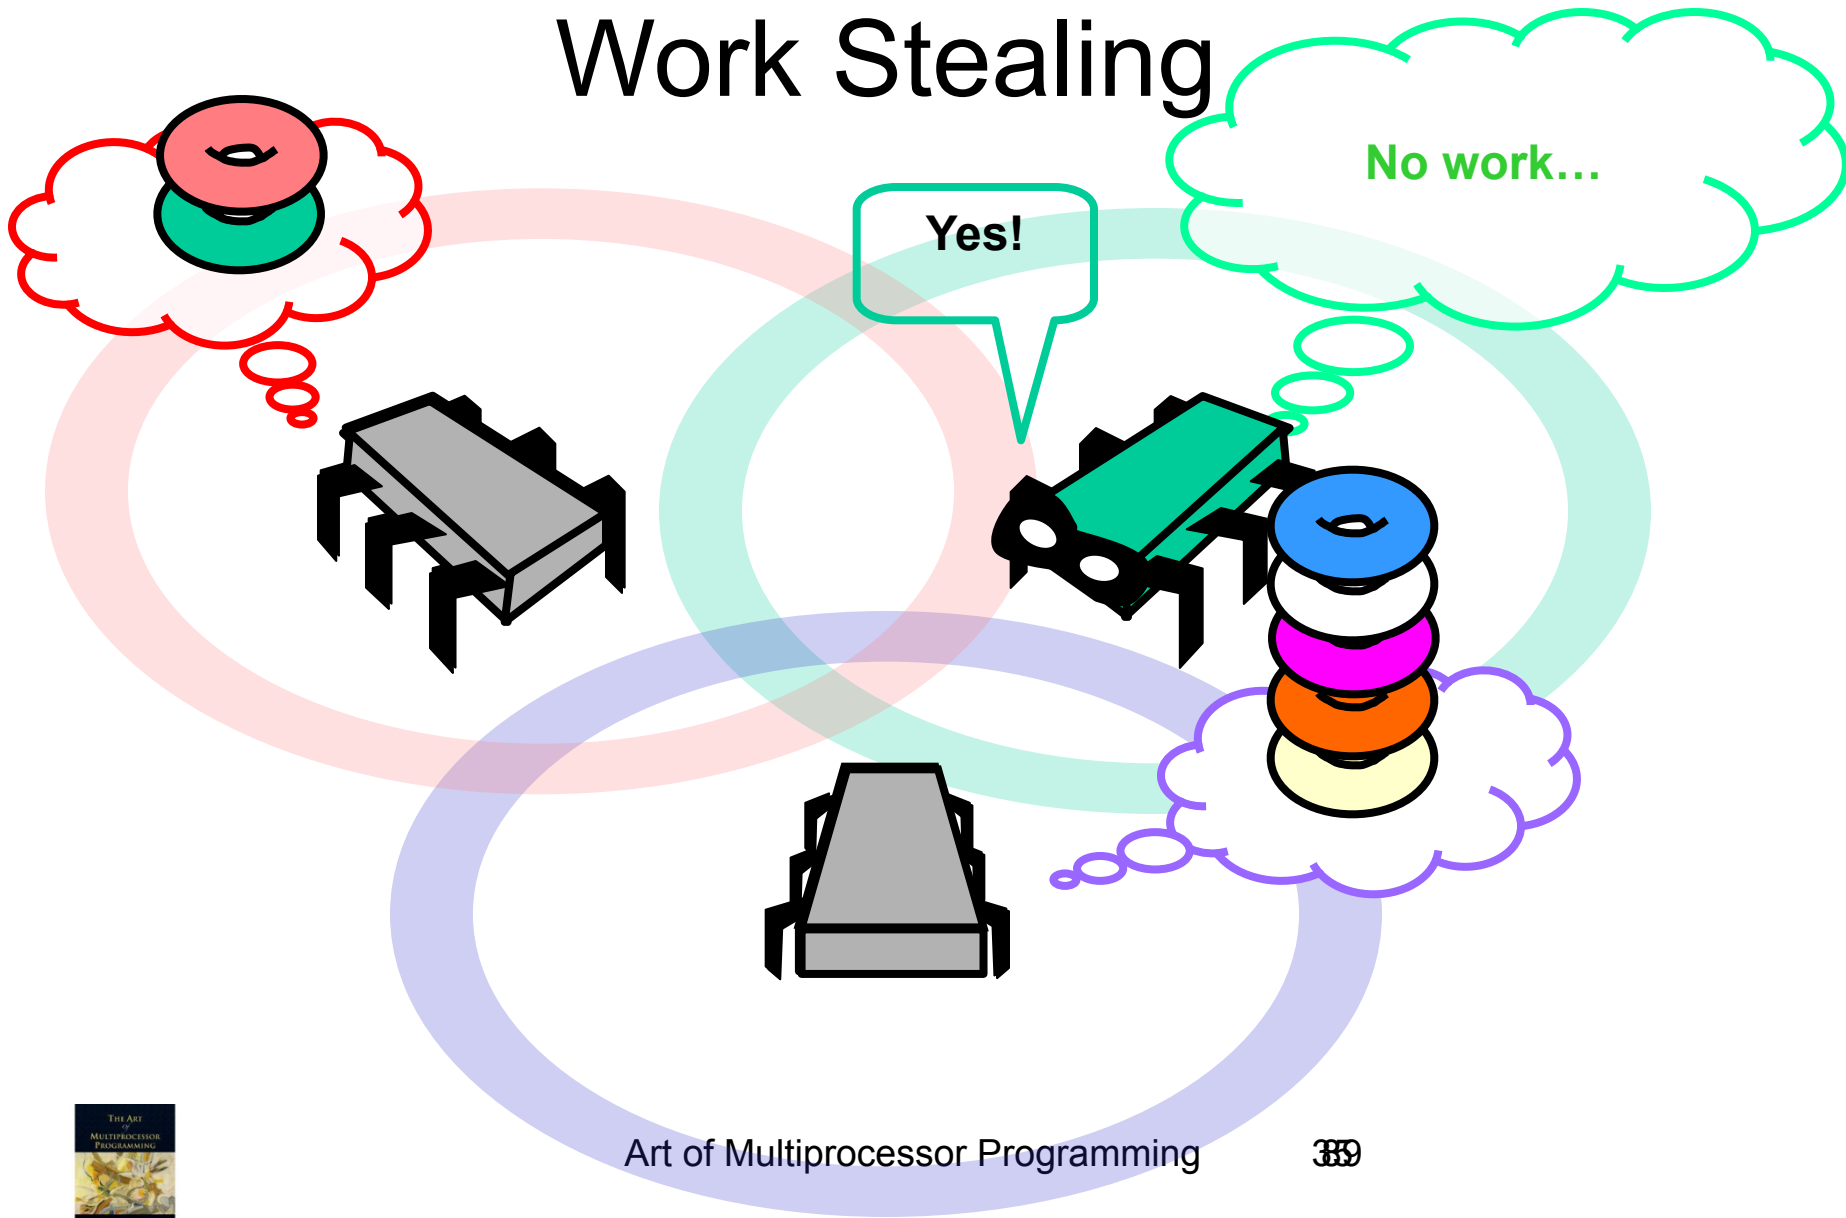
# Work Dealing



# The Problem with Work Dealing

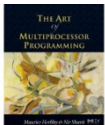


# Work Stealing




# Lock-Free Work Stealing

- Each thread has a pool of ready work
- Remove work without synchronizing
- If you run out of work, steal someone else's
- Choose victim at random




# Local Work Pools


Each work pool is a Double-Ended Queue



# Work DEQueue<sup>1</sup>





1. Double-Ended Queue



# Obtain Work

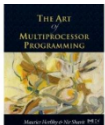

- Obtain work
- Run task until
- Blocks or terminates



popBottom


# New Work

- Unblock node
- Spawn node






# Whatcha Gonna do When the Well Runs Dry?




# Steal Work from Others




Pick random thread's DEQueue



# Steal this Task!

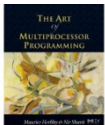


popTop



# Task DEQueue

- Methods
    - pushBottom
    - popBottom
    - popTop
- } Never happen concurrently




# Task DEQueue

- Methods

- pushBottom
- popBottom
- popTop



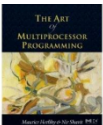
Most common –  
make them fast  
(minimize use of  
CAS)



# Ideal

- Wait-Free
- Linearizable
- Constant time

Fortune Cookie: “It is better to be young, rich and beautiful,  
than old, poor, and ugly”

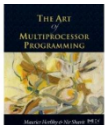


# Compromise

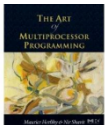

- Method popTop may fail if
  - Concurrent popTop succeeds, or a
  - Concurrent popBottom takes last task



**Blame the victim!**




# Dreaded A Problem

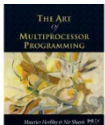






# Dreaded A Problem




# Dreaded A Problem




# Dreaded A Problem




# Dreaded A Problem




# Dreaded A Problem




# Dreaded A Problem



# Dreaded A/P Problem



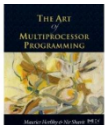
# Fix for Dreaded ABA





# Bounded DEQueue

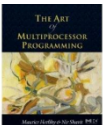
```
public class BDEQueue {  
    AtomicStampedReference<Integer> top;  
    volatile int bottom;  
    Runnable[] tasks;  
    ...  
}
```



# Bounded DEQueue

```
public class BDEQueue {  
    AtomicStampedReference<Integer> top;  
    volatile int bottom;  
    Runnable[] tasks;  
    ...  
}
```


**Index & Stamp  
(synchronized)**



# Bounded DEQueue

```
public class BDEQueue {  
    AtomicStampedReference<Integer> top;  
    volatile int bottom;  
    Runnable[] deq;  
    ...  
}
```


**index of bottom task  
no need to synchronize  
memory barrier needed**



# Bounded DEQueue

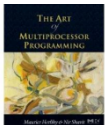
```
public class BDEQueue {  
    AtomicStampedReference<Integer> top;  
    volatile int bottom;  
    Runnable[] tasks;  
    ...  
}
```

**Array holding tasks**



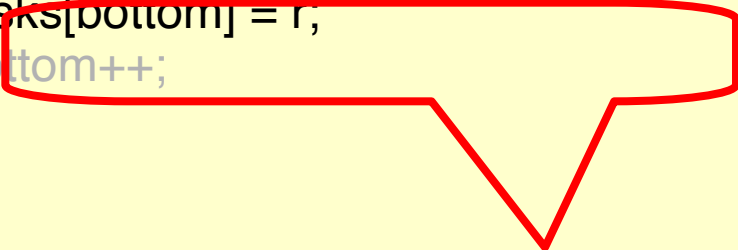
# pushBottom()

```
public class BDEQueue {  
    ...  
    void pushBottom(Runnable r){  
        tasks[bottom] = r;  
        bottom++;  
    }  
    ...  
}
```



# pushBottom()


```
public class BDEQueue {  
    ...  
    void pushBottom(Runnable r){  
        tasks[bottom] = r;  
        bottom++;  
    }  
    ...  
}
```



**Bottom is the index to store  
the new task in the array**

# pushBottom()

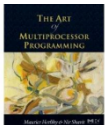
```
public class BDEQueue {  
    ...  
    void pushBottom(Runnable r){  
        tasks[bottom] = r;  
        bottom++;  
    }  
    ...  
}
```



**Adjust the bottom index**

# Steal Work

```
public Runnable popTop() {  
    int[] stamp = new int[1];  
    int oldTop = top.get(stamp), newTop = oldTop + 1;  
    int oldStamp = stamp[0], newStamp = oldStamp + 1;  
    if (bottom <= oldTop)  
        return null;  
    Runnable r = tasks[oldTop];  
    if (top.CAS(oldTop, newTop, oldStamp, newStamp)) return r;  
    return null;  
}
```

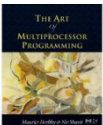




# Steal Work

```
public Runnable popTop() {  
    int[] stamp = new int[1];  
    int oldTop = top.get(stamp), newTop = oldTop + 1;  
    int oldStamp = stamp[0], newStamp = oldStamp + 1;  
    if (bottom <= oldTop)  
        return null;  
    Runnable r = tasks[oldTop];  
    if (top.CAS(oldTop, newTop, oldStamp, newStamp)) return r;  
    return null;  
}
```

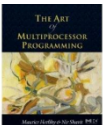
**Read top (value & stamp)**



# Steal Work


```
public Runnable popTop() {  
    int[] stamp = new int[1];  
    int oldTop = top.get(stamp), newTop = oldTop + 1;  
    int oldStamp = stamp[0], newStamp = oldStamp + 1;  
    if (bottom <= oldTop)  
        return null;  
    Runnable r = tasks[oldTop];  
    if (top.CAS(oldTop, newTop, oldStamp, newStamp)) return r;  
    return null;  
}
```

**Compute new value & stamp**



# Steal Work


```
public Runnable popTop() {  
    int[] stamp = new int[1];  
    int oldTop = top.get(stamp), newTop = oldTop + 1;  
    int oldStamp = stamp[0], newStamp = oldStamp + 1;  
    if (bottom <= oldTop)  
        return null;  
    Runnable r = tasks[oldTop];  
    if (top.CAS(oldTop, newTop, oldStamp, newStamp)) return r;  
    return null;  
}
```



**Quit if queue is empty**

# Steal Work

```
public Runnable popTop() {  
    int[] stamp = new int[1];  
    int oldTop = top.get(stamp), newTop = oldTop + 1;  
    int oldStamp = stamp[0], newStamp = oldStamp + 1;  
    if (bottom <= oldTop)  
        return null;  
    Runnable r = tasks[oldTop];  
    if (top.CAS(oldTop, newTop, oldStamp, newStamp)) return r;  
    return null;  
}
```

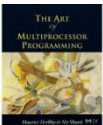


**Try to steal the task**

# Steal Work


```
public Runnable popTop() {  
    int[] stamp = new int[1];  
    int oldTop = top.get(stamp), newTop = oldTop + 1;  
    int oldStamp = stamp[0], newStamp = oldStamp + 1;  
    if (bottom <= oldTop)  
        return null;  
    Runnable r = tasks[oldTop];  
    if (top.CAS(oldTop, newTop, oldStamp, newStamp)) return r;  
    return null;  
}
```

**Give up if  
conflict occurs**



# Take Work


```
Runnable popBottom() {  
    if (bottom == 0) return null;  
    bottom--;  
    Runnable r = tasks[bottom];  
    int[] stamp = new int[1];  
    int oldTop = top.get(stamp), newTop = 0;  
    int oldStamp = stamp[0], newStamp = oldStamp + 1;  
    if (bottom > oldTop) return r;  
    if (bottom == oldTop){  
        bottom = 0;  
        if (top.CAS(oldTop, newTop, oldStamp, newStamp))  
            return r;  
    }  
    top.set(newTop,newStamp); return null;  
    bottom = 0; }
```



# Take Work

```
Runnable popBottom() {  
    f (bottom == 0) return null;  
    bottom--;  
    Runnable r = tasks[bottom];  
    int[] stamp = new int[1];  
    int oldTop = top.get(stamp), newTop = 0;  
    int oldStamp = stamp[0], newStamp = oldStamp + 1;  
    if (bottom > oldTop) return r;  
    if (bottom == oldTop){  
        bottom = 0;  
        if (top.CAS(oldTop, newTop, oldStamp, newStamp))  
            return r;  
    }  
    top.set(newTop, newStamp); return null;  
    bottom = 0; }
```


**Make sure queue is non-empty**



# Take Work

```
Runnable popBottom() {  
    if (bottom == 0) return null;  
    bottom--;  
    Runnable r = tasks[bottom];  
    int[] stamp = new int[1];  
    int oldTop = top.get(stamp), newTop = 0;  
    int oldStamp = stamp[0], newStamp = oldStamp + 1;  
    if (bottom > oldTop) return r;  
    if (bottom == oldTop){  
        bottom = 0;  
        if (top.CAS(oldTop, newTop, oldStamp, newStamp))  
            return r;  
    }  
    top.set(newTop, newStamp); return null;  
    bottom = 0; }  
}
```

**Prepare to grab bottom task**






# Take Work


```
Runnable popBottom() {  
    if (bottom == 0) return null;  
    bottom--;  
    Runnable r = tasks[bottom]:  
    int[] stamp = new int[1];  
    int oldTop = top.get(stamp), newTop = 0;  
    int oldStamp = stamp[0], newStamp = oldStamp + 1;  
    if (bottom > oldTop) return r;  
    if (bottom == oldTop){  
        bottom = 0;  
        if (top.CAS(oldTop, newTop, oldStamp, newStamp))  
            return r;  
    }  
    top.set(newTop, newStamp); return null;  
    bottom = 0; }
```

**Read top, & prepare new values**



# Take Work


```
Runnable popBottom() {  
    if (bottom == 0) return null;  
    bottom--;  
    Runnable r = tasks[bottom];  
    int[] stamp = new int[1];  
    int oldTop = top.get(stamp), newTop = 0;  
    int oldStamp = stamp[0], newStamp = oldStamp + 1;  
    if (bottom > oldTop) return r;  
    if (bottom == oldTop){  
        bottom = 0;  
        if (top.CAS(oldTop, newTop, oldStamp, newStamp))  
            return r;  
    }  
    top.set(newTop, newStamp);  
    bottom = 0;  
}
```



**If top & bottom one or more apart,  
no conflict**

# Take Work

```
Runnable popBottom() {  
    if (bottom == 0) return null;  
    bottom--;  
    Runnable r = tasks[bottom];  
    int[] stamp = new int[1];  
    int oldTop = top.get(stamp), newTop = 0;  
    int oldStamp = stamp[0], newStamp = oldStamp + 1;  
    if (bottom > oldTop) return r;  
    if (bottom == oldTop){  
        bottom = 0;  
        if (top.CAS(oldTop, newTop, oldStamp, newStamp))  
            return r;  
    }  
    top.set(newTop, newStamp); return null;  
    bottom = 0;}
```



**At most one item left**

# Take Work


**Try to steal last task.**

**Reset bottom because the  
DEQueue will be empty  
even if unsuccessful (why?)**

```
Runnable popBottom() {  
    if (bottom == 0) return null;  
    Runnable r = tasks[bottom];  
    int[] stamp = new int[1];  
    int oldStamp = stamp[0], newStamp = oldStamp + 1;  
    if (bottom > oldTop) return r;  
    if (bottom == oldTop){  
        bottom = 0;  
        if (top.CAS(oldTop, newTop, oldStamp, newStamp))  
            return r;  
    }  
    top.set(newTop, newStamp); return null;  
    bottom = 0;}
```

# Take Work


```
Runnable popBottom() {  
    I win CAS  
    if (bottom == 0) return null;  
    bottom--;  
    Runnable r = tasks[bottom];  
    int[] stamp = new int[1];  
    int oldTop = top.get(stamp), newTop = 0;  
    int oldStamp = stamp[0], newStamp = oldStamp + 1;  
    if (bottom > oldTop) return r;  
    if (bottom == oldTop){  
        bottom = 0;  
        if (top.CAS(oldTop, newTop, oldStamp, newStamp))  
            return r;  
    }  
    top.set(newTop, newStamp); return null;  
    bottom = 0;}  
}
```



# Take Work

**If I lose CAS, thief must have won...**

```
Runnable r = tasks[bottom];  
int[] stamp = new int[1];  
int oldTop = top.get(stamp), newTop = 0;  
int oldStamp = stamp[0], newStamp = oldStamp + 1;  
if (bottom > oldTop) return r;  
if (bottom == oldTop){  
    bottom = 0;  
    if (top.CAS(oldTop, newTop, oldStamp, newStamp))  
        return r;  
}
```



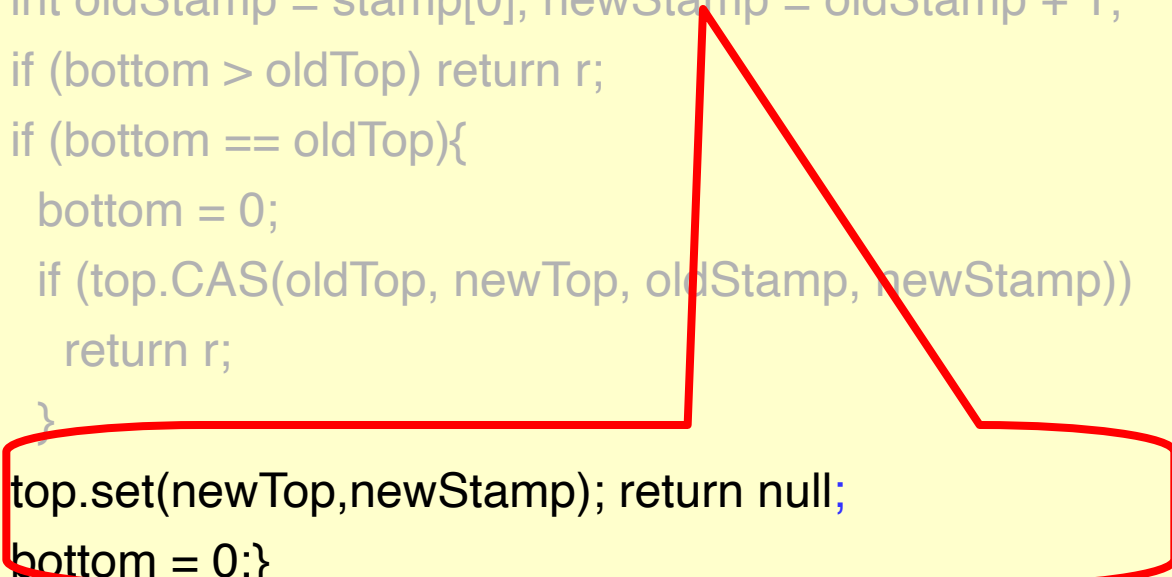
# Take Work

**Failed to get last task  
(bottom could be less than top)**

**Must still reset top and bottom  
since deque is empty**

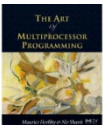
```
Failed to get last task (bottom could be less than top)
Must still reset top and bottom since deque is empty

int[] stamp = new int[r];
int oldTop = top.get(oldStamp);
int oldStamp = stamp[0], newStamp = oldStamp + 1;
if (bottom > oldTop) return r;
if (bottom == oldTop){
    bottom = 0;
    if (top.CAS(oldTop, newTop, oldStamp, newStamp))
        return r;
}
top.set(newTop, newStamp); return null;
bottom = 0;}
```



# Old English Proverb

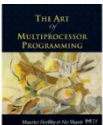
- “May as well be hanged for stealing a sheep as a goat”
- From which we conclude:
  - Stealing was punished severely
  - Sheep were worth more than goats




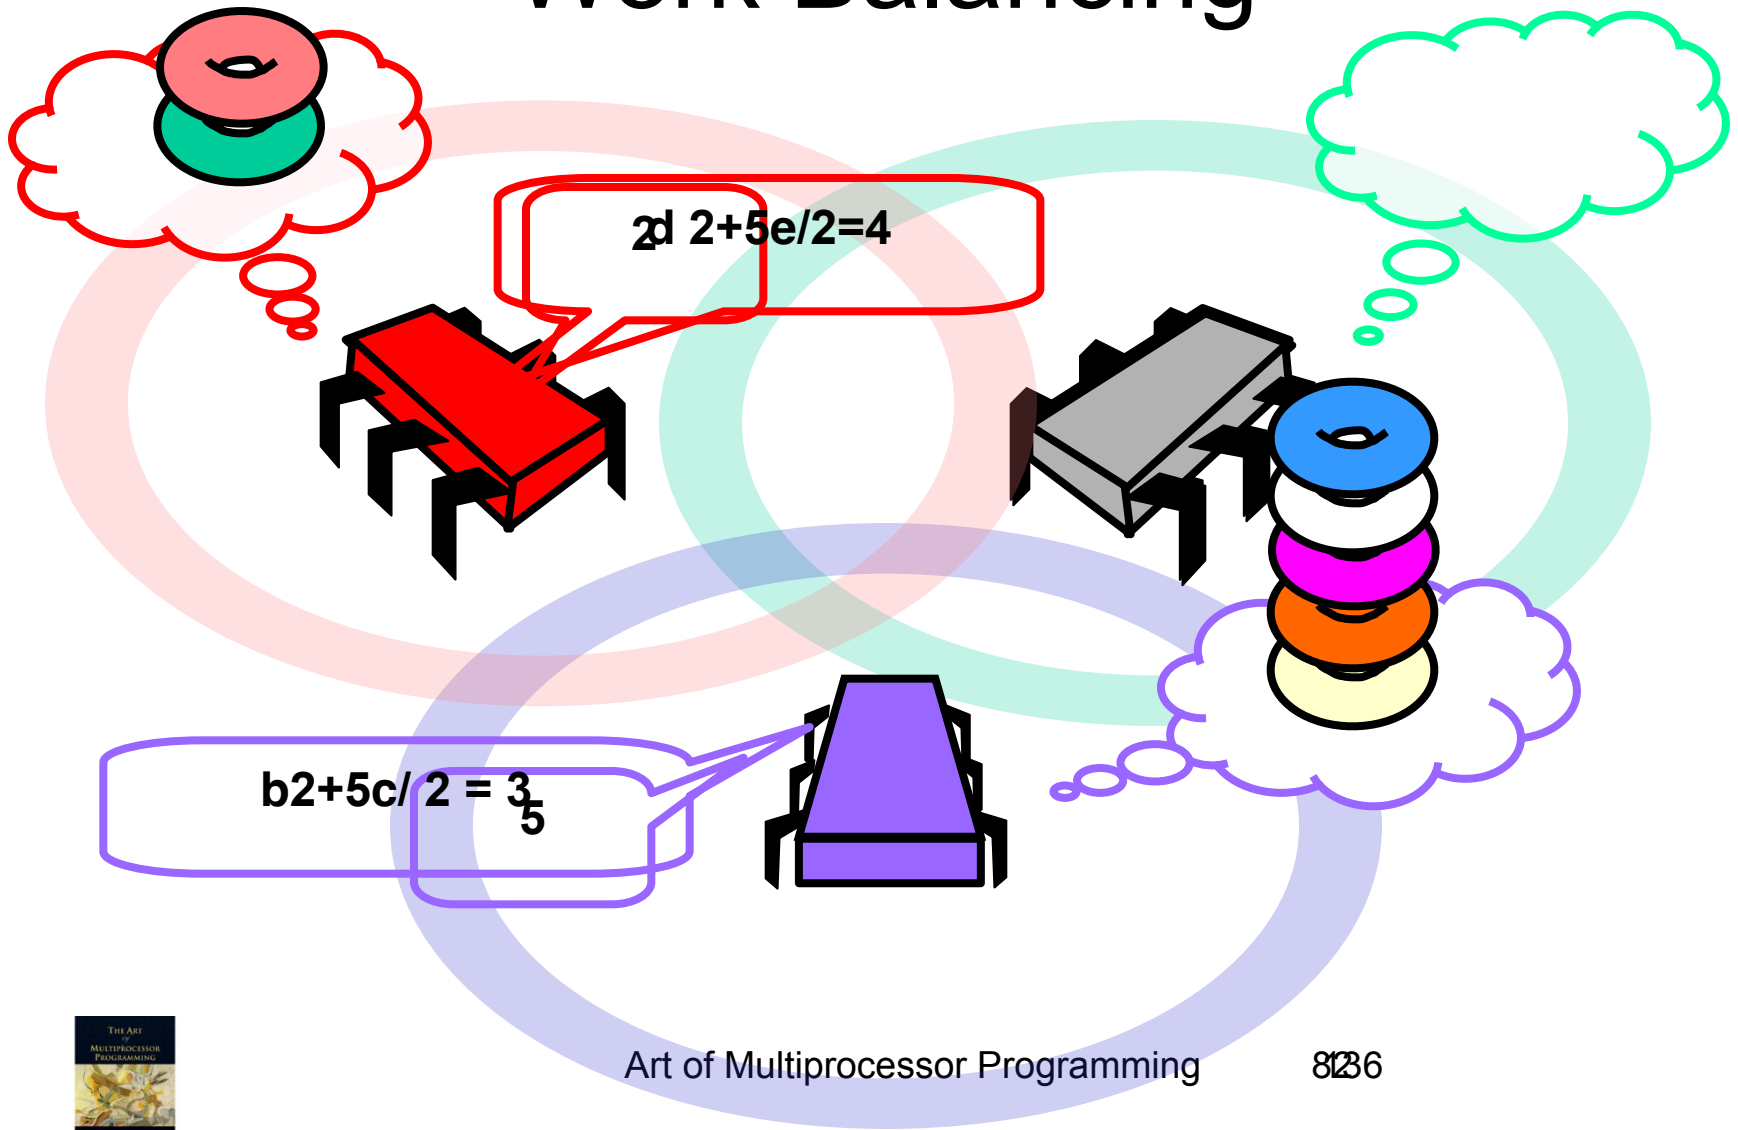


# Variations

- Stealing is expensive
  - Pay CAS
  - Only one task taken
- What if
  - Move more than one task
  - Randomly balance loads?




# Work Balancing



# Work-Balancing Thread

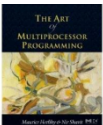
```
public void run() {  
    int me = ThreadID.get();  
    while (true) {  
        Runnable task = queue[me].deq();  
        if (task != null) task.run();  
        int size = queue[me].size();  
        if (random.nextInt(size+1) == size) {  
            int victim = random.nextInt(queue.length);  
            int min = ..., max = ...;  
            synchronized (queue[min]) {  
                synchronized (queue[max]) {  
                    balance(queue[min], queue[max]);  
                }  
            }  
        }  
    }  
}
```



# Work-Balancing Thread

```
public void run() {  
    int me = ThreadID.get();  
    while (true) {  
        Runnable task = queue[me].deq();  
        if (task != null) task.run();  
        int size = queue[me].size();  
        if (random.nextInt(size+1) == size) {  
            int victim = random.nextInt(queue.length);  
            int min = ..., max = ...;  
            synchronized (queue[min]) {  
                synchronized (queue[max]) {  
                    balance(queue[min], queue[max]);  
                }  
            }  
        }  
    }  
}
```


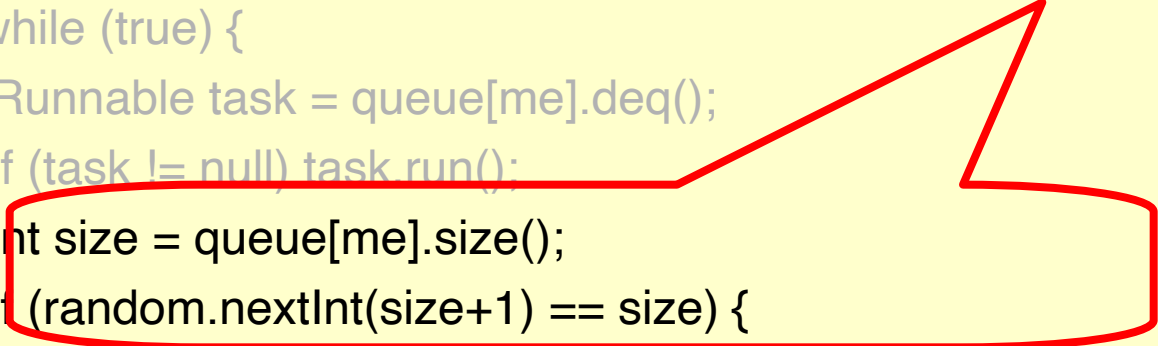
**Keep running tasks**



# Work-Balancing Thread

```
public void run() {  
    int me = ThreadID.get();  
    while (true) {  
        Runnable task = queue[me].deq();  
        if (task != null) task.run();  
        int size = queue[me].size();  
        if (random.nextInt(size+1) == size) {  
            int victim = random.nextInt(queue.length);  
            int min = ..., max = ...;  
            synchronized (queue[min]) {  
                synchronized (queue[max]) {  
                    balance(queue[min], queue[max]);  
                }  
            }  
        }  
    }  
}
```


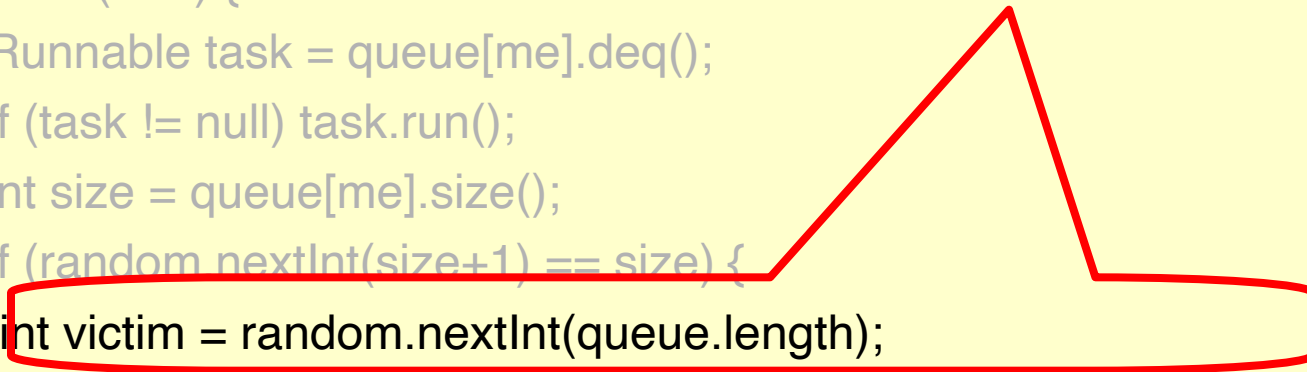
**With probability  
 $1/|queue|$**



# Work-Balancing Thread

```
public void run() {  
    int me = ThreadID.get();  
    while (true) {  
        Runnable task = queue[me].deq();  
        if (task != null) task.run();  
        int size = queue[me].size();  
        if (random.nextInt(size+1) == size) {  
            int victim = random.nextInt(queue.length);  
            int min = ..., max = ...;  
            synchronized (queue[min]) {  
                synchronized (queue[max]) {  
                    balance(queue[min], queue[max]);  
                }  
            }  
        }  
    }  
}
```


**Choose random victim**



# Work-Balancing Thread

```
public void run() {  
    int me = ThreadID.get();  
    while (true) {  
        Runnable task = queue[me].deq();  
        if (task != null) task.run();  
        int size = queue[me].size();  
        if (random.nextInt(size+1) == size) {  
            int victim = random.nextInt(queue.length);  
            int min = ..., max = ...;  
            synchronized (queue[min]) {  
                synchronized (queue[max]) {  
                    balance(queue[min], queue[max]);  
                }  
            }  
        }  
    }  
}
```


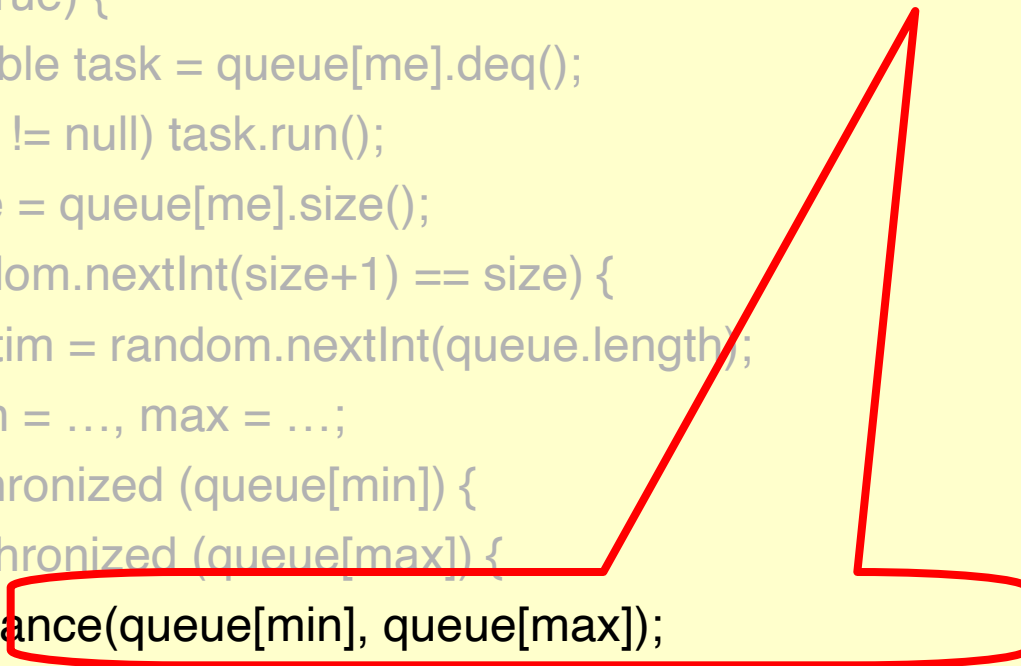
**Lock queues in canonical order**



# Work-Balancing Thread

```
public void run() {  
    int me = ThreadID.get();  
    while (true) {  
        Runnable task = queue[me].deq();  
        if (task != null) task.run();  
        int size = queue[me].size();  
        if (random.nextInt(size+1) == size) {  
            int victim = random.nextInt(queue.length);  
            int min = ..., max = ...;  
            synchronized (queue[min]) {  
                synchronized (queue[max]) {  
                    balance(queue[min], queue[max]);  
                }  
            }  
        }  
    }  
}
```

**Rebalance queues**





# Work Stealing & Balancing

- Clean separation between app & scheduling layer
- Works well when number of processors fluctuates.
- Works on “black-box” operating systems

