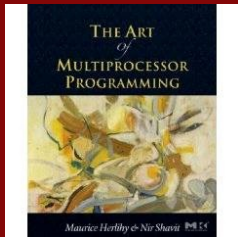
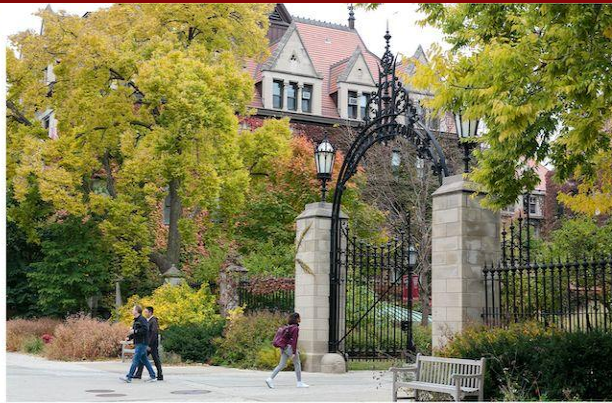


# MPCS 52060 - Parallel Programming

## M6: Advanced Parallel Patterns and Techniques (Part 2)



Original slides from “The Art of Multiprocessor Programming” by Maurice Herlihy & Nir Shavit with modifications by Lamont Samuels

# Data Parallelism

# Ever Wonder ...

When did the term *multicore* become popular?

A multi-core processor is a single computing component with two or more independent actual central processing units, which are the units that read and execute program instructions.)

wikipedia

# Let's Ask Google Ngram!




# Let's Ask Google Ngram!





# Let's Ask Google Ngram!



# Let's Ask Google Ngram!



# WordCount



alpha → 8  
bravo → 3  
charlie → 9  
...  
zulu → 1


easy to do sequentially ...

what about in parallel?




# MapReduce

split text among *mapping* threads




chapter 1



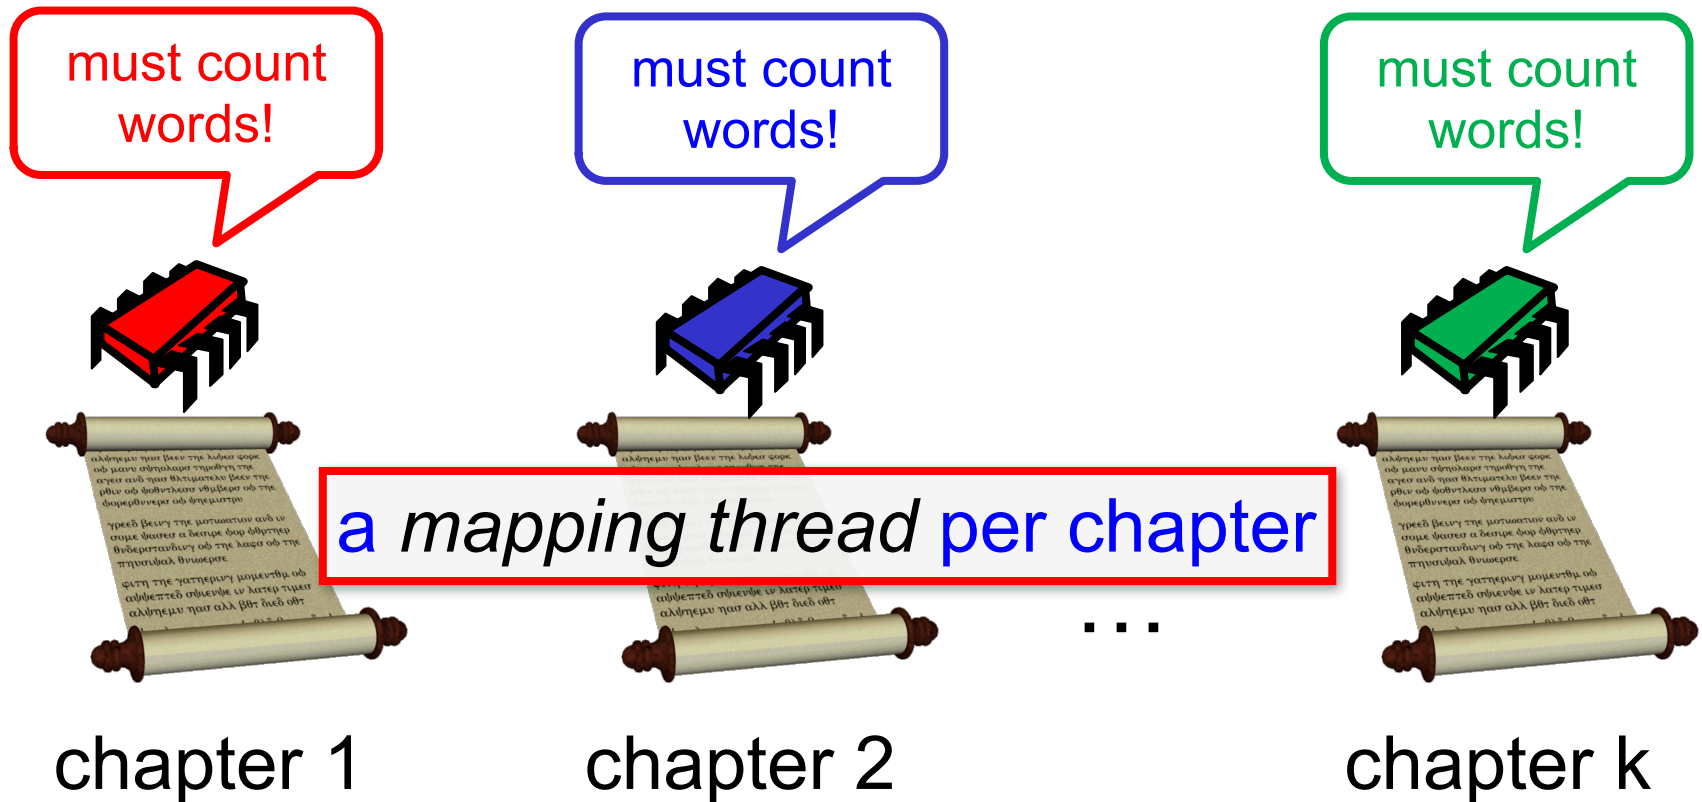
chapter 2

...

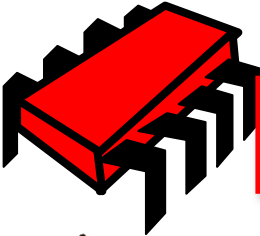


chapter k

# Map Phase



# Map Phase



alpha → 9  
juliet → 2,  
alpha → 1  
tango → 4

each mapper thread produces a *stream* ...  
of *key-value* pairs ...

key: word

value: local count



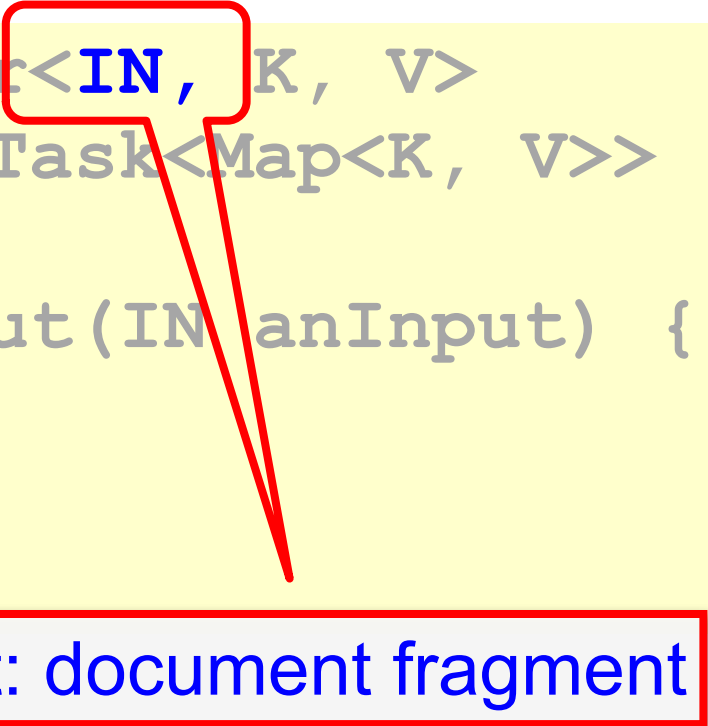
chapter 1

# Mapper Class

```
abstract class Mapper<IN, K, V>  
    extends RecursiveTask<Map<K, V>> {  
    IN input;  
    public void setInput(IN anInput) {  
        input = anInput;  
    }  
}
```

# Mapper Class

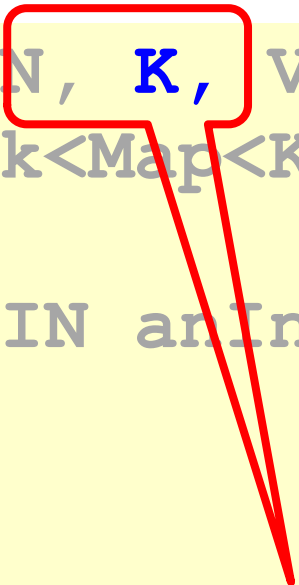
```
abstract class Mapper<IN, K, V>  
    extends RecursiveTask<Map<K, V>> {  
    IN input;  
    public void setInput(IN anInput) {  
        input = anInput;  
    }  
}
```



input: document fragment

# Mapper Class

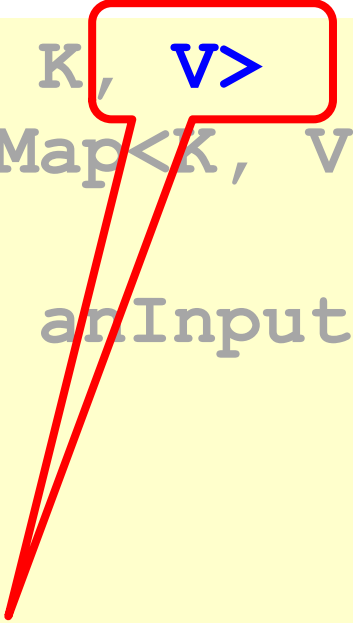
```
abstract class Mapper<IN, K, V>  
    extends RecursiveTask<Map<K, V>> {  
    IN input;  
    public void setInput(IN anInput) {  
        input = anInput;  
    }  
}
```

A red line originates from a red-bordered box around the 'K' in the code snippet and points down to another red-bordered box containing the text 'key: individual word'.

key: individual word

# Mapper Class

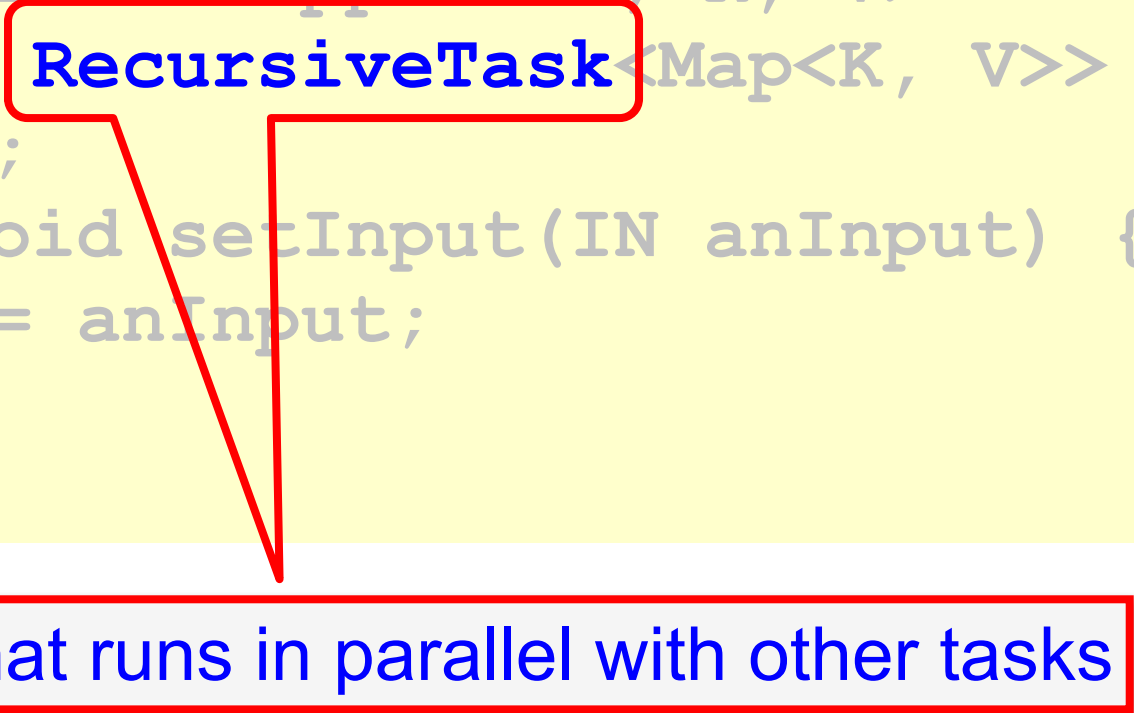
```
abstract class Mapper<IN, K, V>
    extends RecursiveTask<Map<K, V>> {
    IN input;
    public void setInput(IN anInput) {
        input = anInput;
    }
}
```



**value: local count**

# Mapper Class

```
abstract class Mapper<IN, K, V>  
    extends RecursiveTask<Map<K, V>> {  
    IN input;  
    public void setInput(IN anInput) {  
        input = anInput;  
    }  
}
```



a task that runs in parallel with other tasks



# Mapper Class

```
abstract class Mapper<IN, K, V>  
    extends RecursiveTask<Map<K, V>> {  
    IN input;  
    public void setInput(IN anInput) {  
        input = anInput;  
    }  
}
```

produces a map: word → count

# Mapper Class

```
abstract class Mapper<IN, K, V>  
    extends RecursiveTask<Map<K, V>> {  
    IN input;  
    public void setInput(IN anInput) {  
        input = anInput;  
    }  
}
```

initialize input: which document fragment?

# WordCount Mapper

```
class WordCountMapper extends  
  mapreduce.Mapper<  
    List<String>, String, Long  
  > {  
  ...  
}
```

# WordCount Mapper

```
class WordCountMapper extends  
  mapreduce.Mapper<  
    List<String>, String, Long  
  > {  
  ...  
}
```

document fragment is  
list of words

# WordCount Mapper

```
class WordCountMapper extends  
  mapreduce Mapper<  
    List<String>, String, Long  
  > {  
  ...
```

document fragment is  
list of words

map each word ...

# WordCount Mapper

```
class WordCountMapper extends  
  mapreduce Mapper<  
    List<String>, String, Long  
  > {  
  ...  
}
```

document fragment is  
list of words

map each word ...

to its count in  
the fragment

# WordCount Mapper

```
Map<String,Long> compute() {  
    Map<String,Long> map = new HashMap<>();  
    for (String word : input) {  
        map.merge(word,  
                    1L,  
                    (x, y) -> x + y);  
    }  
    return map;  
}
```

# WordCount Mapper

```
Map<String,Long> compute() {
```

```
    Map<String,Long> map = new HashMap<>();  
    for (String word : input) {  
        map.merge(word,  
            1L,  
            (x, y) -> x + y);  
    }
```

the `compute()` method constructs the local word count



# WordCount Mapper

```
Map<String,Long> compute() {  
    Map<String,Long> map = new HashMap<>();  
    for (String word : input) {  
        map.merge(word,  
            1L,  
            (x, y) -> x + y);  
    }  
    return map;  
}
```

**create a map to hold the output**

# WordCount Mapper

```
Map<String,Long> compute() {  
    Map<String,Long> map = new HashMap<>();  
    for (String word : input) {  
        map.merge(word,  
                  1L,  
                  (x, y) -> x + y);  
    }  
    return map;  
}
```

**examine each word in the document fragment**

# WordCount Mapper

```
Map<String,Long> compute() {  
    Map<String,Long> map = new HashMap<>();  
    for (String word : input) {  
        map.merge(word,  
                  1L,  
                  (x, y) -> x + y);  
    }  
    return map;  
}
```


increment that word's  
count in the map

# WordCount Mapper

```
Map<String,Long> compute() {  
    Map<String,Long> map = new HashMap<>();  
    for (String word : input) {  
        map.merge(word,  
                    1L,  
                    (x, y) -> x + y);  
    }  
    return map;  
}
```


when the local count is complete, return the map

# Reduce Phase




alpha → 4  
bravo → 2  
...  
zulu → 1

a *reducer* thread merges mapper outputs




alpha → 2  
juliet → 1  
tango → 1  
...



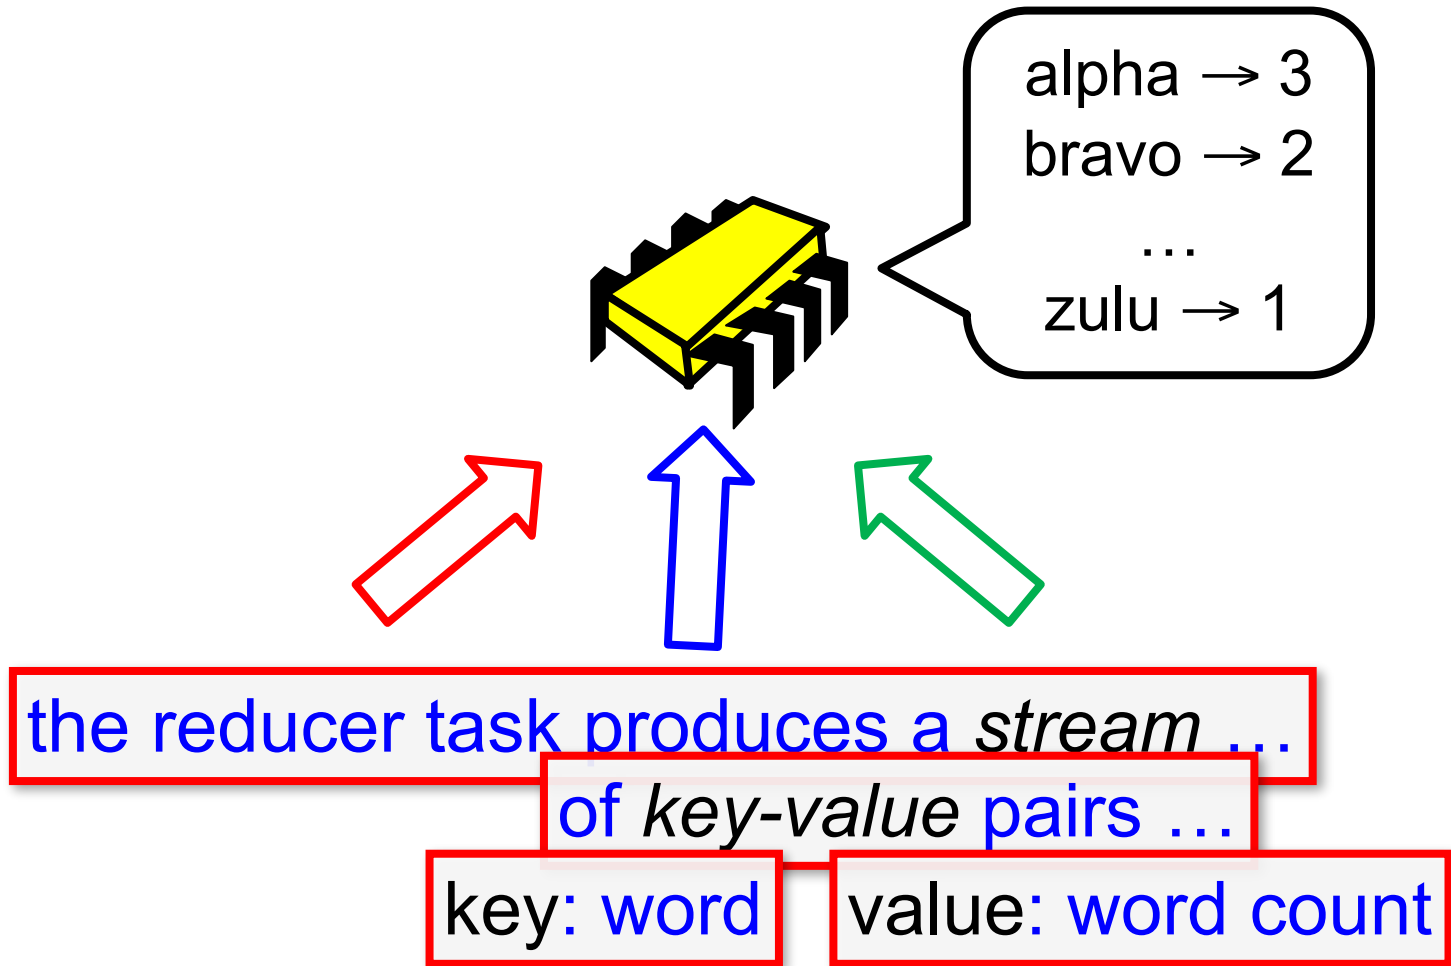
alpha → 1  
foxtrot → 1  
papa → 1  
tango → 1  
...

...



alpha → 1  
oscar → 1,  
bravo → 2...

# Reduce Phase



# Reducer Class

```
abstract class Reducer<K, V, OUT>
    extends RecursiveTask<OUT> {
    K key;
    List<V> valueList;
    public void setInput(
        K aKey,
        List<V> aList) {
        key = aKey;
        valueList = aList;
    }
}
```

# Reducer Class

```
abstract class Reducer<K, V, OUT>  
    extends RecursiveTask<OUT> {  
    K key;  
    List<V> valueList;  
    public void setInput(  
        K aKey,  
        List<V> aList) {  
        key = aKey;  
        valueList = aList;  
    }  
}
```

each reducer is given  
a single key (word)



# Reducer Class

```
abstract class Reducer<K, V, OUT>
    extends RecursiveTask<OUT> {
    K key;
    List<V> valueList;
    public void setInput(
        K aKey,
        List<V> aList) {
        key = aKey;
        valueList = aList;
    }
}
```



and a list of associated values  
(word count per fragment)

# Reducer Class

```
abstract class Reducer<K, V, OUT>
    extends RecursiveTask<OUT> {
    K key;
    List<V> valueList;
    public void setInput(
        K aKey,
        List<V> aList) {
        key = aKey;
        valueList = aList;
    }
}
```

It produces a single summary value  
(the total count for that word)

# WordCount


$$\begin{pmatrix} 0.037 \\ 0.002 \\ 0.045 \\ \dots \\ 0.000 \end{pmatrix}$$

normalizing document wordcount

gives a *fingerprint* vector

# Pseudocode Implementations

- See `m6/concurrent/mapreduce/MapReduce.java`
- See `m6/concurrent/mapreduce/WordCount.java`

# Barriers

# Simple Video Game

- Prepare frame for display
  - By graphics coprocessor
- “soft real-time” application
  - Need at least 35 frames/second
  - OK to mess up rarely

# Simple Video Game

```
while (true) {  
    frame.prepare();  
    frame.display();  
}
```

# Simple Video Game

```
while (true) {  
    frame.prepare() ;  
    frame.display() ;  
}
```

- What about overlapping work?
  - 1<sup>st</sup> thread displays frame
  - 2<sup>nd</sup> prepares next frame



# Two-Phase Rendering

```
while (true) {  
    if (phase) {  
        frame[0].display();  
    } else {  
        frame[1].display();  
    }  
    phase = !phase;  
}
```

```
while (true) {  
    if (phase) {  
        frame[1].prepare();  
    } else {  
        frame[0].prepare();  
    }  
    phase = !phase;  
}
```

# Two-Phase Rendering

```
while (true) {  
  if (phase) {  
    frame[0].display();  
  } else {  
    frame[1].display();  
  }  
  phase = !phase;  
}
```

```
while (true) {  
  if (phase) {  
    frame[1].prepare();  
  } else {  
    frame[0].prepare();  
  }  
  phase = !phase;  
}
```

**even phases**

# Two-Phase Rendering

```
while (true) {  
    if (phase) {  
        frame[0].display();  
    } else {  
        frame[1].display();  
    }  
    phase = !phase;  
}
```


```
while (true) {  
    if (phase) {  
        frame[1].prepare();  
    } else {  
        frame[0].prepare();  
    }  
    phase = !phase;  
}
```

**odd phases**


# Synchronization Problems

- How do threads stay in phase?
- Too early?
  - “we render no frame before its time”
- Too late?
  - Recycle memory before frame is displayed


# Ideal Parallel Computation




# Ideal Parallel Computation



# Real-Life Parallel Computation




# Real-Life Parallel Computation







# Barrier Synchronization



# Barrier Synchronization




# Barrier Synchronization



# Why Do We Care?


- Mostly of interest to
  - Scientific & numeric computation
- Elsewhere
  - Less common in systems programming
  - Still important topic

# Example: Parallel Prefix




**before**

**after**




# Parallel Prefix


**One thread  
Per entry**



# Parallel Prefix: Phase 1



# Parallel Prefix: Phase 2





# Parallel Prefix

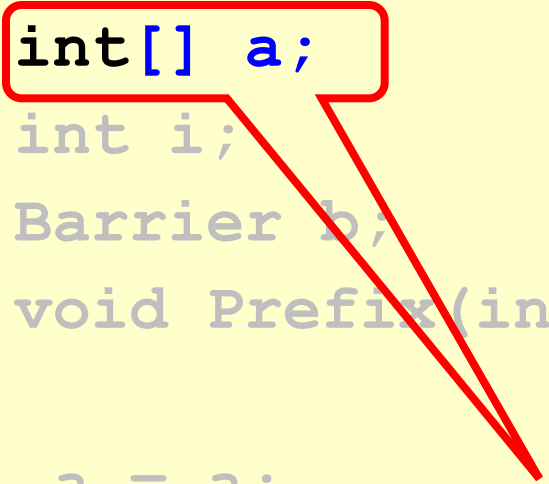
- N threads can compute
  - Parallel prefix
  - Of N entries
  - In  $\log_2 N$  rounds
- What if system is asynchronous?
  - Why we need barriers

# Prefix

```
class Prefix extends Thread {  
    int[] a;  
    int i;  
    Barrier b;  
    void Prefix(int[] a,  
                Barrier b, int i) {  
        a = a;  
        b = b;  
        i = i;  
    }  
}
```

# Prefix

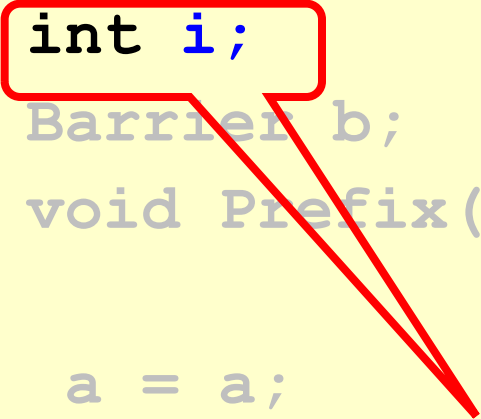
```
class Prefix extends Thread {  
    int[] a;  
    int i;  
    Barrier b;  
    void Prefix(int[] a,  
                Barrier b, int i) {  
        a = a;  
        b = b;  
        i = i;  
    }  
}
```



**Array of input values**

# Prefix

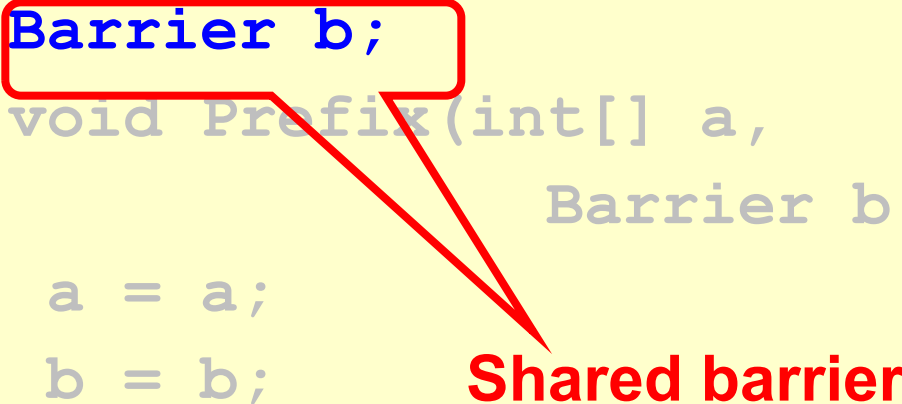
```
class Prefix extends Thread {  
    int[] a;  
    int i;  
    Barrier b;  
    void Prefix(int[] a,  
                Barrier b, int i) {  
        a = a;  
        b = b;  
        i = i;  
    }  
}
```



**Thread index**

# Prefix

```
class Prefix extends Thread {  
    int[] a;  
    int i;  
    Barrier b;  
    void Prefix(int[] a,  
                Barrier b, int i) {  
        a = a;  
        b = b;  
        i = i;  
    }  
}
```

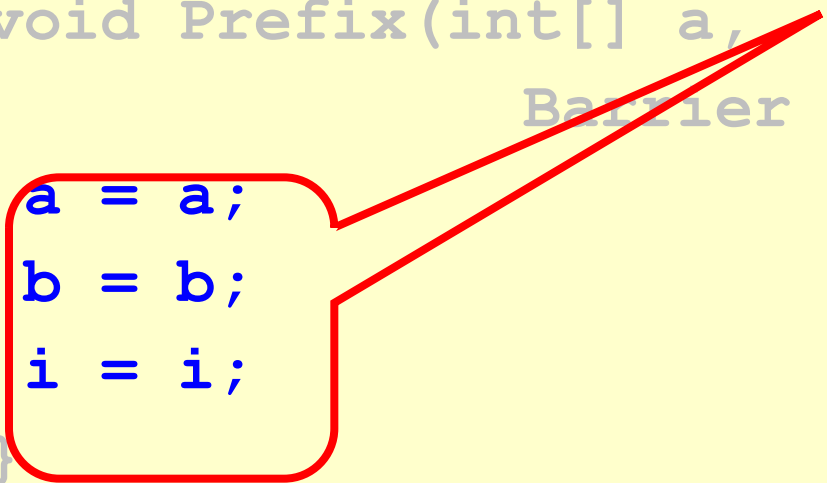


**Shared barrier**

# Prefix

```
class Prefix extends Thread {  
    int[] a;  
    int i;  
    Barrier b;  
    void Prefix(int[] a,  
                Barrier b, int i) {  
        a = a;  
        b = b;  
        i = i;  
    }  
}
```

**Initialize fields**



```
a = a;  
b = b;  
i = i;
```

# Where Do the Barriers Go?

```
public void run() {  
    int d = 1, sum = 0;  
    while (d < N) {  
        if (i >= d)  
            sum = a[i-d];  
        if (i >= d)  
            a[i] += sum;  
        d = d * 2;  
    }  
}
```

# Where Do the Barriers Go?

```
public void run() {  
    int d = 1, sum = 0;  
    while (d < N) {  
        if (i >= d)  
            sum = a[i-d];  
        b.await();  
        if (i >= d)  
            a[i] += sum;  
        d = d * 2;  
    }  
}
```

**Make sure everyone reads  
before anyone writes**



# Where Do the Barriers Go?

```
public void run() {  
    int d = 1, sum = 0;  
    while (d < N) {  
        if (i >= d)  
            sum = a[i-d];  
        b.await();  
        if (i >= d)  
            a[i] += sum;  
        b.await();  
        d = d * 2;  
    }  
}
```

**Make sure everyone reads  
before anyone writes**

**Make sure everyone writes  
before anyone reads**

# Barrier Implementations

- Cache coherence
  - Spin on locally-cached locations?
  - Spin on statically-defined locations?
- Latency
  - How many steps?
- Symmetry
  - Do all threads do the same thing?

# Barriers

```
public class Barrier {  
    AtomicInteger count;  
    int size;  
    public Barrier(int n) {  
        count = AtomicInteger(n);  
        size = n;  
    }  
    public void await() {  
        if (count.getAndDecrement() == 1) {  
            count.set(size);  
        } else {  
            while (count.get() != 0);  
        }  
    }  
}
```

# Barriers

```
public class Barrier {  
    AtomicInteger count;  
    int size;  
    public Barrier(int n) {  
        count = AtomicInteger(n);  
        size = n;  
    }  
    public void await() {  
        if (count.getAndDecrement() == 1) {  
            count.set(size);  
        } else {  
            while (count.get() != 0);  
        }  
    }  
}
```

**Number of threads  
not yet arrived**

# Barriers

```
public class Barrier {  
    AtomicInteger count;  
    int size;  
    public Barrier(int n) {  
        count = AtomicInteger(n);  
        size = n;  
    }  
    public void await() {  
        if (count.getAndDecrement() == 1) {  
            count.set(size);  
        } else {  
            while (count.get() != 0);  
        }  
    }  
}
```

**Number of threads participating**

# Barriers

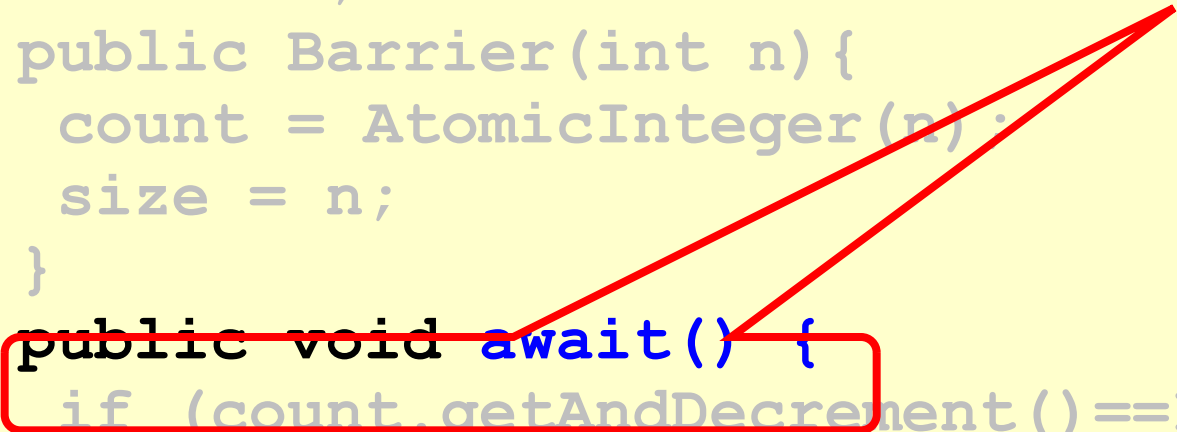
```
public class Barrier {  
    AtomicInteger count;  
    int size;  
    public Barrier(int n) {  
        count = AtomicInteger(n);  
        size = n;  
    }  
    public void await() {  
        if (count.getAndDecrement() == 1) {  
            count.set(size);  
        } else {  
            while (count.get() != 0);  
        }  
    }  
}
```

**Initialization**

# Barriers

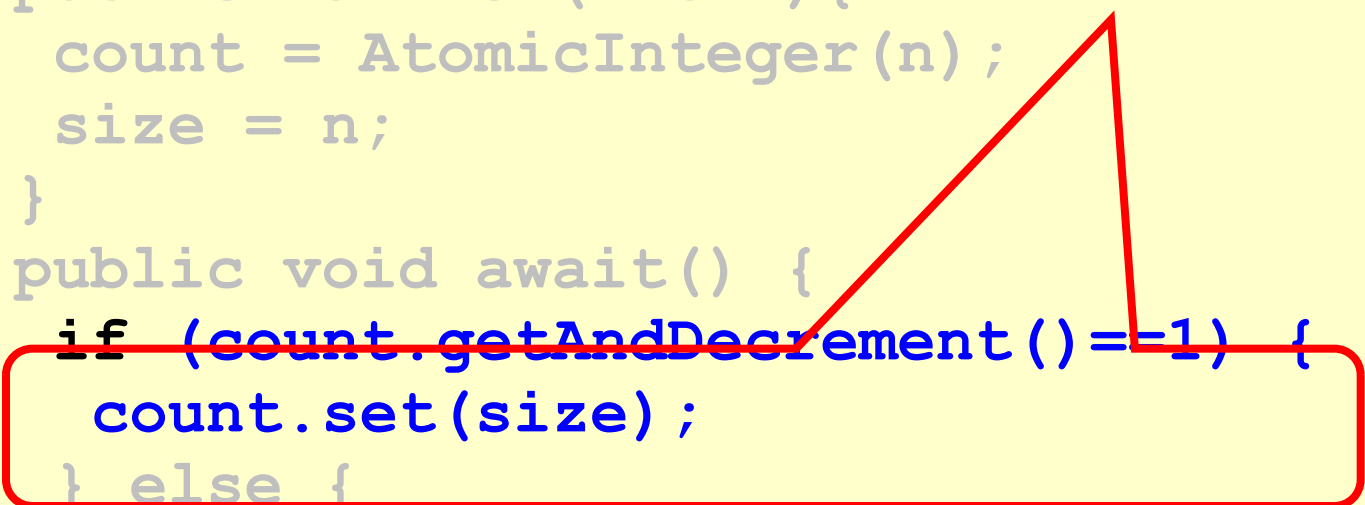
```
public class Barrier {  
    AtomicInteger count;  
    int size;  
    public Barrier(int n) {  
        count = AtomicInteger(n);  
        size = n;  
    }  
    public void await() {  
        if (count.getAndDecrement() == 1) {  
            count.set(size);  
        } else {  
            while (count.get() != 0);  
        }  
    }  
}
```

**Principal method**



# Barriers

```
public class Barrier {  
    AtomicInteger count; If I'm last, reset fields  
    int size; for next time  
    public Barrier(int n) {  
        count = AtomicInteger(n);  
        size = n;  
    }  
    public void await() {  
        if (count.getAndDecrement() == 1) {  
            count.set(size);  
        } else {  
            while (count.get() != 0);  
        }}}}
```

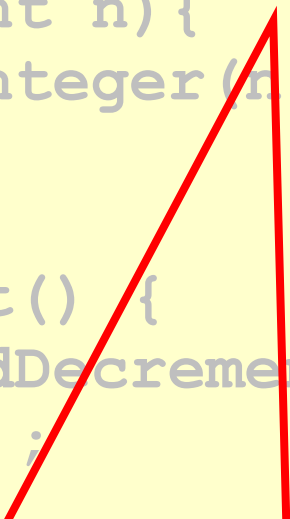




# Barriers

```
public class Barrier {
    AtomicInteger count;
    int size;
    public Barrier(int n) {
        count = AtomicInteger(n);
        size = n;
    }
    public void await() {
        if (count.getAndDecrement() == 1) {
            count.set(size);
        } else {
            while (count.get() != 0);
        }
    }
}
```

Otherwise, wait for everyone else



# Barriers

```
public class Barrier {  
    AtomicInteger count;  
    int size;  
    public Barrier(int n) {  
        count = AtomicInteger(n);  
        size = n;  
    }  
    public void await() {  
        if (count.getAndDecrement() == 1) {  
            count.set(size);  
        } else {  
            while (count.get() != 0);  
        }  
    }  
}
```

**What's wrong with this protocol?**

# Reuse

```
Barrier b = new Barrier(n);
```

```
while ( mumble() ) {
```

```
    work();
```

```
    b.await();
```

```
}
```

do work

synchronize

repeat


# Barriers

```
public class Barrier {  
    AtomicInteger count;  
    int size;  
    public Barrier(int n) {  
        count = AtomicInteger(n);  
        size = n;  
    }  
    public void await() {  
        if (count.getAndDecrement() == 1) {  
            count.set(size);  
        } else {  
            while (count.get() != 0);  
        }  
    }  
}
```

# Barriers

```
public class Barrier {  
    AtomicInteger count;  
    int size;  
    public Barrier(int n) {  
        count = AtomicInteger(n);  
        size = n;  
    }  
    public void await() {  
        if (count.getAndDecrement() == 1) {  
            count.set(size);  
        } else {  
            while (count.get() != 0);  
        }  
    }  
}
```

Waiting for  
Phase 1 to finish



# Barriers

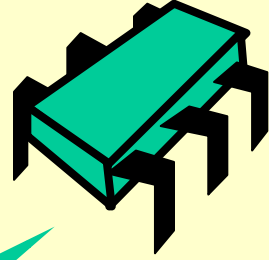
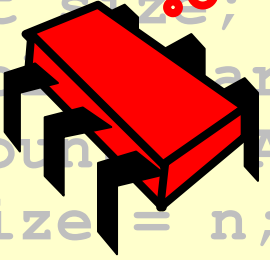
```
public class Barrier {
    AtomicInteger count;
    int size;

    public Barrier(int n) {
        count = new AtomicInteger(n);
        size = n;
    }

    public void await() {
        if (count.getAndDecrement() >= 1) {
            count.set(size);
        } else {
            while (count.get() != 0);
        }
    }
}
```

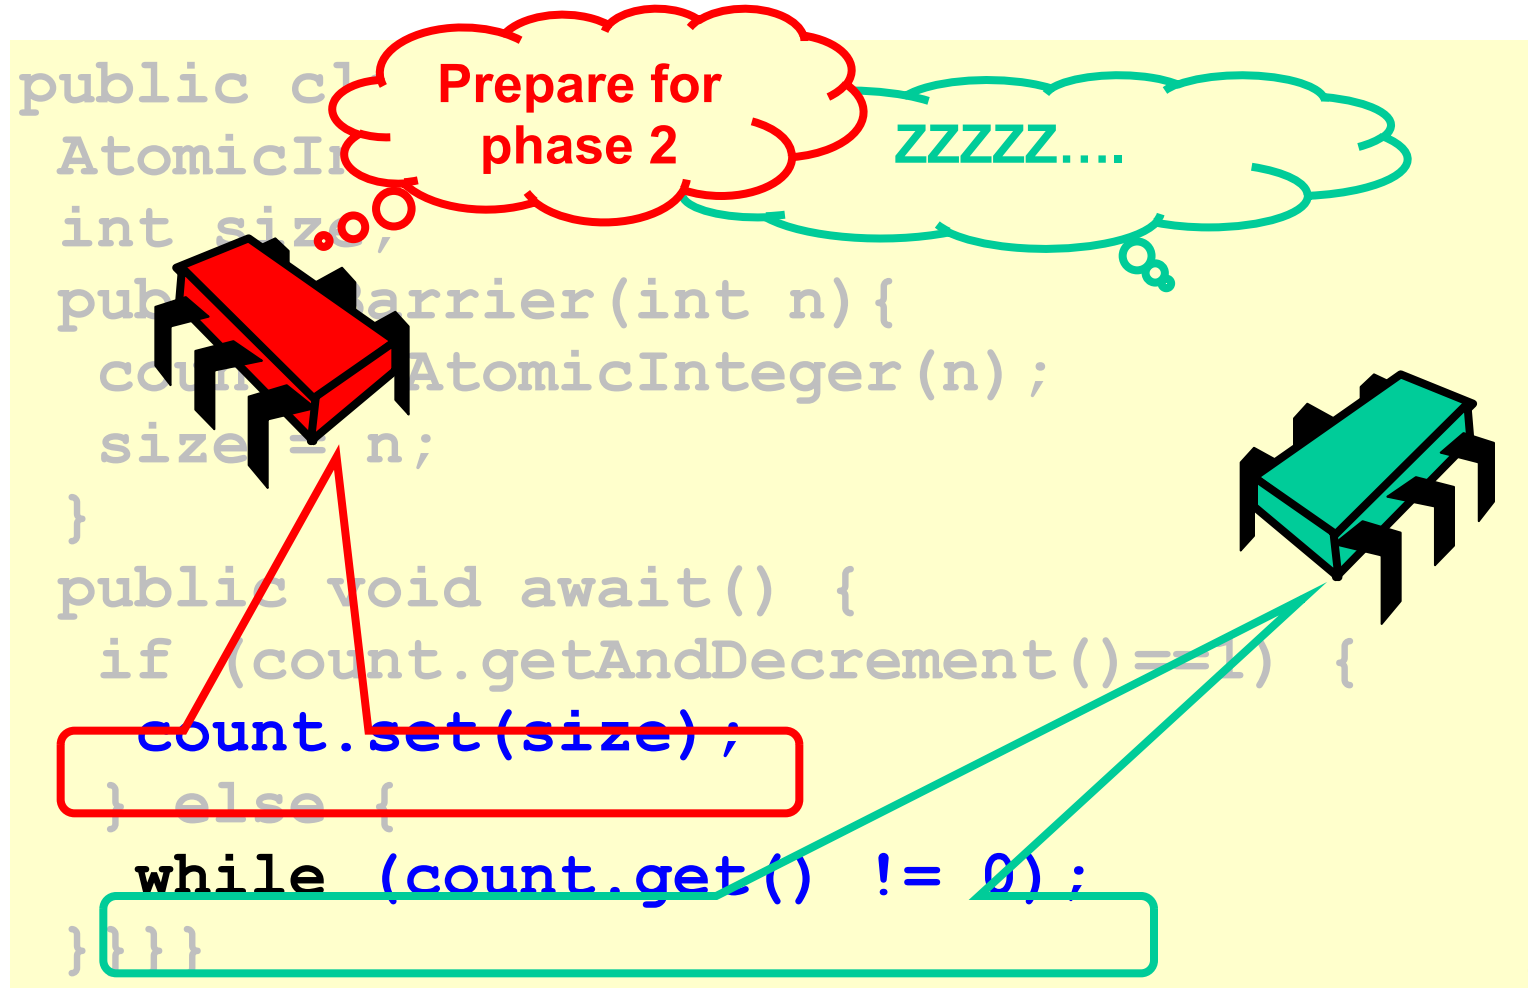
Phase 1 is so over

Waiting for Phase 1 to finish



The diagram illustrates the execution of a barrier. A red block (Phase 1) is shown with a red thought bubble indicating it is over. A green block (Phase 2) is shown with a green thought bubble indicating it is waiting for Phase 1 to finish. The code snippet shows the implementation of the barrier, with the 'if' and 'while' loops highlighted by red and green boxes respectively.

# Barriers



# Uh-Oh

```
public class Barrier {  
    AtomicInteger count;  
    int size;  
    public Barrier(int n) {  
        count = new AtomicInteger(0);  
        size = n;  
    }  
    public void await() {  
        if (count.getAndDecrement() == 1) {  
            count.set(size);  
        } else {  
            while (count.get() != 0);  
        }  
    }  
}
```

Waiting for  
Phase 2 to finish

Waiting for  
Phase 1 to finish

Oo.



# Basic Problem

- One thread “wraps around” to start phase 2
- While another thread is still waiting for phase 1
- One solution:
  - Always use two barriers

# Sense-Reversing Barriers

```
public class Barrier {
    AtomicInteger count;
    int size;
    volatile boolean sense = false;
    threadSense = new ThreadLocal<boolean>...

    public void await {
        boolean mySense = threadSense.get();
        if (count.getAndDecrement()==1) {
            count.set(size); sense = mySense
        } else {
            while (sense != mySense) {}
        }
        threadSense.set(!mySense) } } }
```

# Sense-Reversing Barriers

Completed odd or  
even-numbered  
phase?

```
public class Barrier {
    AtomicInteger count;
    int size;
    volatile boolean sense = false;
    ThreadLocal<boolean> threadSense = new ThreadLocal<boolean>();

    public void await () {
        boolean mySense = threadSense.get();
        if (count.getAndDecrement() == 1) {
            count.set(size); sense = mySense;
        } else {
            while (sense != mySense) {}
        }
        threadSense.set(!mySense);
    }
}
```

# Sense-Reversing Barriers

```
public class Barrier {  
    AtomicInteger count;  
    int size;  
    volatile boolean sense = false;  
    threadSense = new ThreadLocal<boolean>...  
  
    public void await {  
        boolean mySense = threadSense.get();  
        if (count.getAndDecrement()==1) {  
            count.set(size); sense = mySense  
        } else {  
            while (sense != mySense) {}  
        }  
        threadSense.set(!mySense) } } }
```

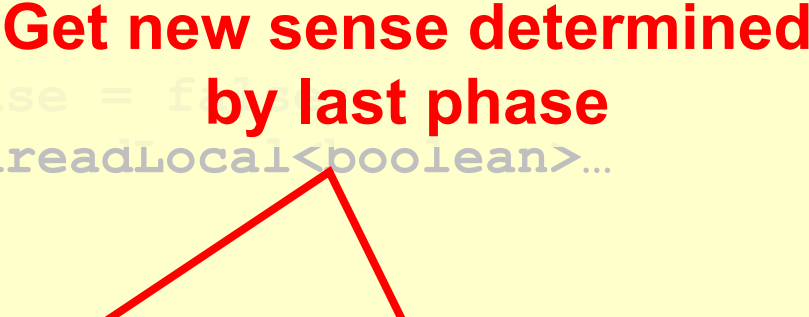
**Store sense for next phase**

# Sense-Reversing Barriers

```
public class Barrier {
    AtomicInteger count;
    int size;
    volatile boolean sense = false;
    ThreadLocal<boolean> threadSense = new ThreadLocal<boolean>();

    public void await {
        boolean mySense = threadSense.get();
        if (count.getAndDecrement()--1) {
            count.set(size); sense = mySense
        } else {
            while (sense != mySense) {}
        }
        threadSense.set(!mySense) } } }
```

**Get new sense determined by last phase**

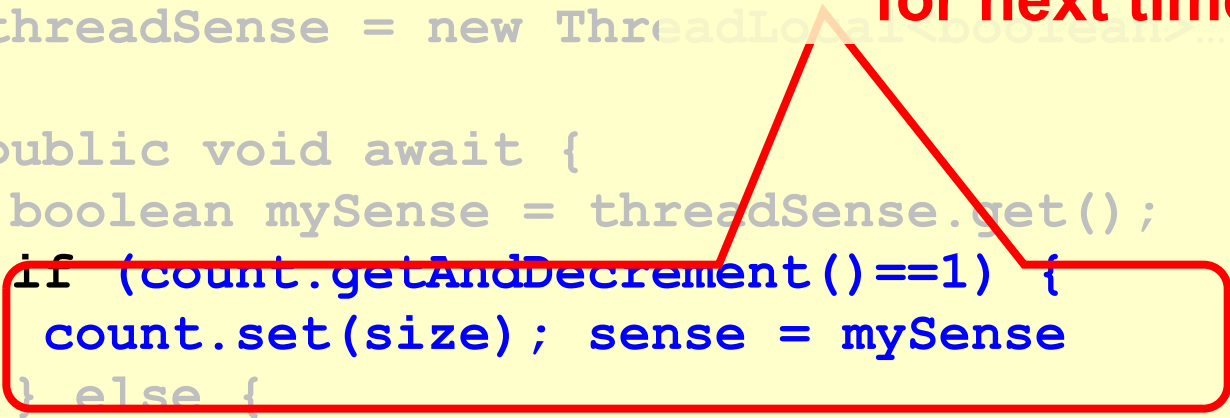


# Sense-Reversing Barriers

```
public class Barrier {
    AtomicInteger count;
    int size;
    volatile boolean sense = false;
    threadSense = new ThreadLocal<boolean>();

    public void await {
        boolean mySense = threadSense.get();
        if (count.getAndDecrement() == 1) {
            count.set(size); sense = mySense
        } else {
            while (sense != mySense) {}
        }
        threadSense.set(!mySense) } } }
```

**If I'm last, reverse sense  
for next time**



# Sense-Reversing Barriers

```
public class Barrier {  
    AtomicInteger count;  
    int size;  
    volatile boolean sense = false;  
    threadSense = new ThreadLocal<boolean>...
```

Otherwise, wait for  
sense to flip

```
public void await {  
    boolean mySense = threadSense.get();  
    if (count.getAndDecrement()==1) {  
        count.set(size); sense = mySense  
    } else {  
        while (sense != mySense) {}  
    }  
    threadSense.set(!mySense) } }
```

# Sense-Reversing Barriers

```
public class Barrier {  
    AtomicInteger count;  
    int size;  
    volatile boolean sense = false;  
    ThreadLocal<boolean> threadSense = new ThreadLocal<boolean>();  
  
    public void await () {  
        boolean mySense = threadSense.get();  
        if (count.getAndDecrement() == 1) {  
            count.set(size); sense = mySense;  
        } else {  
            while (sense != mySense) {}  
        }  
        threadSense.set(!mySense);  
    }  
}
```

**Prepare sense for  
next phase**

